

IPC

Message Queues

- Interprocess communication
 - pipes
 - signals
 - very restrictive
 - need better IPC mechanisms
 - message queues
 - semaphores
 - shared memory

- Interprocess Communication (IPC)

- message queues

- “selective write”
- “selective read”
- FIFO queue
- data structure maintained by kernel

- semaphores

- synchronization
- “communicate small amounts of data”

- shared memory

- communication via a shared data space
- random access

-unrelated processes *on the same system*

- IPC resources
 - creator
 - owner
 - access permissions
 - established upon creation
 - can be modified via system calls

ipcs utility

```
linux# ipcs
```

```
----- Shared Memory Segments -----
```

key	shmid	owner	perms	bytes	nattch	status
0x00000000	25198594	root	666	247264	3	


```
----- Semaphore Arrays -----
```

key	semid	owner	perms	nsems	status
0x00000000	65537	root	666	4	
0x00000000	98306	root	666	16	
0x00000000	131075	root	666	16	
0x00000000	163844	root	666	16	

```
----- Message Queues-----
```

key	msqid	owner	perms	used-bytes	messages
-----	-------	-------	-------	------------	----------


One shared memory segment attached (shared) by three processes



Four sets of semaphores all owned by root



No message queues are currently allocated



- **ipcs utility**

```
ipcs -s -l 622592
```

```
Semaphore Array semid=622592
```

```
uid=33   gid=33   cuid=0   cgid=0
```

```
mode=0600, access_perms=0600
```

```
nsems = 1
```

```
otime = Mon Sep 26 10:49:00 2005
```

```
ctime = Sun Sep 25 20:12:11 2005
```

semnum	value	ncount	zcount	pid
0	0	9	0	24077

-a	all
-m	shared memory
-q	message queues
-s	semaphores

-c	creator
-l	limits
-p	process id
-t	time
-u	summary

- limits
 - set when kernel is generated

```
# ipcs -l
----- Shared Memory Limits -----
max number of segments = 4096
max seg size (kbytes) = 32768
max total shared memory (pages) = 2097152
min seg size (bytes) = 1

----- Semaphore Limits -----
max number of arrays = 128
max semaphores per array = 250
max semaphores system wide = 32000
max ops per semop call = 32
semaphore max value = 32767

----- Messages: Limits -----
max queues system wide = 16
max size of message (bytes) = 8192
default max size of queue (bytes) = 16384
```

- **ipcrm utility**
 - **remove resource (owner)**

-M shmkey

Mark the shared memory segment associated with key shmkey for removal. This marked segment will be destroyed after the last detach.

-m shmid

Mark the shared memory segment associated with id shmid for removal. This marked segment will be destroyed after the last detach.

-Q msgkey

Remove the message queue associated with key msgkey from the system.

-q msqid

Remove the message queue associated with the id msqid from the system.

-S semkey

Remove the semaphore set associated with key semkey from the system.

-s semid

Removes the semaphore set associated with id semid from the system.

- **ipcs system calls**

Functionality	Message queue	Semaphore	Shared memory
Allocate IPC resource; gain access to an existing IPC resource	msgget	semget	shmget
Control an IPC resource: obtain/modify status information, remove the resource	msgctl	semctl	shmctl
IPC operations: send/receive messages, perform semaphore operations, attach/free a shared memory segment	msgsnd msgrcv	semop	shmat shmdt

```
struct ipc_perm {
    __key_t    __key          /* Key */
    __uid_t    uid;           /* Owner's user ID. */
    __gid_t    gid;           /* Owner's group ID */
    __uid_t    cuid;          /* Creator's user ID */
    __gid_t    cgid;          /* Creator's group ID */
    unsigned short int mode;   /* Access permissions */
    unsigned short int __pad1;
    unsigned short int __seg;  /* Sequence number */
    unsigned short int __pad2;
    unsigned long int __unused1;
    unsigned long int __unused2;
};
```

- get system calls

- msgget
 - semget
 - shmget
-
- allocate new resource or gain access to existing resource
 - permission
 - returns ipc identifier
 - used to reference the resource
 - index to IPC permission structure

```
struct ipc_perm {  
    __key_t    __key        /* Key                */  
    __uid_t    uid;         /* Owner's user ID.  */  
    __gid_t    gid;         /* Owner's group id   */
```

- common arguments
 - key
 - used to generate an IPC identifier
 - one-to-one relation
 - **ftok** library function
 - unrelated processes can generate same key

```
key_t ftok(char *pathname, char proj);
```

- common arguments
 - `IPC_PRIVATE`
 - create IPC resource with unique IPC identifier
 - no other process creating/accessing an IPC resource will have same IPC identifier
 - in related processes
 - parent creates IPC resource
 - child inherits
 - exec pass as argument or environment variable
 - non-related processes
 - server creates
 - send key to clients

- Common arguments (cont.)
 - message flag
 - Access permissions
 - Lower nine bits of flag used for permissions

Table 6.4 Required Permissions for IPC System Calls.

Permissions Required	Message Queues	Semaphores	Shared Memory
write (alter)	msgsnd place message in the queue	semop increase or decrease a semaphore value	shmat to write to the shared memory segment
	msgctl write out modified IPC status information	semctl set the value of one semaphore or a whole set; write out modified IPC status information	shmctl write out modified IPC status information
read	msgrcv obtain message from queue	semop block until a semaphore becomes 0	shmat read from the shared memory segment
	msgctl to retrieve IPC status information	semctl to retrieve IPC status information	shmctl to retrieve IPC status information

- flags can be ORed
 - `IPC_CREAT`
 - create if it does not exists
 - if present and not created with `IPC_PRIVATE` returns IPC identifier
 - `IPC_CREAT` | `IPC_EXCL`
 - create exclusive
 - fails if resource exists

–ctl system calls

- msgctl
- semctl
- shmctl

- act upon permissions

- IPC_STAT
 - status information
- IPC_SET
 - set owner/group/mode
- IPC_RMID
 - destroy and remove

IPC_SET and IPC_RMID only by owner, creator or superuser

– operations system calls

- `msgsnd/msgrcv`
 - blocks on write-full read-empty
 - until `signal` received, can read/write or resource removed
 - overwrite with `IPC_NOWAIT`
- `semop`
 - semaphore operations
 - sets and test semaphore values
 - blocks when attempt to decrement a semaphore currently at 0 or waiting for a semaphore to become 0
- `shmat/shmdt`
 - map/attach shared memory blocks
 - unmap/detach shared memory blocks
 - non-blocking

- message queues
 - Creation

```
int msgget (key_t key, int msgflg);
```

returns:

nonnegative integer queue identifier associated with key
used to reference the message queue

key:

- specified directly
- using `ftok`
- used to produce unique identifier

— msgflg

- low order bits
 - access permissions
- additional flags ORed
 - IPC_CREAT
 - IPC_EXCL
- new message queue is created
 - IPC_PRIVATE used as key
 - use IPC_CREAT flag and no queue exists
 - returns `id` if queue exists and IPC_EXCL was not specified

● Example

```
/* Message queue generation */
#define _GNU_SOURCE
#include <stdio>
#include <unistd.h>
#include <linux/limits.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
using namespace std;
const int MAX=5;
int
main( ){
    FILE *fin;
    char  buffer[PIPE_BUF], proj = 'A';
    int   i, n, mid[MAX];
    key_t key;
    for (i = 0; i < MAX; ++i, ++proj) {
        key = ftok(".", proj);
        if ((mid[i] = msgget(key, IPC_CREAT | 0660)) == -1) {
            perror("Queue create");
            return 1;
        }
    }
    fin = popen("ipcs", "r");
    while ((n = read(fileno(fin), buffer, PIPE_BUF)) > 0)
        write(fileno(stdout), buffer, n);
    pclose(fin);
    for (i = 0; i < MAX; ++i )
        msgctl(mid[i], IPC_RMID, (struct msqid_ds *) 0);
    return 0;
}
```

```
gcos:test dcanas$ ./msg
```

```
IPC status from <running system> as of Tue Sep 25 13:29:53 EDT 2012
```

T	ID	KEY	MODE	OWNER	GROUP
---	----	-----	------	-------	-------

Message Queues:

q	262144	0x410297f6	--rw-rw----	dcanas	staff
q	131073	0x420297f6	--rw-rw----	dcanas	staff
q	131074	0x430297f6	--rw-rw----	dcanas	staff
q	131075	0x440297f6	--rw-rw----	dcanas	staff
q	131076	0x450297f6	--rw-rw----	dcanas	staff

T	ID	KEY	MODE	OWNER	GROUP
---	----	-----	------	-------	-------

Shared Memory:

m	393216	0x53414e44	--rw-rw-rw-	dcanas	staff
---	--------	------------	-------------	--------	-------

T	ID	KEY	MODE	OWNER	GROUP
---	----	-----	------	-------	-------

Semaphores:

```
gcos:test dcanas$ ipcs
```

```
IPC status from <running system> as of Tue Sep 25 13:32:28 EDT 2012
```

T	ID	KEY	MODE	OWNER	GROUP
---	----	-----	------	-------	-------

Message Queues:

T	ID	KEY	MODE	OWNER	GROUP
---	----	-----	------	-------	-------

Shared Memory:

m	393216	0x53414e44	--rw-rw-rw-	dcanas	staff
---	--------	------------	-------------	--------	-------

T	ID	KEY	MODE	OWNER	GROUP
---	----	-----	------	-------	-------

Semaphores:

- Creation
 - system message-queue data structure
 - msqid_ds
 - maintained by system

```
struct msqid_ds {
    struct ipc_perm msg_perm;
    struct msg *msg_first;           /* first message on queue, unused */
    struct msg *msg_last;           /* last message in queue, unused */
    __kernel_time_t msg_stime;      /* last msgsnd time */
    __kernel_time_t msg_rtime;      /* last msgrcv time */
    __kernel_time_t msg_ctime;      /* last change time */
    unsigned long msg_lbytes;        /* Reuse junk fields for 32 bit */
    unsigned long msg_lqbytes;      /* ditto */
    unsigned short msg_cbytes;       /* current # of bytes on queue */
    unsigned short msg_qnum;         /* number of messages in queue */
    unsigned short msg_qbytes;       /* max number of bytes on queue */
    __kernel_ipc_pid_t msg_lspid;    /* pid of last msgsnd */
    __kernel_ipc_pid_t msg_lrpid;    /* last receive pid */
};
```

- creation
 - set user/group id
 - `msg_first/last` point to first/last message in queue
 - linked list of messages

```
struct msg {  
    struct msg    *msg_next;    /* ptr to next message on q */  
    long          msg_type;     /* message type */  
    ushort       msg_ts;       /* message text size */  
    short        msg_spot;      /* address of text message */  
};
```

- Message Queue structure




Figure 6.5 A message queue with N items.

- Message Queue Control

```
int msgctl( int msqid, int cmd, struct msqid_s *buf);
```

- IPC_STAT

- returns current values of msqid_ds in buf

- IPC_SET

- modify values

- IPC_RMID

- removes message queue

```
/* Display message queue status information */
#include <iostream>
#include <cstdio>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
using namespace std;
int main () {
    int      mid;
    key_t    key;
    struct msqid_ds buf;
    key = ftok(".", 'z');
    if ((mid = msgget(key, IPC_CREAT | 0660 )) == -1) {
        perror("Queue create:");
        return 1;
    }
    msgctl(mid, IPC_STAT, &buf);
    cout << "Message Queue *Permissions*  Structure information" << endl;
    cout << "Owner's user ID      \t" << buf.msg_perm.uid << endl;
    cout << "Owner's group ID      \t" << buf.msg_perm.gid << endl;
    cout << "Creator's user ID   \t" << buf.msg_perm.cuid << endl;
    cout << "Creator's group ID  \t" << buf.msg_perm.cgid << endl;
    cout << "Access mode in HEX\t" << hex << buf.msg_perm.mode << endl;
    cout << "\nAdditional Selected Message Queue Structure Information\n";
    cout << "Current # of bytes on queue      \t" << dec << buf.msg_cbytes << endl;
    cout << "Current # of messages on queue  \t" << dec << buf.msg_qnum << endl;
    cout << "Maximum # of bytes on queue     \t" << dec << buf.msg_qbytes << endl;
    msgctl(mid, IPC_RMID, (struct msqid_ds *) 0 );
    return 0;
}
```

Message Queue *Permissions* Structure information

Owner's user ID	501
Owner's group ID	501
Creator's user ID	501
Creator's group ID	501
Access mode in HEX	1b0

1b0 => 0001 1011 0000 => 000 110 110 000 Permission (lower nine bits) => 660
--

Additional Selected Message Queue Structure Information

Current # of bytes on queue	0
Current # of messages on queue	0
Maximum # of bytes on queue	16384

● Message Queue Operations

- send and receive messages

```
struct msgbuf {  
    long int mtype;          /* type of message */  
    char mtext[1];          /* text of message */  
};
```

— mtype

- type of message
 - long integer
 - used by **msgrcv** to selectively retrieve messages

— mtext

- body of the message
 - any valid structure
 - long integer
 - followed by 0 or more bytes

- sending messages

```
int msgsnd( int msqid, struct msgbuf *msgp,  
            size_t msgsz, int msgflg );
```

- msgp

- pointer to message

- msgsz

- size of message
 - size of message structure - **size** of message type

- msgflg

- actions if limits reached

- **IPC_NOWAIT**

- return and message not sent

- 0

- block

- not beyond limits, receive a signal, message queue removed

- receiving messages

```
int msgrcv( int msqid, struct msgbuf *msgp,  
            size_t msgsz, long msgtyp, int msgflg );
```

- msgp

- pointer to receiving buffer
 - long integer
 - text

- msgsz

- maximum size of message

- msgflg

- actions (if message not in queue or message too large)
 - **IPC_NOWAIT**
 - message type requested not in queue
 - **MSG_NOERROR**
 - silently truncate
 - **MSG_EXCEPT**

– msgtyp

- type of message

msgtyp	Action
0	retrieves the first message of any msgtyp
> 0	retrieves the first message equal to msgtyp if MSG_EXCEPT is not specified. If MSG_EXCEPT is specified, the first message that is not equal to msgtyp
< 0	retrieve the first message of the lowest type less than or equal to absolute value of msgtyp

- Example

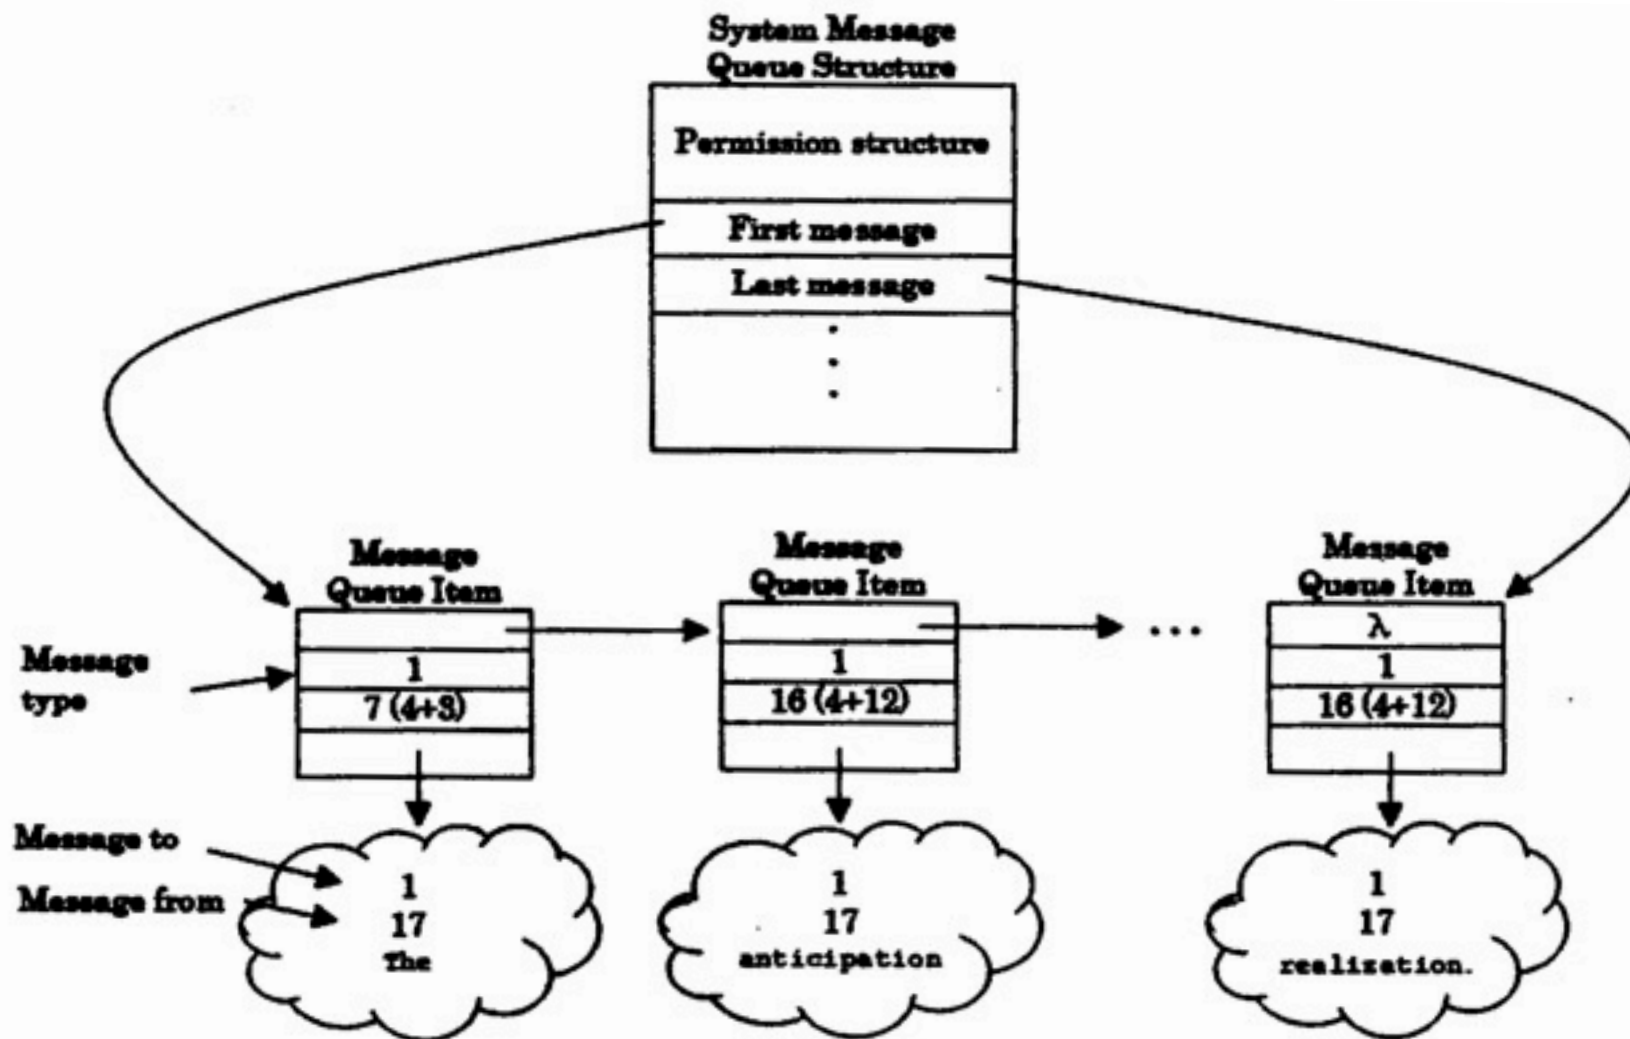


Figure 6.7 Conceptual view of message queue after the client has sent all seven messages

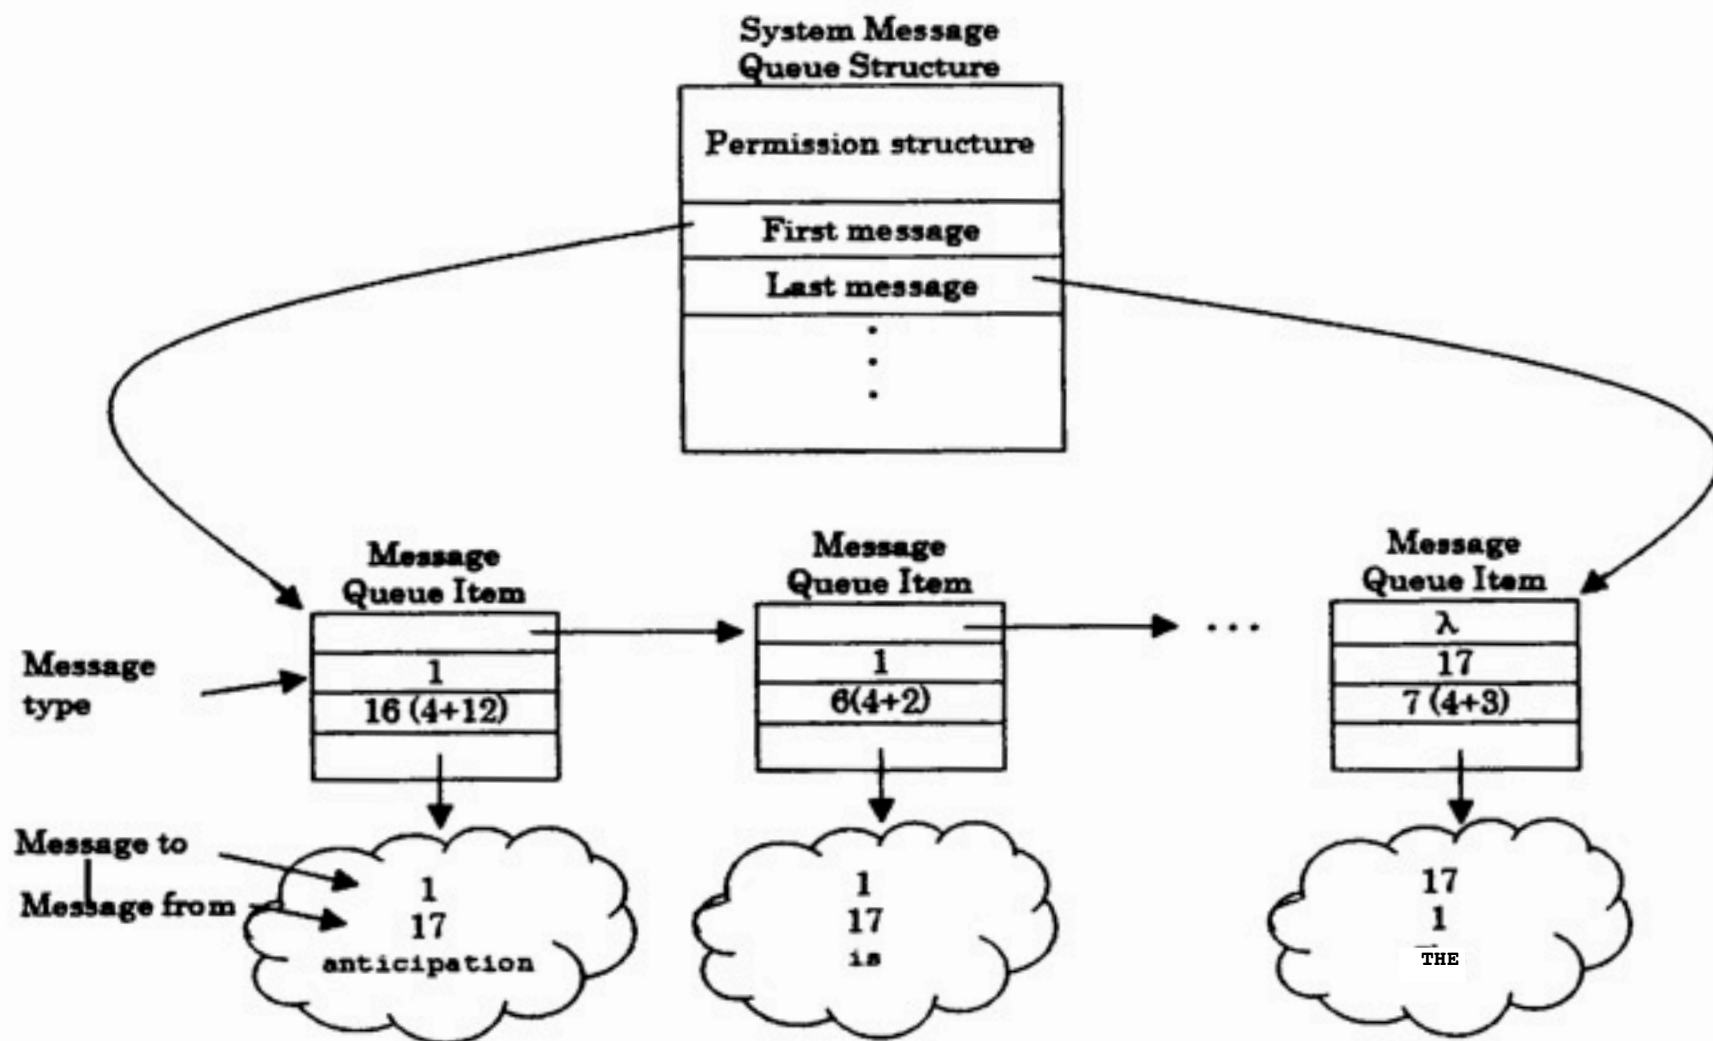


Figure. 6.8 Conceptual view of message queue after the first client message has been processed.