



Semaphores



- Semaphores
 - Synchronization
 - Avoid starvation
 - Avoid deadlock
 - Atomic
 - Integer value
 - Resource counter



synchronize operations on non-shareable resources

- atomic operation
- avoids busy wait
- two operations
 - P & V
 - on non-negative variables
 - semaphores



Edsger W. Dijkstra
1930–2002

Dijkstra, E. W. (March 1968). "Letters to the editor: go to statement considered harmful". *Communications of the ACM* **11** (3): 147–148. doi:[10.1145/362929.362947](https://doi.org/10.1145/362929.362947). ISSN 0001-0782.

- P & V
 - only valid operations on a semaphore variable
 - operations on same semaphore are mutually exclusive
 - sequential effect in no particular order
 - any waiting process may be selected

P(Sem) :

- if possible decrements Sem by 1
- else process must wait until Sem > 0

V(Sem) :

- if there is a process waiting -> **release**
- else increment Sem by 1

- Behavior

```
syscall P(int Sem)
```

```
{  
    if ( Sem > 0 ) Sem--;  
    else block on Sem;  
}
```

```
syscall V(int Sem)
```

```
{  
    if (process waiting on Sem) wake_up a process;  
    else Sem++;  
}
```

```
process p1 ()
{
    while(true)
    {
        non critical code;
        P(Sem);
        Critical Section;
        V(Sem);
        non critical code;
    }
}
process p2 ()
{
    while(true)
    {
        non critical code;
        P(Sem);
        Critical Section;
        V(Sem);
        non critical code;
    }
}
```

- Binary semaphores

- may only have values of 0 or 1
 - 1 semaphore NOT BUSY
 - 0 semaphore IS BUSY

- Counting semaphores
- Make take any integer value
- $\text{sem} > 0$
 - number of processes that may still continue execution after a $P(\text{sem})$ operation
- $\text{sem} == 0$
 - no more processes may continue after the execution of a $P(\text{sem})$ operations
- $\text{sem} < 0$
 - number of processes waiting on semaphore sem

- Counting semaphores

```
P(int Sem)
{
    Sem--;
    if (Sem < 0 ) block on Sem;
}
```

```
V(int Sem)
{
    Sem++;
    if ( Sem <= 0 ) wake up a process;
}
```

- **System V implementation**
 - A *set* of nonnegative integer
 - System defined maximum
 - Maintained by the kernel
- **Creation**
 - unique data structure
 - created/maintained by kernel
 - implemented as an integer
 - Stored in kernel space
 - Accessed through system calls
 - *Atomic*

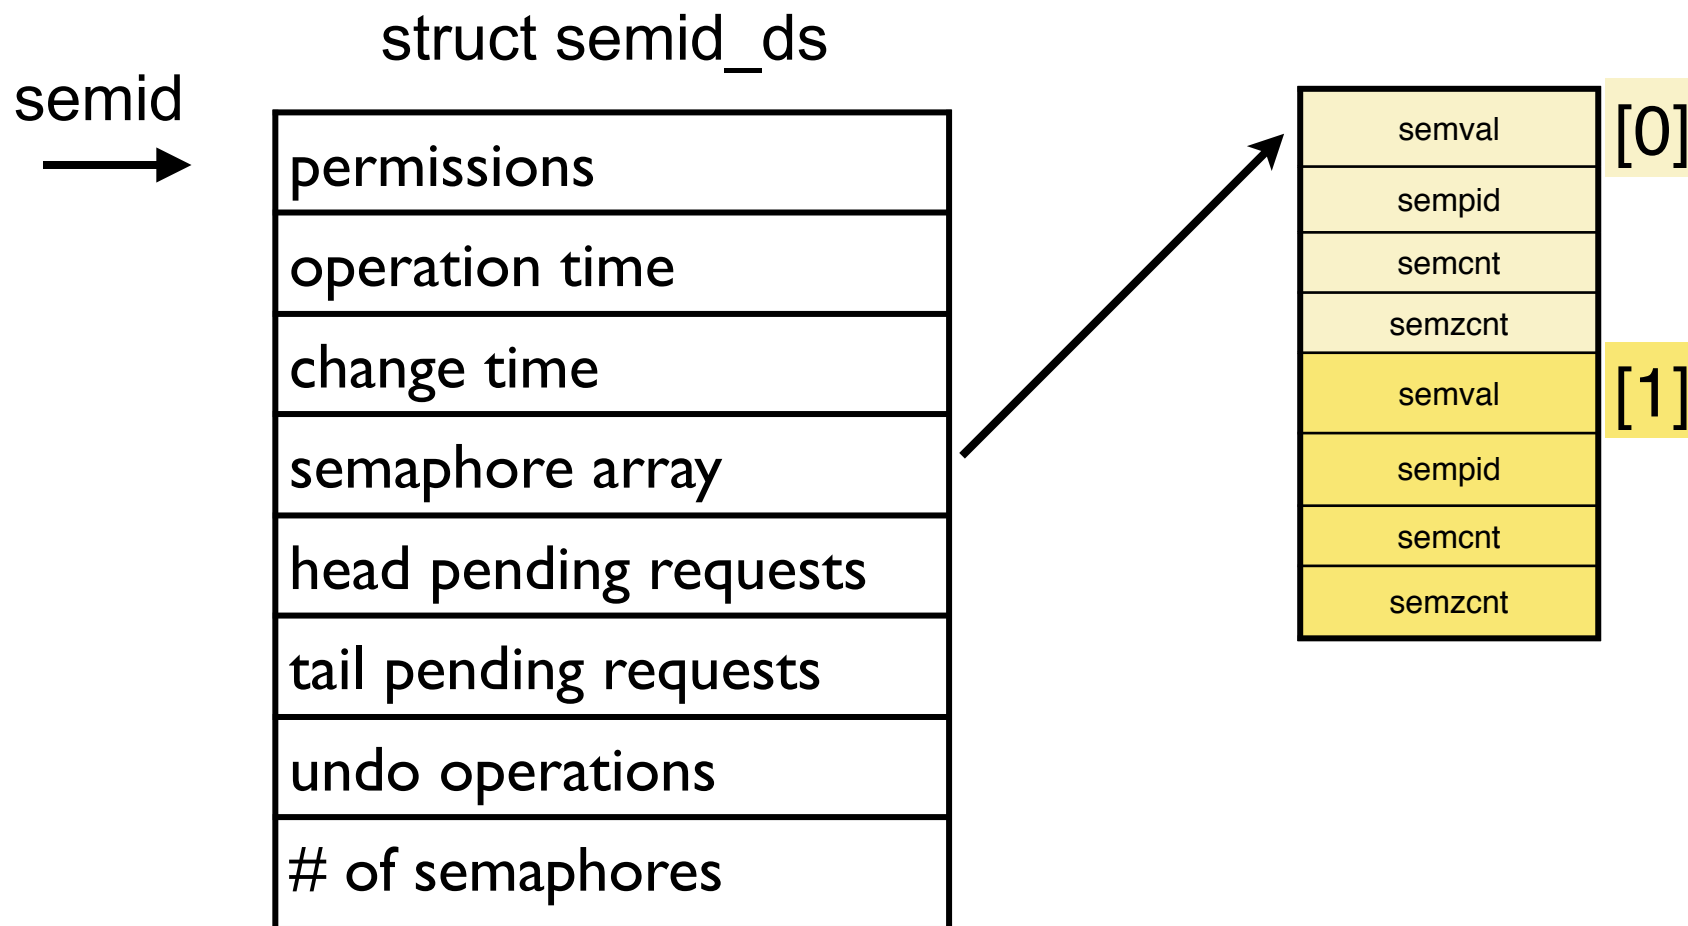
- For every set of semaphores kernel maintains following structure

```
#include <sys/types.h>
#include <sys/ipc.h>          /* defines ipc_perm structure*/
struct semid_ds {
    struct ipc_perm    sem_perm;    /* operations permission */
    __time_t           sem_otime;    /* time of last semop      */
    __time_t           sem_ctime;    /* time of last change     */
    struct sem         *sem_base;    /* ptr to 1st semaphore in set */
    struct sem_queue    *sem_pending; /* pending operations*/
    struct sem_queue    *sem_pending_last; /* last pending operation*/
    struct sem_undo     *undo;        /* undo requests on this set */
    unsigned long int    sem_nsems;   /* # of semaphores in set */
};
```

Each member is described by following structure

```
struct sem {
    ushort semval;    /* semaphore value, nonnegative */
    short  sempid;    /* pid of last operation         */
    ushort semncnt;   /* # awaiting semval > cval      */
    ushort semzcnt;   /* # awaiting semval = 0         */
};
```

- Kernel data structure for a semaphore set



```

struct    sem  {
    ushort    semval;    /* semaphore value, nonnegative */
    short     sempid;    /* pid of last operation        */
    ushort    semcnt;    /* # awaiting semval > cval     */
    ushort    semzcnt;    /* # awaiting semval = 0       */
};
  
```

- **Kernel data structure for a semaphore set**
 - `sem_pending`
 - Linked list of pending semaphore operations
 - `sem_pending_last`
 - end of list
 - `sem_undo`
 - undo requested semaphore operation
 - process exits without releasing semaphore operation
 - used with `SEM_UNDO` flag
 - helps reduce deadlock possibilities

● Semaphore creation/access

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semget (key_t key,  int nsems,  int semflag);
```

returns semid

nsems: number of semaphores (creation only)

semflag:

Numeric	Symbolic	Description
0400	SEM_R	Read by owner
0200	SEM_A	Alter by owner
0040	SEM_R >> 3	Read by group
0020	SEM_A >> 3	Alter by group
0004	SEM_R >> 6	Read by world
0002	SEM_A >> 6	Alter by world
	IPC_CREAT	
	IPC_EXCL	

- `semget`
 - `key`
 - to generate unique semaphore id
 - `IPC_PRIVATE`
 - `nsems`
 - number of semaphores
 - only used for creation
 - `semflag`
 - permission
 - `IPC_CREAT`
 - semaphore array is not initialized
 - `IPC_EXCL`


```
/* Creating sets of semaphores */
#include <iostream>
#include <stdio>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
using namespace std;
int
main( ){
    int      sem1, sem2, sem3;
    key_t    ipc_key;
    ipc_key = ftok(".", 'S');
    if ((sem1 = semget(ipc_key, 3, IPC_CREAT | 0660)) == -1) {
        perror("semget: IPC_CREAT | 0660");
    }
    printf("sem1 identifier %d\n", sem1);
    if ((sem2 = semget(ipc_key, 3, IPC_CREAT|IPC_EXCL|0600)) == -1) {
        perror("semget: IPC_CREAT | IPC_EXCL | 0660");
    }
    printf("sem2 identifier %d\n", sem12);
    if ((sem3 = semget(IPC_PRIVATE, 3, 0600)) == -1) {
        perror("semget: IPC_PRIVATE");
    }
    printf("sem3 identifier %d\n", sem3);
    return 0;
}
```

```
$/a.out
sem1 identifier 2818048
semget: IPC_CREAT | IPC_EXCL | 0600: File exists
sem2 identifier -1
sem3 identifier 65537
```

```
$ ./a.out
sem1 identifier 2818048
semget: IPC_CREAT | IPC_EXCL | 0600: File exists
sem2 identifier -1
sem3 identifier 65538
```

```
$ ipcs -s
```

T	ID	KEY	MODE	OWNER	GROUP
Semaphores:					
s	2818048	0x530253c3	--ra-ra----	dcanas	staff
s	65537	0x00000000	--ra-----	dcanas	staff
s	65538	0x00000000	--ra-----	dcanas	staff

● Semaphore control

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semctl (int semid, int semnum, int cmd, union semnum arg);
```

use cmd :

IPC_RMID to remove semaphore
IPC_STAT returns current value in arg
IPC_SET to set semaphore value - initialization
 set semaphore to arg.val

```
union semnum {
    int      val;                /* used for SETVAL only          */
    struct semid_ds *buf;        /* used for IPC_STAT and IPC_SET */
    ushort int *array            /* used for GETALL & SETALL      */
    struct seminfo *_buf;        /* buffer for IPC_INFO           */
} arg;
```

- **arg values**
 - an integer used with SETVAL to indicate a specific values for a particular semaphore
 - a reference to semid_ds structure for information returned by IPC_STAT or IPC_SET
 - a reference to an array of type unsigned integers used to either initialize the set or as a return location for GETALL
 - a reference to a seminfo structure when IPC_INFO is requested

- Semaphore Control (cmd values)
 - on specific set
 - IPC_STAT
 - values of sem_id structure returned in fourth argument
 - IPC_SET
 - permissions, uid, gid information
 - as specified in fourth argument
 - IPC_RMID
 - remove
 - on entire set
 - GETALL
 - returns values of the semaphore set
 - SETALL
 - sets values of the semaphore set

- Semaphore Control
 - on individual semaphores
 - GETVAL
 - returns value of individual semaphore referenced by semnun
 - SETVAL
 - sets value of individual semaphore referenced by semnun
 - GTPID
 - GETCNT
 - returns number of processes waiting on semaphore to increase value
 - GETZCNT
 - returns number of processes waiting on semaphore to become 0

● Semaphore Operations

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
```

```
int semop (int semid,  struct sembuf *sops,  unsigned int nsops);
```

pointer sops points to an array of the following structure:

```
struct sembuf {
    unsigned short  sem_num;      /* semaphore #          */
    short int       sem_op;       /* semaphore operation  */
    short int       sem_flag;     /* operations flag      */
};
```

nsops : number of elements in array of sembuf

Each element specifies an operation for one particular semaphore value specified by sem_num

- semaphore control
 - sem_flag
 - IPC_NOWAIT
 - return immediately if operation cannot be performed
 - no other semaphore sets are modified
 - SEM_UNDO
 - if IPC_NOWAIT not specified
 - undo operations if a blocked operation subsequently fails

● Semaphore Operations

1. If `sem_op` is positive, the value of `sem_op` is added to the semaphore's current value. This corresponds to the release of resources that a semaphore controls.
2. If `sem_op` is zero, the caller wants to wait until the semaphore's value becomes zero.
3. If `sem_op` is negative, the caller wants to wait until the semaphore's value becomes greater or equal to the absolute value of `sem_op`. Then the absolute value of `sem_op` is subtracted from the semaphore's current value. This corresponds to the allocation of resources

- actions when semop is positive

flag set	action taken by semop
	add sem_op to semval
SEM_UNDO	add sem_op to semval and update the undo counter for the semaphore

- actions when semop is zero

condition	flag set	action taken by semop
<code>semval == 0</code>		Return immediately
<code>semval != 0</code>	<code>IPC_NOWAIT</code>	Return -1 immediately and set <code>errno</code> to <code>EAGAIN</code>
<code>semval != 0</code>		Increment <code>semzcnt</code> for the semaphore and wait until: - <code>semval == 0</code> , then adjust <code>semzcnt</code> and return - <code>semid</code> is removed, then return -1 and set <code>errno</code> to <code>EIDRM</code> -a signal is caught, then adjust <code>semzcnt</code> and set <code>errno</code> to <code>EINTR</code>

- actions when semop is negative

condition	flag set	action taken by semop
<code>semval >= abs(semop)</code>		Subtract <code>abs(sem_op)</code> from <code>semval</code>
<code>semval >= abs(semop)</code>	<code>SEM_UNDO</code>	Subtract <code>abs(sem_op)</code> from <code>semval</code> and update the undo counter for the semaphore
<code>semval < abs(semop)</code>		increment <code>semcnt</code> for the semaphore and wait until - <code>semval >= abs(semop)</code> , then adjust <code>semcnt</code> and subtract as noted in the previous two rows of table - <code>semid</code> is removed, then return -1 and set <code>errno</code> to <code>EIDRM</code> -a signal is caught, then adjust <code>semcnt</code> and set <code>errno</code> to <code>EINTR</code>
<code>semval < abs(semop)</code>	<code>IPC_NOWAIT</code>	return -1 immediately and set <code>errno</code> to <code>EAGAIN</code>

- operations on a set of semaphores

