

# TCP/IP Networks

---

CSC 348-648



Spring 2013

## Network Review


---

- Physical layer
  - Concerns the physical transmission bits
  - Example digital encoding is +5 volts for a 1 and -5 volts for a 0
- Data link layer
  - Frames data and controls line access
  - Assumes all computers are *local* (**no** forwarding)
  - Example data link protocol is CSMA/CD
- Network layer
  - Allows datagrams (packets) to be forwarded to other networks
  - A router receives packets on one line and transmits on another
  - Example network layer is Internet Protocol (IP)

## IP Operation

Assume a user has data to send over the network

- IP receives data from upper layer
- Creates **packets** containing data
- Sends packets to *next-hop* (forward to next machine)
- Process repeats until the destination is reached



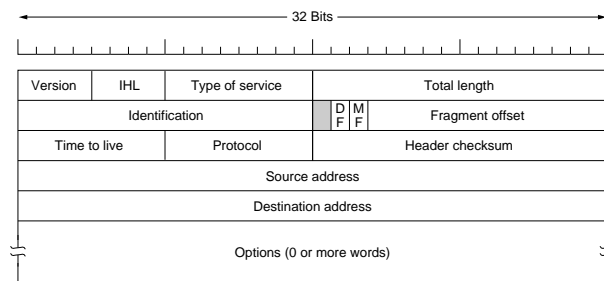
*How does the next-hop (router) know who the packet belongs to and who is the next hop?*

## IP Protocol Datagrams

- Datagram (packet) consists of a header and data



- Header consists of: 20 byte *fixed part* and an *optional part*



- Header contains different *control fields*
- We will focus on the address

## IPv4 Addresses


---

- Every host or router (actually interface) has a unique IP address
  - IP addresses are 32 bits long (IP version 4)
- *Dotted-decimal notation* is used to represent an address
  - Each byte is represented via a decimal number
  - $193.32.216.9 \Rightarrow [11000001\ 00100000\ 11011000\ 00001001]$
- Addresses are hierarchical and encode two numbers
  - **Network** - identifies the network
  - **Host** - identifies the machine in the network

## IP Network Example


---

Consider one router and seven hosts (*one address per interface*)



- Three hosts at bottom have similar addresses,  $223.1.1.x$ 
  - The leftmost 24 bits they share is the **network** portion
  - Remaining 8 bits is the **host** portion
- Hosts of  $223.1.1.x$  form a network, interconnected via a LAN
  - The network address is  $223.1.1.0/24$

- The /24 is also called the **network mask** or **network prefix** (indicates the leftmost 24 bits are the network address)
- Any additional host attached to this network must have a unique address of the form 223.1.1.*x*
- Other networks have a similar structure



- N.B. the router interconnects different networks

## IPv4 Address Classes

The original Internet architecture defined 5 different IP address classes

- This is also known as **classful addressing**
- Given an IP address you can determine the network mask

<div style="text-align: center;"> <span style="font-size: 1.2em;">←</span> <span style="font-size: 1.2em;">→</span> </div> <div style="text-align: center; margin-top: 5px;">32 Bits</div>				
Class				Range of host addresses
A	0	Network	Host	1.0.0.0 to 127.255.255.255
B	10	Network	Host	128.0.0.0 to 191.255.255.255
C	110	Network	Host	192.0.0.0 to 223.255.255.255
D	1110	Multicast address		224.0.0.0 to 239.255.255.255
E	11110	Reserved for future use		240.0.0.0 to 247.255.255.255

## IP Addresses with Special Meanings

---

0 0																																This host
0 0				...				0 0				Host																A host on this network				
1 1																																Broadcast on the local network
Network								1 1 1 1				...				1 1 1 1																Broadcast on a distant network
127				(Anything)																												Loopback

- 0.0.0.0 only used by a host when booting
- All zeroes for the network number, refers to the local network
  - If 223.1.1.0/24 is the network and I am 223.1.1.52, locally I can be reached using 0.0.0.52

- Address of all ones is the broadcast address for the local network

*What is the dotted-decimal address?*

- Address with the proper network number, and all ones for the host number allows host to broadcast to a different network
  - If 223.1.1.0/24 is a distant network, then 223.1.1.255 broadcasts to all hosts at the network
  - This will be used to create a simple Denial of Service (DoS) attack
- 127.x.y.z is reserved for loop-back testing
  - Packet is never placed on the network, processed locally

## Routing Tables

---

How does a source host send a datagram to a destination host?


- The IP layer maintains a **routing table** in memory
  - Remember, routing tables are *next hop* oriented
  - Multiple hop paths are not recorded
- Each entry in the routing table has the following information<sup>a</sup>
  1. Destination address, either *host* or *network* address
  2. IP address of the *next-hop router*
  3. Flags specifying if next hop is host or network
  4. Identification of the interface the datagram should be passed to (e.g. multiple Ethernet cards attached)

---

<sup>a</sup>Abbreviated list of items, more later.

## Example Routing Tables

---



- In the diagram, each interface (Ethernet card) is labeled (in red)
- For example, the router has 3 interfaces (eth0, eth1, and eth2)
  - Each interface must be uniquely identified, since it attaches a unique network

- An abbreviated routing table for host A would be

Routing Table for A		
Destination	Next Hop	Interface
223.1.1.0/24		eth0
223.1.2.0/24	223.1.1.4	eth0
223.1.3.0/24	223.1.1.4	eth0

- First entry indicates 223.1.1.0/24 is the local network
- The second and third entries indicate datagrams for destinations on network 223.1.2.0/24 or 223.1.3.0/24 must be sent to 223.1.1.4
- eth0 is the Ethernet interface (only one card on A)

*Each network is represented with one entry, how many would be required if each host had a separate entry?*

- An abbreviated routing table for the router would be

Routing Table for Router		
Destination	Next Hop	Interface
223.1.1.0/24		eth0
223.1.2.0/24		eth1
223.1.3.0/24		eth2

- First entry indicates 223.1.1.0/24 is local on eth0
- Second entry indicates 223.1.2.0/24 is local on eth1
- Third entry indicates 223.1.3.0/24 is local on eth2

## IP Routing Steps


---

- IP routing performs the following actions
  1. Search routing table for complete destination address, if found send packet to the next-hop entry
  2. Search routing table for an entry that matches the destination network number, if found send packet to the next-hop entry
    - Must take into account possible subnet mask
  3. Search for *default* entry, if found send to next-hop router
- IP search order is, host address → host network → default
- If all the steps fail, then the datagram is not deliverable

## Routing Example: A → B

---

Assume A (223.1.1.1) sends datagram to B (223.1.1.3)



- There is no host entry for 223.1.1.3
- There is a network entry for 223.1.1.0/24
- A link layer frame (containing the datagram) is created and addressed to the link layer address of 223.1.1.3

*We are at layer 3, how do we get a layer 2 address?*

- Ethernet frame is sent and received by host B



## Routing Example: A → E

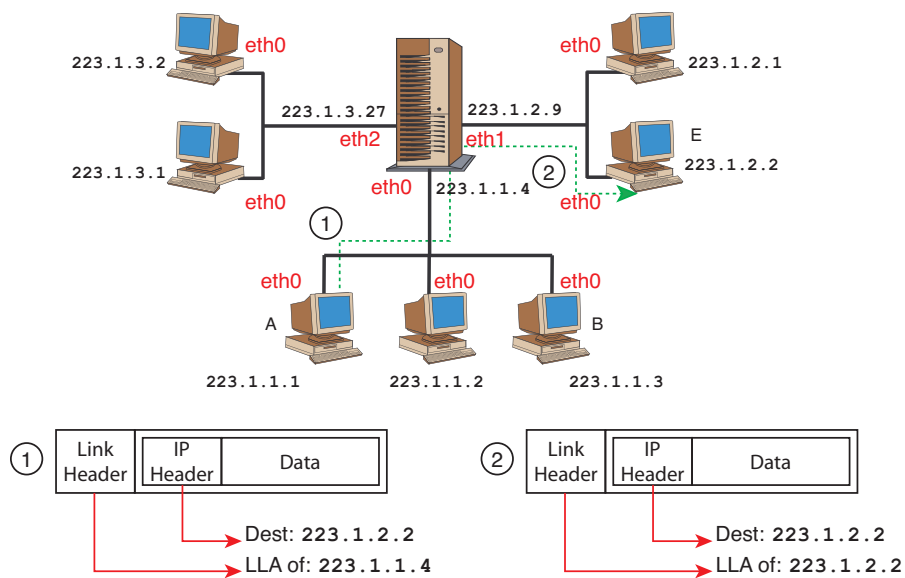
---

Assume A (223.1.1.1) sends datagram to E (223.1.2.2)

Routing Table for A			Routing Table for Router		
Destination	Next Hop	Interface	Destination	Next Hop	Interface
223.1.1.0/24		eth0	223.1.1.0/24		eth0
223.1.2.0/24	223.1.1.4	eth0	223.1.2.0/24		eth1
223.1.3.0/24	223.1.1.4	eth0	223.1.3.0/24		eth2

- Host A finds entry for 223.1.2.0/24 network
  - Requires sending packet to 223.1.1.4
- Host A creates and sends link-layer frame (containing datagram) addressed to the link-layer address of 223.1.1.4
  - Therefore, the next-hop entry is used for the link-layer address
  - IP destination address remains unchanged

- Router 223.1.1.4 receives frame and removes datagram
  - Destination address is 223.1.2.2
  - Router is allowed to forward datagrams
- Router finds entry for 223.1.2.0/24 network
  - This is directly connected via eth1
  - Datagram will be forwarded
- Router creates and sends link-layer frame (containing datagram) addressed to the link-layer address of 223.1.2.2 on eth1
- Frame received by host E, datagram removed and processed
- N.B. operation of host and router are equivalent, except routers are allowed to forward datagrams




## Static versus Dynamic IP Routing

- Already know how IP packets are routed using routing tables
  - *How were the entries generated... statically or dynamically?*
- In static routing entries are manually adjusted
  - Acceptable for small networks
- For larger networks, dynamically change table entries
  - Routes should change based on network conditions
  - Allow routers to pass route information to one another
  - Use variations of Bellman-Ford and Dijkstra's
- N.B. This will **not** change the way IP datagrams are routed, just how/when the routing table contents change

## Layer 3 Addresses versus Layer 2 Addresses

---

- Sending IP datagrams requires knowledge of the link-layer address
  - The IP datagram is placed inside a link-layer frame then sent




- Hosts are attached to the network via an interface card
  - The interface card is a layer 1 and 2 device
  - Can only understand **LAN addresses** (also called **hardware address** or **MAC address**)

- For example, every Ethernet board has a 48-bit Ethernet address
  - In DOS run `winipcfg` to determine the Ethernet address
  - Card manufacturers given a *block* of Ethernet addresses
  - Assignments given in RFC 1700 or a more up-to-date at <http://standards.ieee.org/develop/regauth/oui/public.html>
- Similar to IP, hardware address of all 1's is for broadcast

*How do IP addresses get mapped to data-link addresses?*


## Address Resolution

- Address resolution - Mapping between two different address forms
  - IP addresses and data-link addresses
- Two different protocols in IP
  - Address Resolution Protocol (ARP)
  - Reverse Address Resolution Protocol (RARP)



## ARP

- Dynamic mapping of IP address → hardware address [RFC 826]
- Assume A needs the hardware address of host B



- Host A broadcasts an ARP request
  - Ethernet address of all 1's is the broadcast address
  - Host asks *"Who owns IP address 160.1.1.2?"*
- Host B would reply with its hardware address

- ARP is intended for broadcast (shared medium) networks
  - Not needed for point-to-point links
- Broadcast at boot-time
  - When a machine connects, broadcast its mapping (ARP reply)
  - All other hosts on network (or subnet) will know the IP and hardware address of the new machine

## ARP Optimization

---

There are various ways to improve the efficiency of ARP

- ARP Tables
  - Keep an ARP table for most recently used IP addresses
  - Entries can time-out
- Hey, this is lots of fun...
  - When your laptop is first connected to the network, enter the DOS command, `arp -a`
  - Ping a machine on our network, `ping www.wfu.edu`
  - Issue the `arp` command again, should see hardware (physical) address for the machine
    - This may be a proxy ARP, how can you tell?*
  - Wait a few minutes and issue the `arp` command again

## ARP Security

---

- ARP introduces a security problem

*How could this protocol be used to obtain important information from users?*

- A common way around this security issue is to use **static ARP**
  - Permanent ARP table created
  - *So what is ARPish about this solution...*

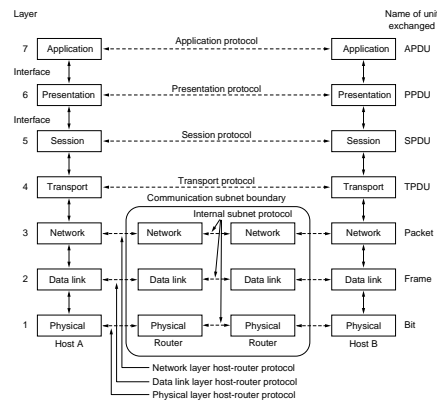
## Transport Layer

---

- Reliable, cost-effective data transport from source to destination
  - End-to-end protocol, only implemented at hosts
- *Transport entity* provides services to upper layers
  - Type of service, QoS, data transfer, connection management, and flow control
- Two types of service available
  - Connection-oriented - Establishment, maintenance, and termination of connection
  - Connectionless - Unreliable service (delivery not guaranteed), reduces the overhead associated with transport layer

*Network layer also has connection-oriented and connectionless, why do we specify it here?*

# Network and Transport Layers




- The network layer is part of the communication subnet
  - What if the network-layer is connection-oriented but unreliable?
  - What happens if a router crashes?
  - *Provide a reliable service over an unreliable network*

## Transport Addressing

- When an application process wishes to set-up a connection to a remote application process, it must specify which one
    - 152.17.140.92 provides `http` and `ssh` services
    - Therefore network layer address is **not** sufficient
  - In the Internet, transport services identified using **port numbers**
    - Some port numbers are *well-known*, 80 is for `http`
    - DOS command `netstat -a` shows ports in use
- So what is a port scan?*

## Internet Transport Layer

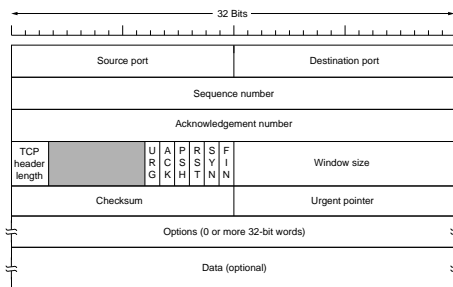
- There are two distinct Internet transport-layer protocols
  - User Datagram Protocol (UDP)
  - Transmission Control Protocol (TCP)




- UDP provides unreliable connectionless service
- TCP provides reliable connection-oriented service

## Transport Layer Packets

- Packets have a header and data
  - Headers for TCP and UDP are different

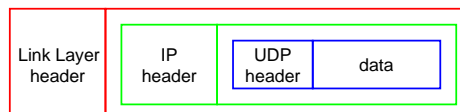


TCP header



UDP header

- At the physical layer the data looks like





## When is UDP Preferred

---

UDP has the following benefits, making it better for some applications

- No connection establishment
  - TCP requires a three-way handshake before sending data
  - UDP sends data immediately, no initial connection delay
- No connection state
  - TCP requires state information about send/receive buffers, congestion-control, sequence numbers... *per connection*
  - UDP does not require any of this state information
- Small packet header overhead
  - UDP only adds 8 bytes of header information
- Unregulated send rate
  - Send as quickly as desired
  - Has introduced the idea of *TCP friendly* applications

## When is TCP Preferred

---

- Transport Control Protocol (TCP) provides connection-oriented, reliable service [RFC 793, 1122, 1323, 2018, 2581]
- All connections TCP connections are full-duplex, point-to-point
  - Requires three-way handshake to establish connection
- Reliable transport service
  - Deliver all data without error and proper order
- Congestion control mechanism
  - Attempts to limit each connection to *fair share* of bandwidth (*Crouse says “très bon, l'équité est tout”*)
  - **No** minimum bandwidth guaranteed
- TCP connection is a byte stream, not a message stream
  - Message boundaries are not preserved
  - Data is buffered before sent (additional delay)

## Port Numbers

---


- Transport addresses are used to identify processes
  - For Internet transport layers, this called a **port number**
  - **Both** source and destination processes have a port number
- Source and destination ports *together* uniquely identify a process

*Do we really need source and destination port numbers to identify a process (single process at one end)?*
- Port numbers are 16 bits, ranging from 0 to 65535
- Numbers ranging from 0 to 1023 are **well-known port numbers**
  - Reserved for use by well-known protocols
  - http is 80, ftp is 21, complete list at RFC 1700

## Web Server Example

---

- Consider a web server and two stations connected via the Internet



- Web server runs http over port 80
  - Any station wishing to connect will use the IP address 152.140.12.25 and port 80
- Let station A connect to the web server
  - A free port is used for source ( $> 1023$ ), for example 1123
  - Destination port is 80

- Web server creates a new process for each request (Unix fork)
  - Allows server to connect to multiple users simultaneously

*What happens if station B connects to the web server? Station B will use the same destination port 80, how does the web server differentiate between station A and station B requests?*

*What happens if station B connects to the web server and happens to select the same source port number as A? How does the web server differentiate between station A and station B requests?*

## Domain Name System

---

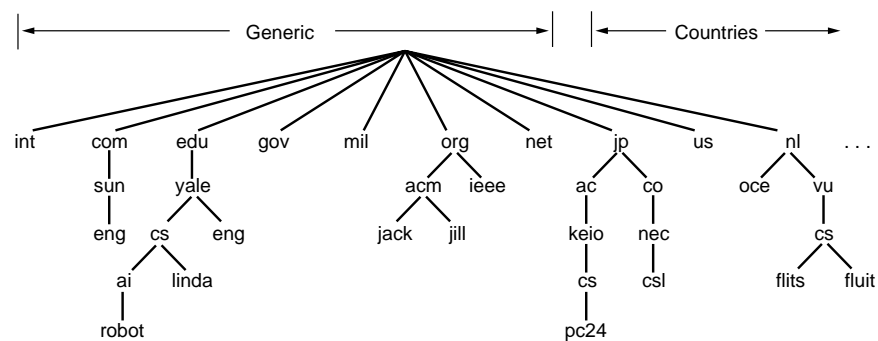
- We can always use IP (network) addresses to refer to hosts
  - IP addresses are difficult to remember, pluf@152.17.48.111

*What happens if the machine address changes?*

- However, the network only understands IP addresses
  - Need a method to map *names* to IP addresses
- Domain Name System (DNS) provides this mapping service
  - Hierarchical, domain naming scheme
  - Given a text (ASCII) name provides the correct IP address
  - DNS protocol uses UDP for small queries and TCP for queries where the answer is over 512 bytes (*any security issues?*)

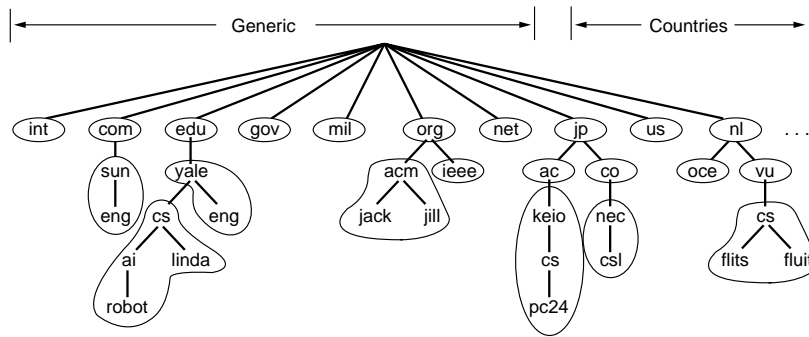
IP Address	Machine Name
152.17.48.111	mail.wfu.edu
152.17.48.77	www.wfu.edu

## DNS Name Space



- Hierarchical scheme of names, can be depicted as a tree
  - 200 top-level domains, each domain covers many hosts
  - Two categories, generic and countries
- The domain manager registers new names, for example .wfu.edu


## Name Servers



- DNS *name server* provides the name and IP mapping
  - A server could store the entire database, but it is not scalable
  - Will use the name hierarchy to provide scalability
- DNS is divided into non-overlapping zones
  - Zone is part of the overall tree and contains name servers
  - Name servers know about names in their zones

## Name Resolution

- Assume flits.cs.vu.nl needs to resolve linda.cs.yale.edu
  1. flits.cs.vu.nl first searches local file/cache first
  2. If not found, asks its local name server cs.vu.nl
  3. If cs.vu.nl does not know, it asks edu-server.net
  4. If edu-server.net does not know, it asks cs.yale.edu
  5. edu-server.net should know and sends the IP



- Once flits.cs.vu.nl gets the IP it stores it in cache
  - This is an example of a **recursive query**, where searches are performed on behalf of the original requester
  - Alternatively, IP of the next server to ask could have been sent

## DNS Caching

---

- DNS requests are cached
  - Quick response for repeated translations
  - Other queries may reuse parts of lookup records for domains
- DNS negative queries are cached
  - Do not have to repeat past mistakes (misspellings)
- Cached data periodically times out
  - Lifetime (TTL) of data controlled by owner of data
  - TTL passed with every record