

# Approximate Inference in Bayesian Networks

Chapter 14 – Part II

# Inference by Sampling

- Key Idea
  - Draw  $N$  samples from a sampling distribution  $S$
  - Compute an approximate posterior  $P'$
  - Show that this converges to the true probability  $P$
- If we could sample from a variable's (posterior probability), we could estimate its (posterior) probability

$X$	<i>count</i>
$x_1$	$n_1$
$\vdots$	$\vdots$
$x_k$	$n_k$
<i>total</i>	$m$

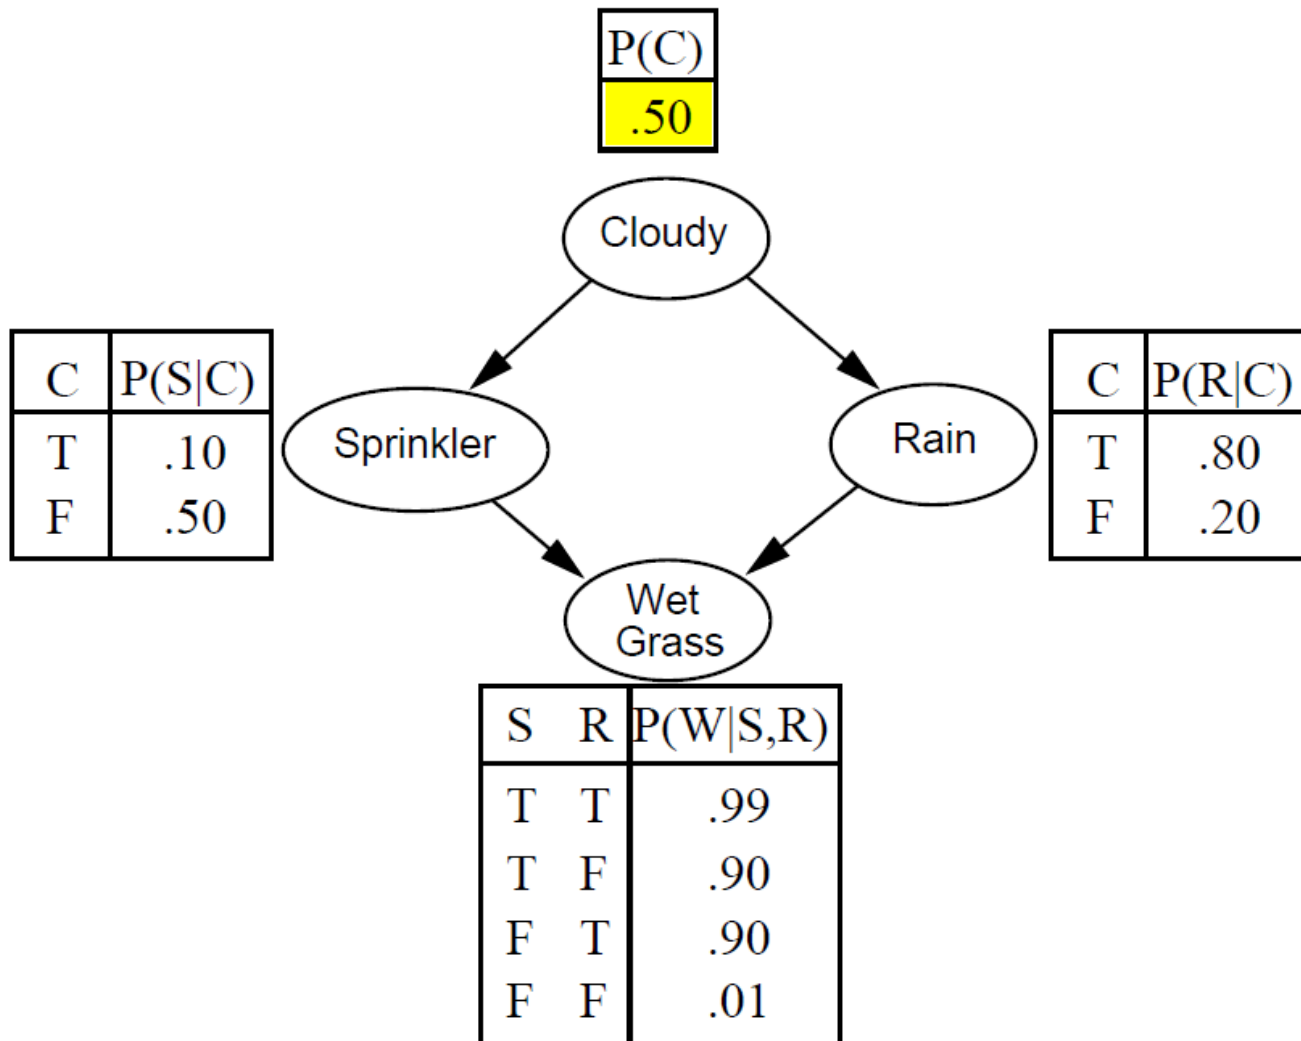
 $\Leftrightarrow$ 

$X$	<i>probability</i>
$x_1$	$n_1/m$
$\vdots$	$\vdots$
$x_k$	$n_k/m$

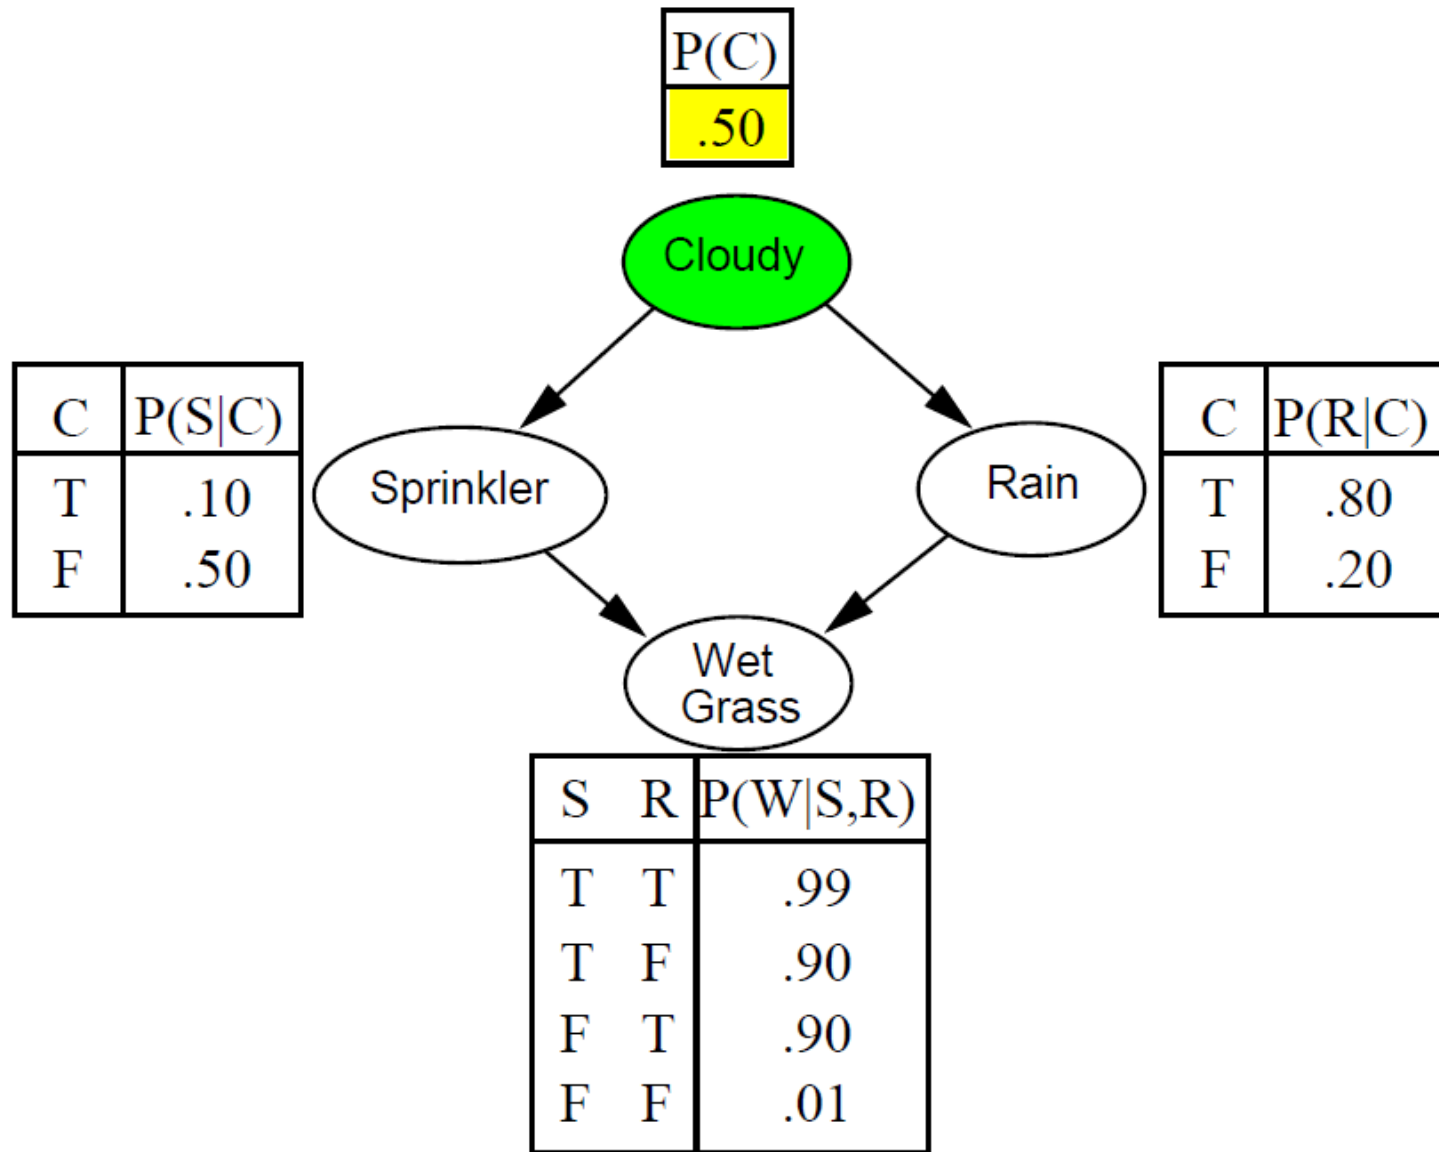
# Sampling from an empty network

```
function PRIOR-SAMPLE( $bn$ ) returns an event sampled from  $bn$   
  inputs:  $bn$ , a belief network specifying joint distribution  $\mathbf{P}(X_1, \dots, X_n)$   
   $\mathbf{x} \leftarrow$  an event with  $n$  elements  
  for  $i = 1$  to  $n$  do  
     $x_i \leftarrow$  a random sample from  $\mathbf{P}(X_i \mid \text{parents}(X_i))$   
    given the values of  $\text{Parents}(X_i)$  in  $\mathbf{x}$   
  return  $\mathbf{x}$ 
```

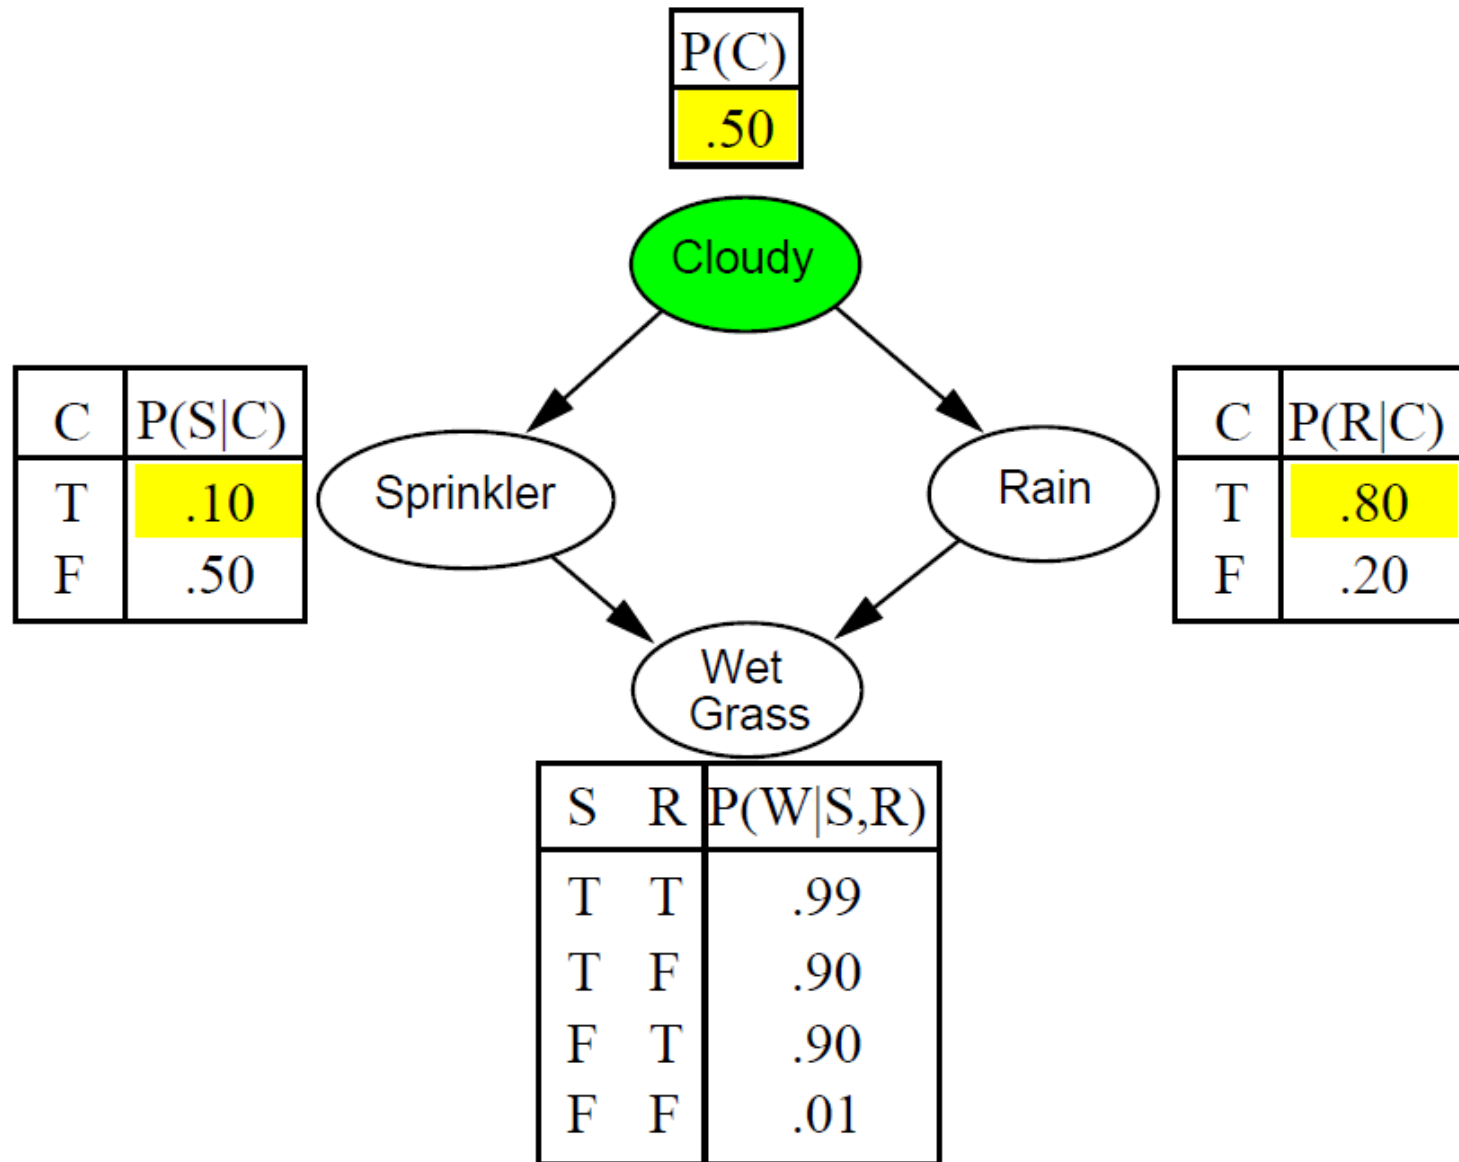
# Example



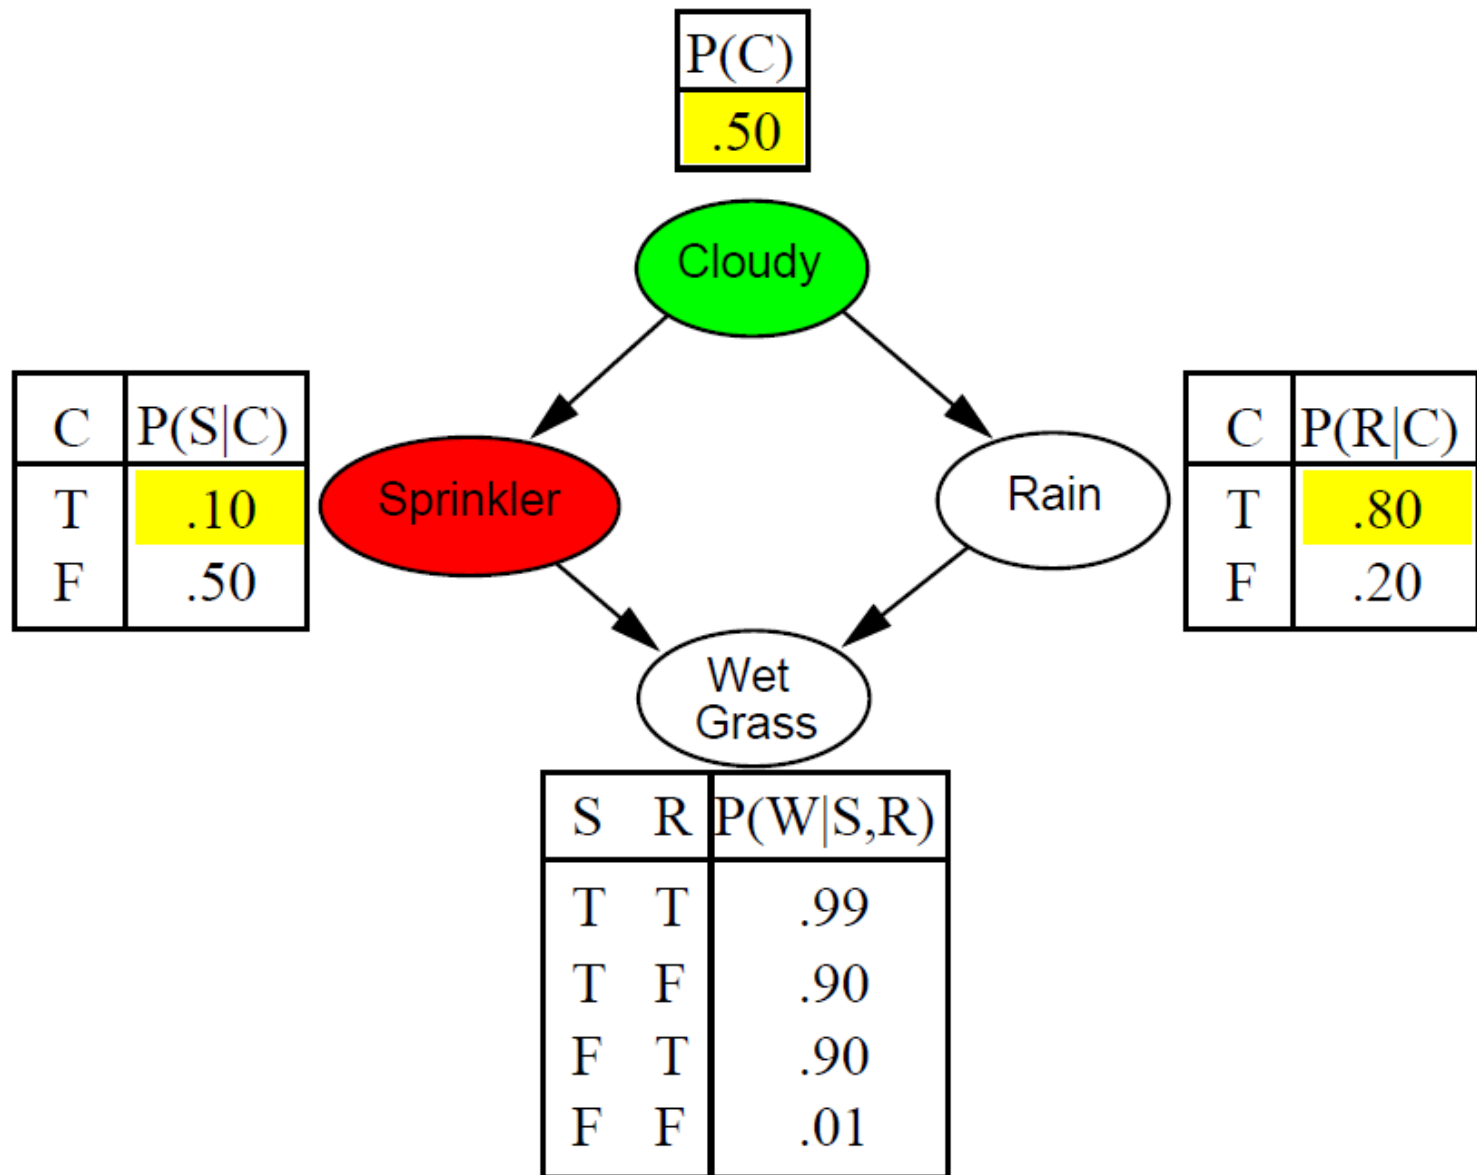
# Example



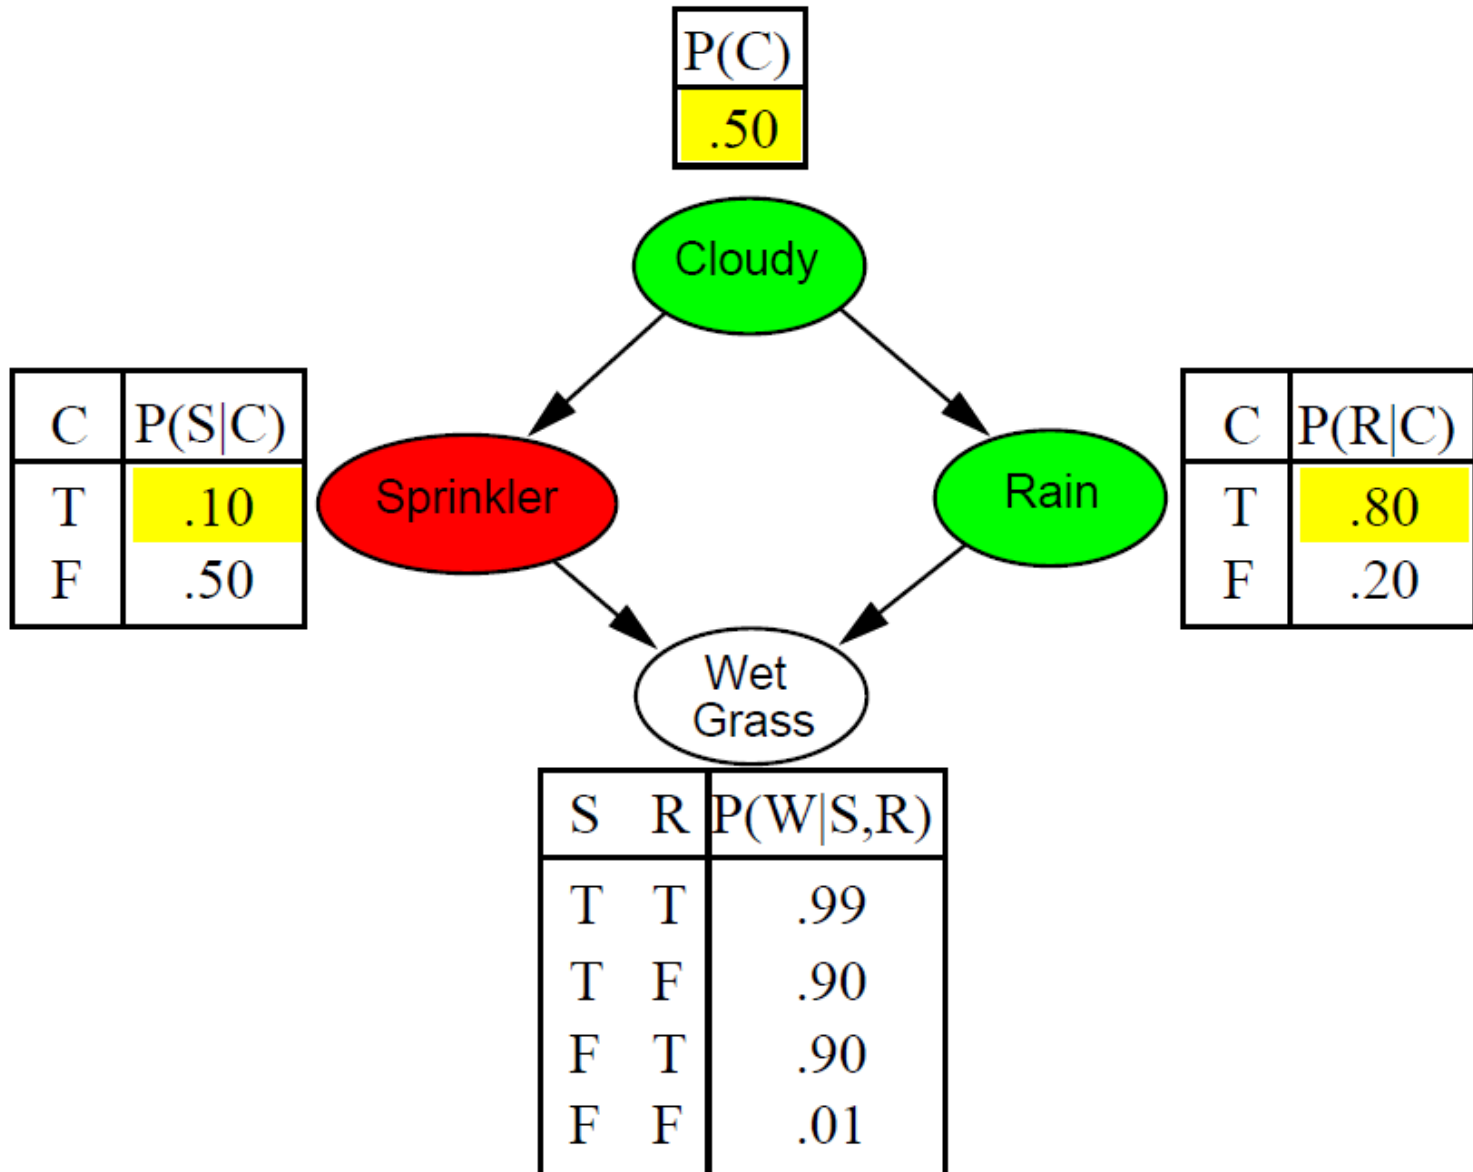
# Example



# Example

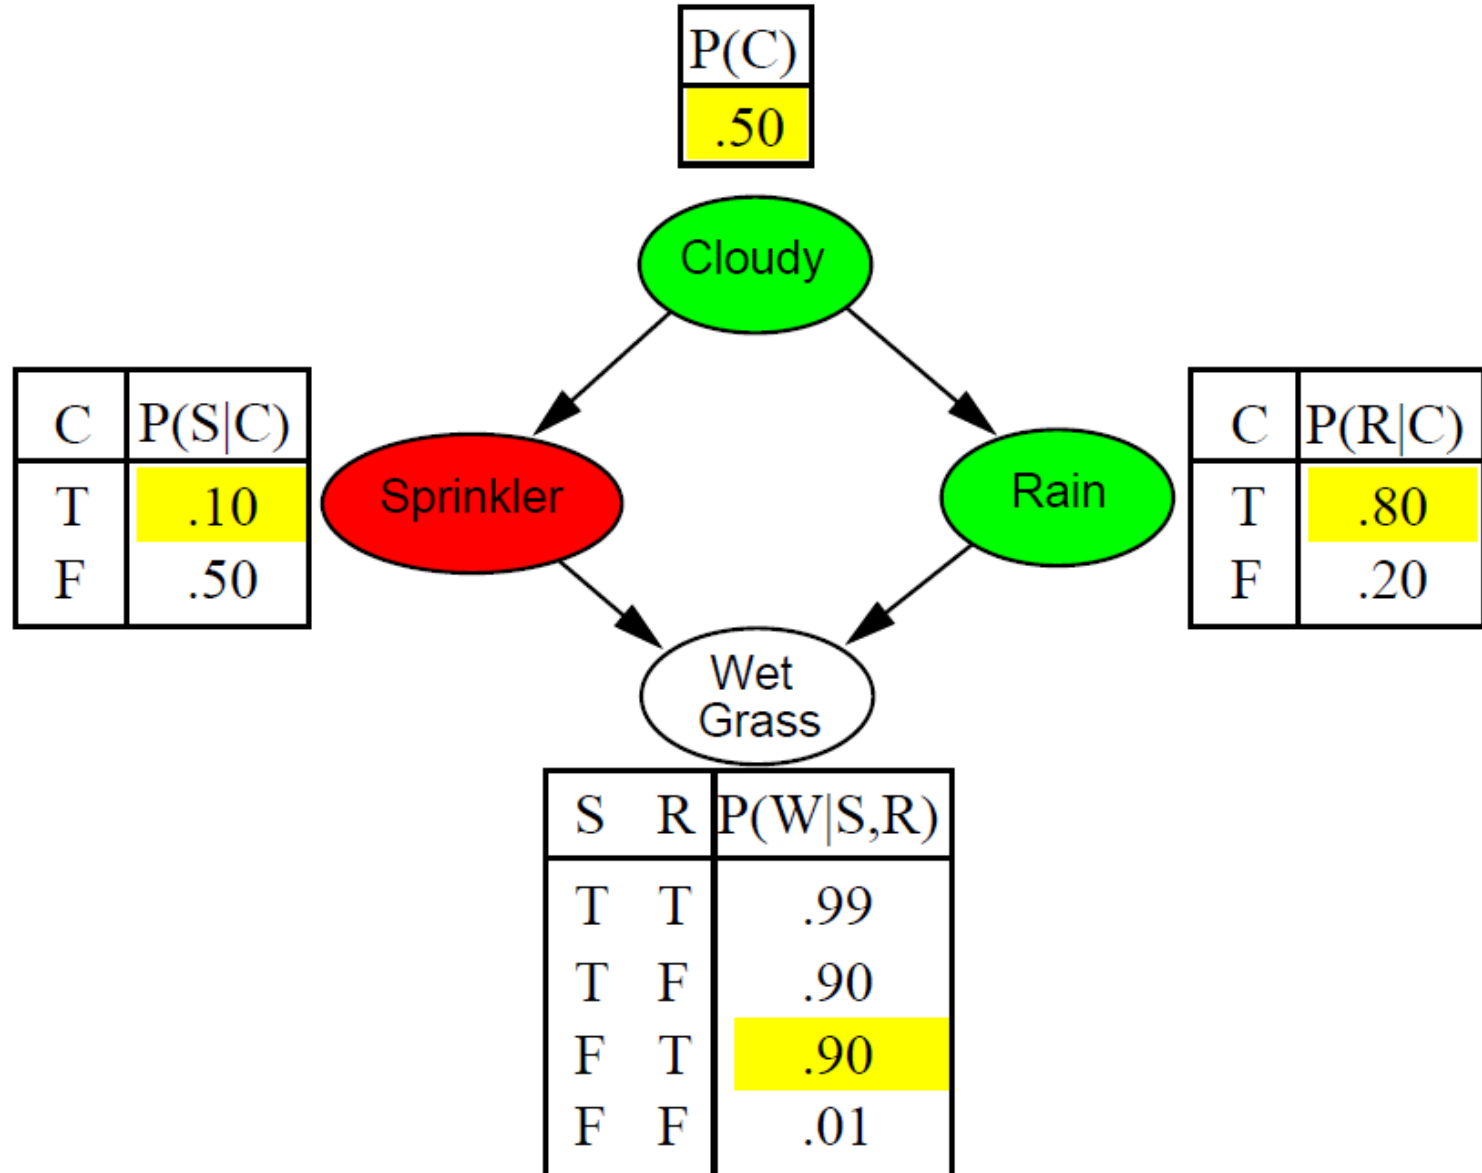


# Example

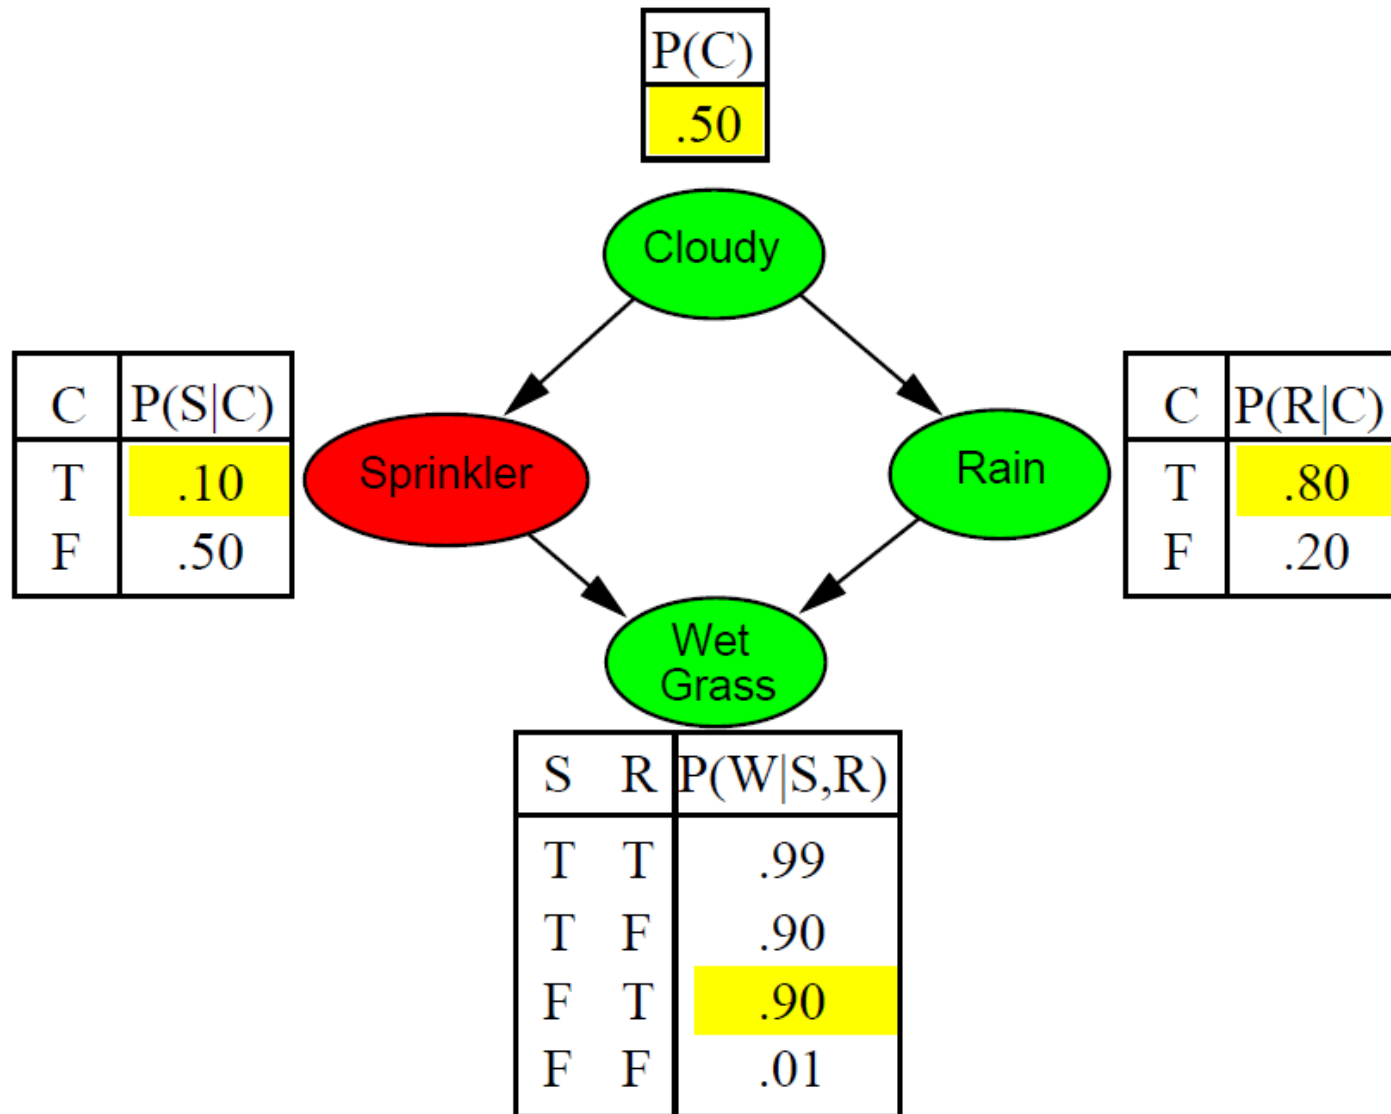




# Example



# Example



# Sampling from an empty network

Probability that PRIORSAMPLE generates a particular event

$$S_{PS}(x_1 \dots x_n) = \prod_{i=1}^n P(x_i | \text{parents}(X_i)) = P(x_1 \dots x_n)$$

i.e., the true prior probability

$$\text{E.g., } S_{PS}(t, f, t, t) = 0.5 \times 0.9 \times 0.8 \times 0.9 = 0.324 = P(t, f, t, t)$$

Let  $N_{PS}(x_1 \dots x_n)$  be the number of samples generated for event  $x_1, \dots, x_n$

Then we have

$$\begin{aligned} \lim_{N \rightarrow \infty} \hat{P}(x_1, \dots, x_n) &= \lim_{N \rightarrow \infty} N_{PS}(x_1, \dots, x_n) / N \\ &= S_{PS}(x_1, \dots, x_n) \\ &= P(x_1 \dots x_n) \end{aligned}$$

That is, estimates derived from PRIORSAMPLE are consistent

Shorthand:  $\hat{P}(x_1, \dots, x_n) \approx P(x_1 \dots x_n)$

# Rejection Sampling

$\hat{P}(X|e)$  estimated from samples agreeing with  $e$

```
function REJECTION-SAMPLING( $X, e, bn, N$ ) returns an estimate of  $P(X|e)$ 
  local variables:  $N$ , a vector of counts over  $X$ , initially zero
  for  $j = 1$  to  $N$  do
     $x \leftarrow$  PRIOR-SAMPLE( $bn$ )
    if  $x$  is consistent with  $e$  then
       $N[x] \leftarrow N[x] + 1$  where  $x$  is the value of  $X$  in  $x$ 
  return NORMALIZE( $N[X]$ )
```

E.g., estimate  $P(Rain|Sprinkler = true)$  using 100 samples

27 samples have  $Sprinkler = true$


Of these, 8 have  $Rain = true$  and 19 have  $Rain = false$ .

$\hat{P}(Rain|Sprinkler = true) = \text{NORMALIZE}(\langle 8, 19 \rangle) = \langle 0.296, 0.704 \rangle$

Similar to a basic real-world empirical estimation procedure

# Rejection Sampling - Example

Observe  $Sm = true, Re = true$



	Ta	Fi	Al	Sm	Le	Re	
$s_1$	false	true	false	true	false	false	<b>x</b>
$s_2$	false	true	true	true	true	true	<b>✓</b>
$s_3$	true	false	true	false	—	—	<b>x</b>
$s_4$	true	true	true	true	true	true	<b>✓</b>
...							
$s_{1000}$	false	false	false	false	—	—	<b>x</b>

$$P(sm) = 0.02$$

$$P(re|sm) = 0.32$$

How many samples are rejected?

How many samples are used?

# Analysis of Rejection Sampling

$$\begin{aligned}\hat{\mathbf{P}}(X|\mathbf{e}) &= \alpha \mathbf{N}_{PS}(X, \mathbf{e}) && \text{(algorithm defn.)} \\ &= \mathbf{N}_{PS}(X, \mathbf{e}) / N_{PS}(\mathbf{e}) && \text{(normalized by } N_{PS}(\mathbf{e}) \text{)} \\ &\approx \mathbf{P}(X, \mathbf{e}) / P(\mathbf{e}) && \text{(property of PRIORSAMPLE)} \\ &= \mathbf{P}(X|\mathbf{e}) && \text{(defn. of conditional probability)}\end{aligned}$$

Hence rejection sampling returns consistent posterior estimates

Problem: hopelessly expensive if  $P(\mathbf{e})$  is small

$P(\mathbf{e})$  drops off exponentially with number of evidence variables!

# Likelihood weighting aka Importance Sampling

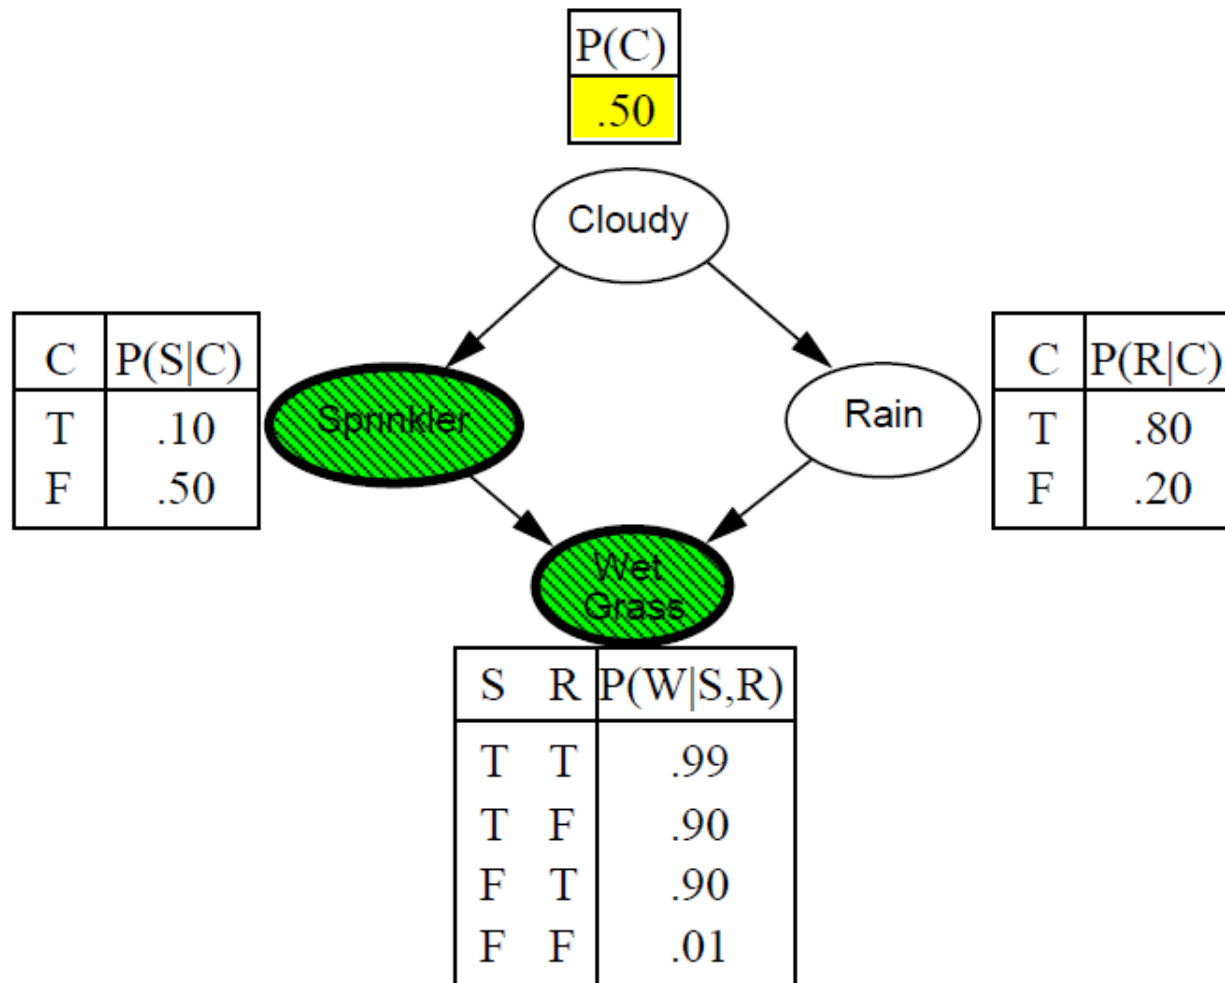
Idea: fix evidence variables, sample only nonevidence variables,  
and weight each sample by the likelihood it accords the evidence

```
function LIKELIHOOD-WEIGHTING( $X, e, bn, N$ ) returns an estimate of  $P(X|e)$   
  local variables:  $W$ , a vector of weighted counts over  $X$ , initially zero  
  for  $j = 1$  to  $N$  do  
     $x, w \leftarrow$  WEIGHTED-SAMPLE( $bn$ )  
     $W[x] \leftarrow W[x] + w$  where  $x$  is the value of  $X$  in  $x$   
  return NORMALIZE( $W[X]$ )
```

---

```
function WEIGHTED-SAMPLE( $bn, e$ ) returns an event and a weight  
   $x \leftarrow$  an event with  $n$  elements;  $w \leftarrow 1$   
  for  $i = 1$  to  $n$  do  
    if  $X_i$  has a value  $x_i$  in  $e$   
      then  $w \leftarrow w \times P(X_i = x_i \mid \text{parents}(X_i))$   
      else  $x_i \leftarrow$  a random sample from  $P(X_i \mid \text{parents}(X_i))$   
  return  $x, w$ 
```

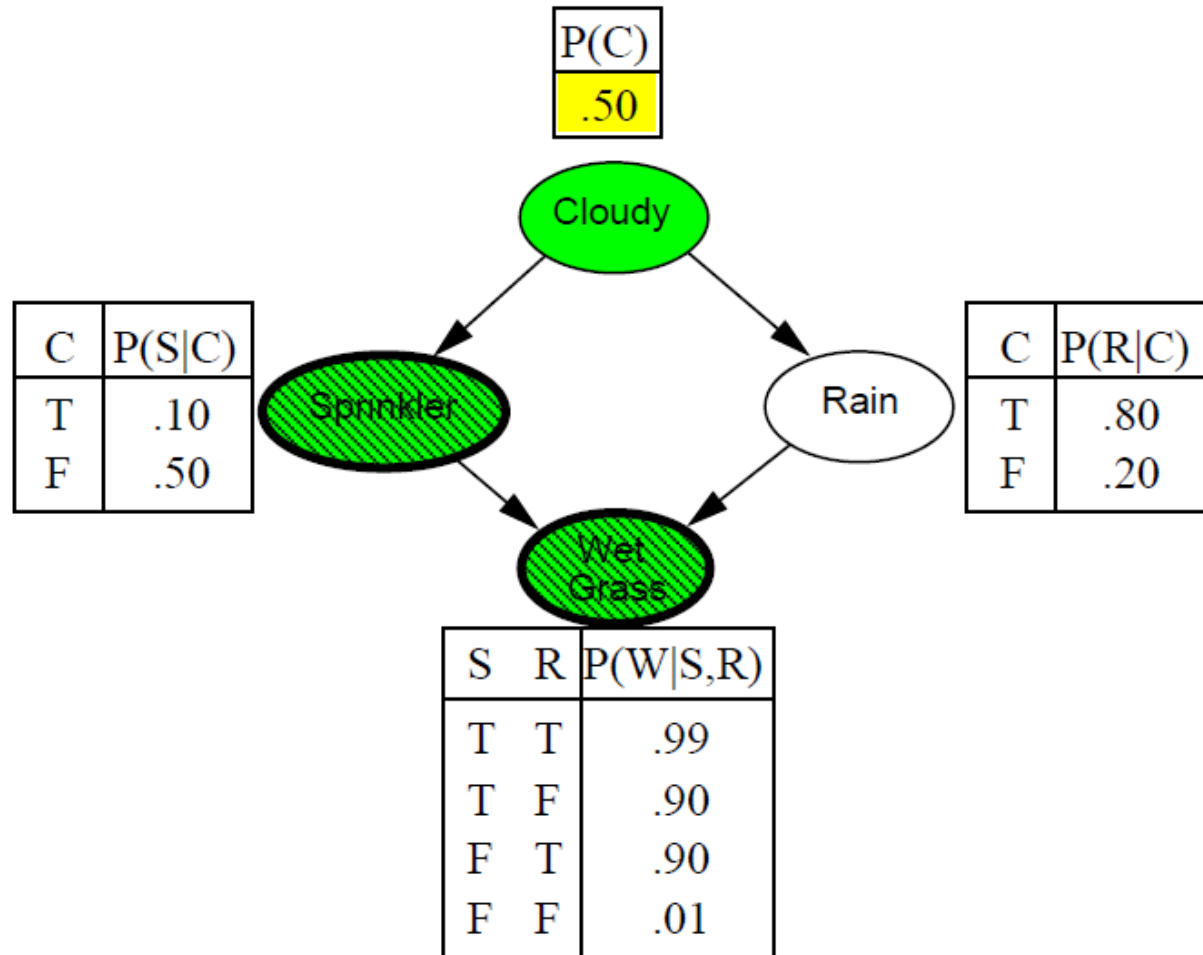
# Likelihood Weighting Example



$$w = 1.0$$

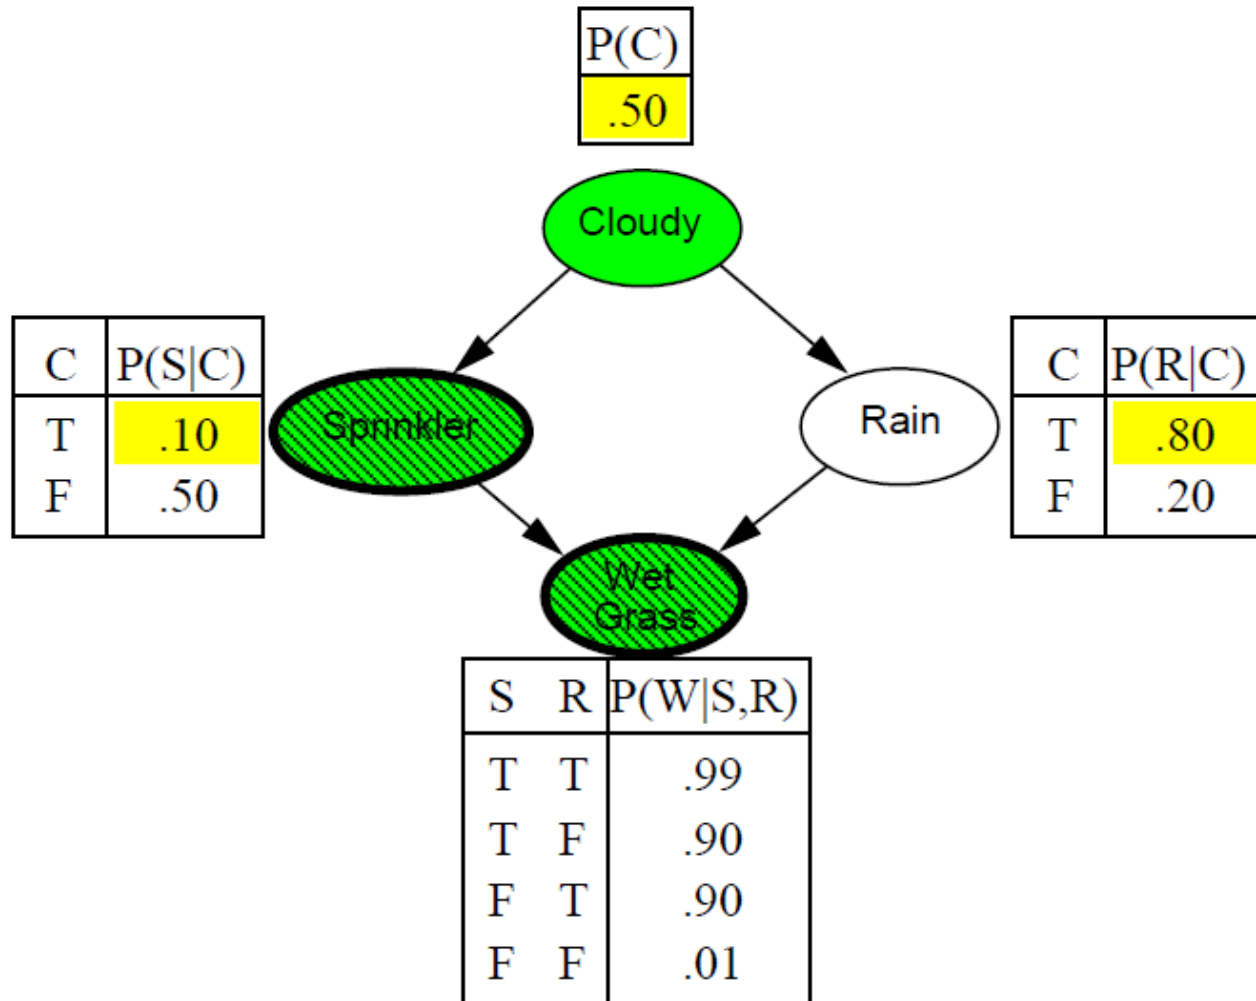


# Likelihood Weighting Example



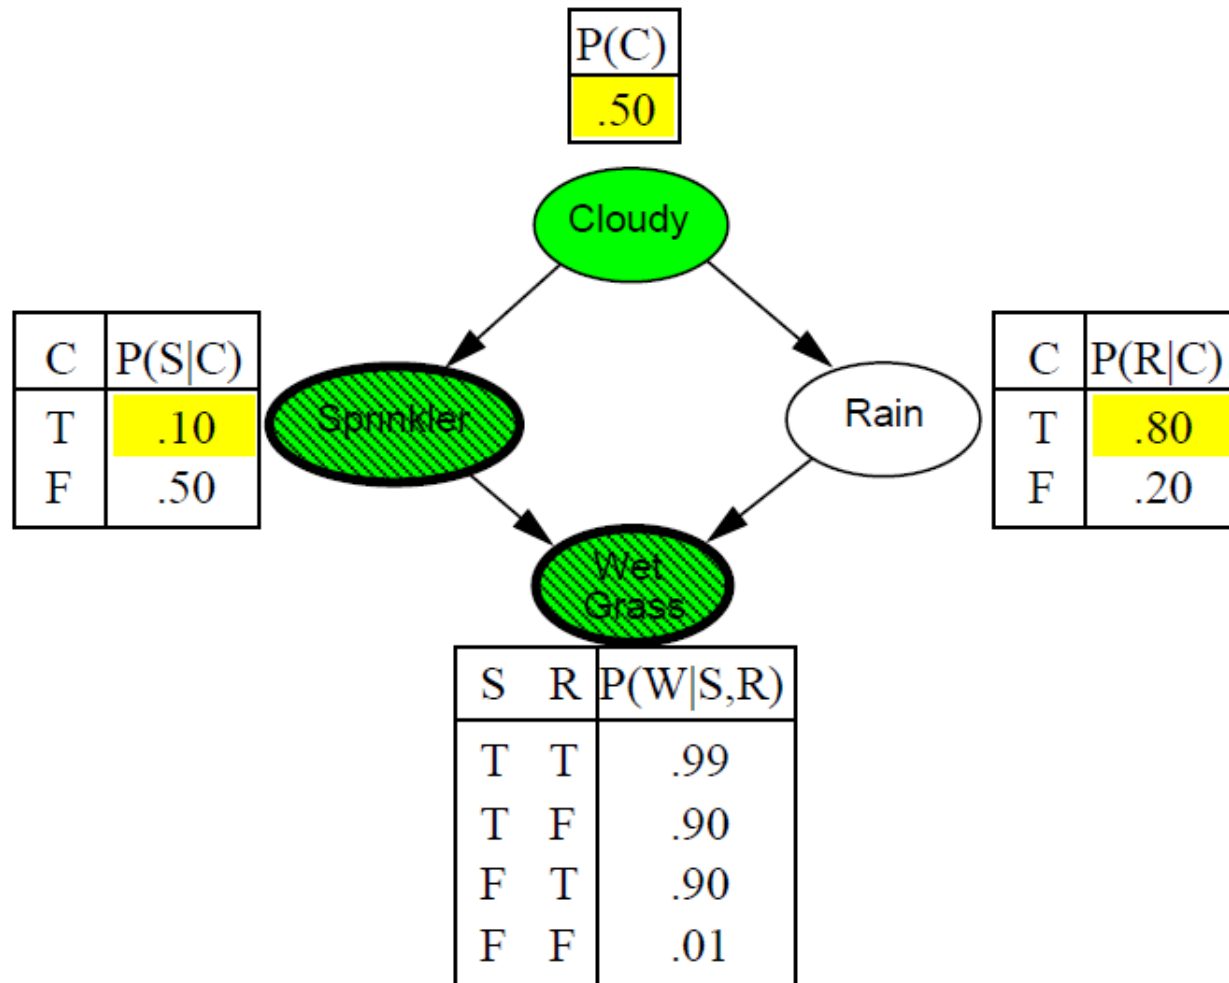
$$w = 1.0$$

# Likelihood Weighting Example



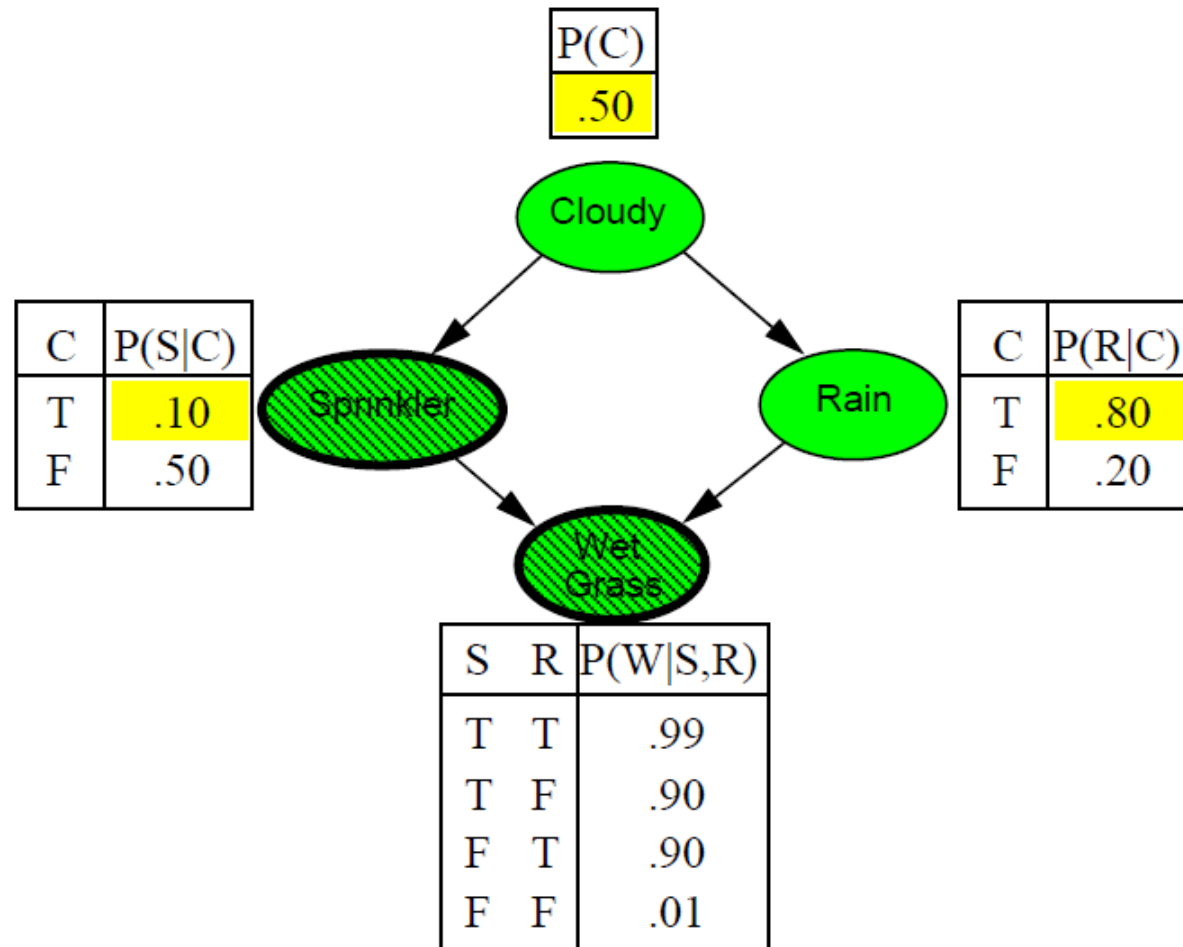
$$w = 1.0$$

# Likelihood Weighting



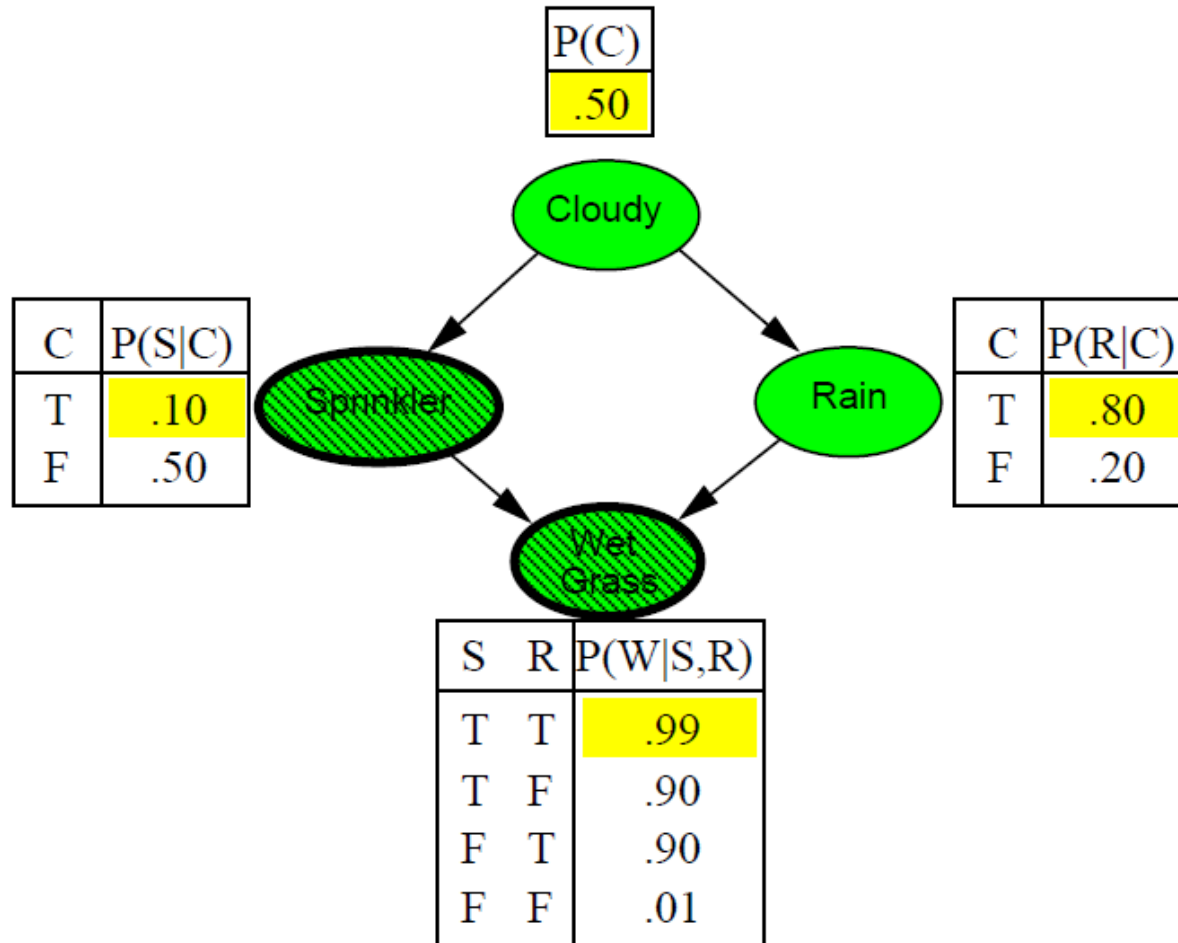
$$w = 1.0 \times 0.1$$

# Likelihood Weighting



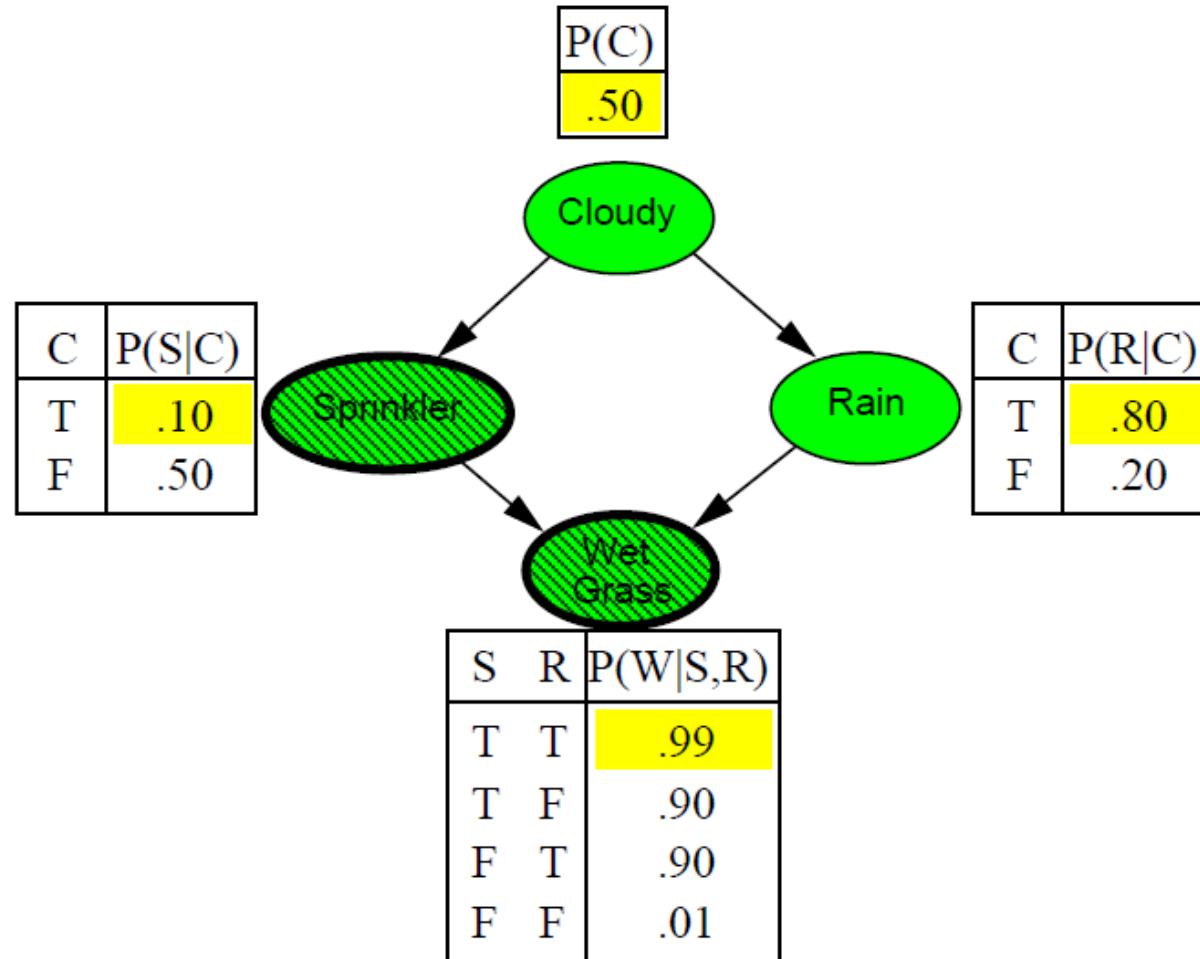
$$w = 1.0 \times 0.1$$

# Likelihood Weighting



$$w = 1.0 \times 0.1$$

# Likelihood Weighting




$$w = 1.0 \times 0.1 \times 0.99 = 0.099$$

# Analysis of Likelihood Weighting

Sampling probability for WEIGHTEDSAMPLE is

$$S_{WS}(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^l P(z_i | \text{parents}(Z_i))$$

Note: pays attention to evidence in **ancestors** only  
 $\Rightarrow$  somewhere “in between” prior and posterior distribution



Weight for a given sample  $\mathbf{z}, \mathbf{e}$  is


$$w(\mathbf{z}, \mathbf{e}) = \prod_{i=1}^m P(e_i | \text{parents}(E_i))$$

Weighted sampling probability is

$$\begin{aligned} & S_{WS}(\mathbf{z}, \mathbf{e}) w(\mathbf{z}, \mathbf{e}) \\ &= \prod_{i=1}^l P(z_i | \text{parents}(Z_i)) \prod_{i=1}^m P(e_i | \text{parents}(E_i)) \\ &= P(\mathbf{z}, \mathbf{e}) \text{ (by standard global semantics of network)} \end{aligned}$$

Hence likelihood weighting returns consistent estimates  
but performance still degrades with many evidence variables  
because a few samples have nearly all the total weight

# Likelihood Weighting Example



	Ta	Fi	Al	Le	Weight
$s_1$	true	false	true	false	$0.01 \times 0.01$
$s_2$	false	true	false	false	$0.9 \times 0.01$
$s_3$	false	true	true	true	$0.9 \times 0.75$
$s_4$	true	true	true	true	$0.9 \times 0.75$
...					
$s_{1000}$	false	false	true	true	$0.01 \times 0.75$

$$P(sm|fi) = 0.9$$

$$P(sm|\neg fi) = 0.01$$

$$P(re|le) = 0.75$$

$$P(re|\neg le) = 0.01$$



# Markov Chain Monte Carlo

“State” of network = current assignment to all variables.

Generate next state by sampling one variable given Markov blanket

Sample each variable in turn, keeping evidence fixed


```
function MCMC-Ask( $X, \mathbf{e}, bn, N$ ) returns an estimate of  $P(X|\mathbf{e})$ 
  local variables:  $\mathbf{N}[X]$ , a vector of counts over  $X$ , initially zero
                    $\mathbf{Z}$ , the nonevidence variables in  $bn$ 
                    $\mathbf{x}$ , the current state of the network, initially copied from  $\mathbf{e}$ 

  initialize  $\mathbf{x}$  with random values for the variables in  $\mathbf{Y}$ 
  for  $j = 1$  to  $N$  do
    for each  $Z_i$  in  $\mathbf{Z}$  do
      sample the value of  $Z_i$  in  $\mathbf{x}$  from  $P(Z_i|mb(Z_i))$ 
        given the values of  $MB(Z_i)$  in  $\mathbf{x}$ 
       $\mathbf{N}[x] \leftarrow \mathbf{N}[x] + 1$  where  $x$  is the value of  $X$  in  $\mathbf{x}$ 
  return NORMALIZE( $\mathbf{N}[X]$ )
```

Can also choose a variable to sample at random each time

# Markov Chain

With *Sprinkler = true*, *WetGrass = true*, there are four states:



Wander about for a while, average what you see

# MCMC

Estimate  $\mathbf{P}(Rain|Sprinkler = true, WetGrass = true)$

Sample *Cloudy* or *Rain* given its Markov blanket, repeat.  
Count number of times *Rain* is true and false in the samples.

E.g., visit 100 states

31 have *Rain = true*, 69 have *Rain = false*


$$\begin{aligned}\hat{\mathbf{P}}(Rain|Sprinkler = true, WetGrass = true) \\ = \text{NORMALIZE}(\langle 31, 69 \rangle) = \langle 0.31, 0.69 \rangle\end{aligned}$$

Theorem: chain approaches stationary distribution:  
long-run fraction of time spent in each state is exactly  
proportional to its posterior probability

# Sampling the Markov Network

Markov blanket of *Cloudy* is  
*Sprinkler* and *Rain*

Markov blanket of *Rain* is  
*Cloudy*, *Sprinkler*, and *WetGrass*



Probability given the Markov blanket is calculated as follows:

$$P(x'_i | mb(X_i)) = P(x'_i | parents(X_i)) \prod_{Z_j \in Children(X_i)} P(z_j | parents(Z_j))$$

Easily implemented in message-passing parallel systems, brains

Main computational problems:

- 1) Difficult to tell if convergence has been achieved
- 2) Can be wasteful if Markov blanket is large:

$P(X_i | mb(X_i))$  won't change much (law of large numbers)

# Summary of Inference Methods

Exact inference by variable elimination:

- polytime on polytrees, NP-hard on general graphs
- space = time, very sensitive to topology

Approximate inference by LW, MCMC:

- LW does poorly when there is lots of (downstream) evidence
- LW, MCMC generally insensitive to topology
- Convergence can be very slow with probabilities close to 1 or 0
- Can handle arbitrary combinations of discrete and continuous variables