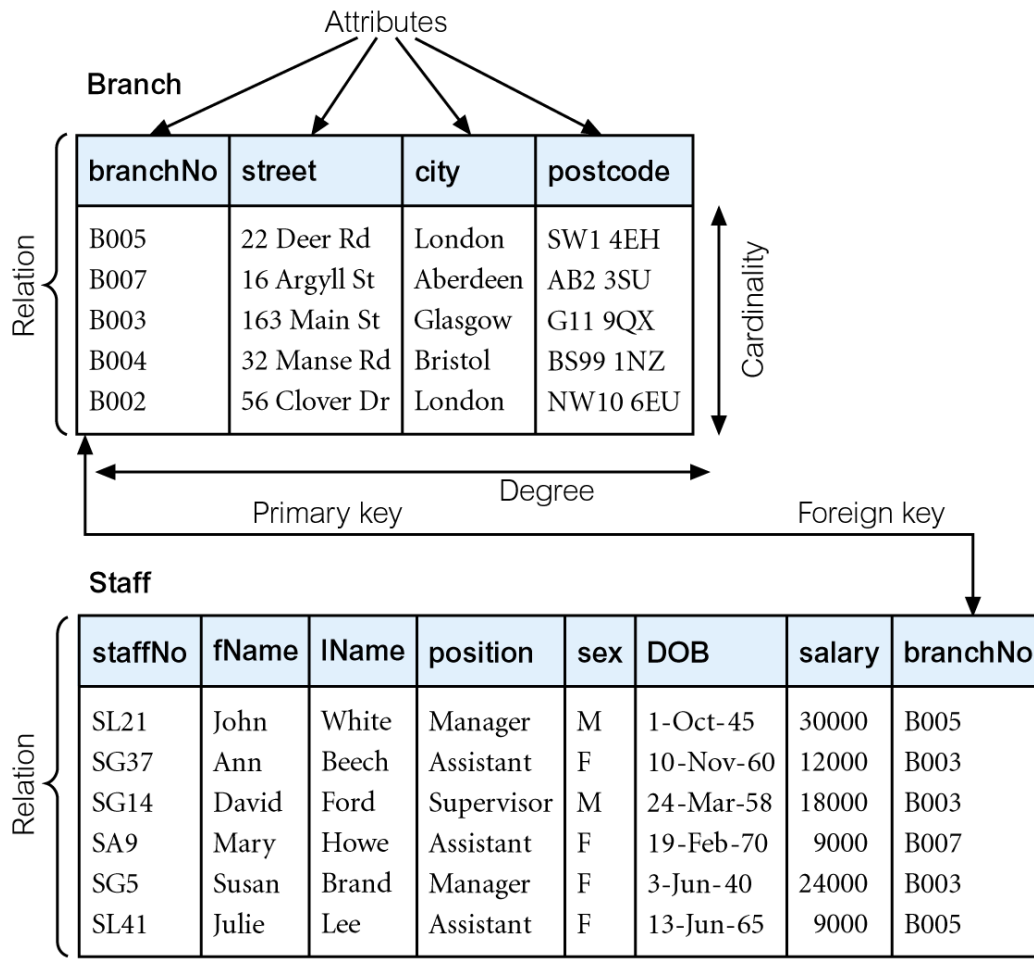


Terminology Questions Relations, Schemas,

Terminology

- Base relation: A relation we *explicitly define* in our database
 - We will later see *derived relations*, which are built from applying operations to base relations
- Relation schema: A *named* relation defined by a *set of attribute and domain name pairs*: (A1:D1, A2:D2, A3:D3)
 - Defines the structure of one table (one relation)
 - Example: Staff(staffNumber:int, name:string, salary:double)
 - Similar to a class, not to an object
- Relational database schema:
 - A collection of relation schema, where all names are unique
 - This defines the database – i.e. what is the group of tables (relations) that make up the database

Super Keys, Candidate Keys, Primary Keys



Superkey: any attribute/set Of attributes that uniquely identify

Candidate key: a minimal superkey

Primary key: One chosen candidate key

For BRANCH relation:

- branchNo, street, city, postcode
- branchNo, street, city
- branchNo, street, postcode
- branchNo, city, postcode
- branchNo, street
- branchNo, city
- branchNo, postcode

branchNo *

street, city, postcode

street, postcode

street, city (?)

Mathematical View of Relations

Definition

- Let D_i represent a domain (a finite source of values)
 - Potentially very large [all integers]
 - Or a narrow definition {1,4,5}
- We can take Cartesian product of two domains $D_1 \times D_2$ resulting in a set of ordered pairs $\{(d1, d2) \mid d1 \in D_1, d2 \in D_2\}$
 - Extends to arbitrary number of domains
 $\prod_{i=1}^n D_i$. Result are tuples of degree n .
-- Number of possible tuples is $\prod_{i=1}^n |D_i|$ == product of sizes of each domain

Definition

- Assume that we can have some property of interest that allows us to pick a subset of the tuples from our CP of domains:
 - $\{(d_1, d_2) \mid d_1 \text{ in } D_1 \text{ and } d_1 \text{ is even, } d_2 \text{ in } D_2 \text{ and } d_2 \text{ is odd}\}$
 - This gives us a subset of all possible tuples, termed a *relation*

Properties

- Since $D_1.. D_N$ are sets themselves (unique elements), their CP cannot generate repeated tuples [no repeats in tables]
- Since the subset of CP defining our relation is itself a set, its internal order is not important [table rows can be re-arranged]

Properties

- Tuples are ordered. [suggests that table columns can't be re-ordered]
 - Assume we re-order domains in schema though?
- Cardinality is a term often used to represent the size of a set [can employ on a table as well]

Translating E/R to Relational
(Chapter 17, pages 442-449)
“Logical Database Design
Methodology for the Relational
Model – Step 2.1”

Entities

- Each entity in E/R gets its own table in entity relationship, with columns for each attribute
 - Primary key for E/R model is primary key for relational model



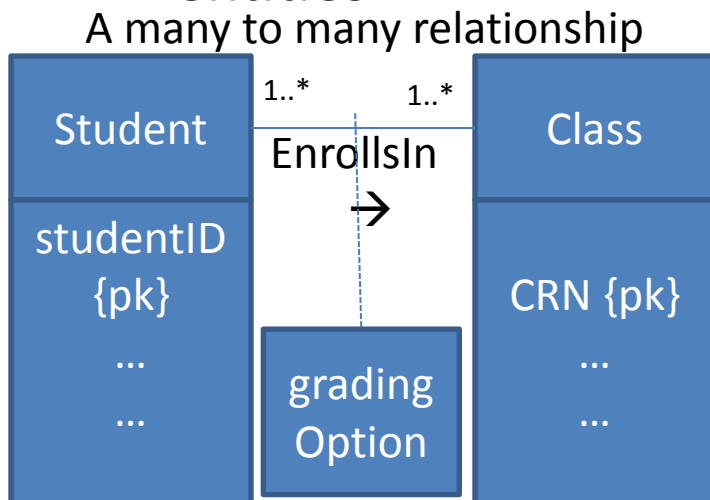
Student			
<u>studentID</u>	gpa	firstName	lastName

Relationships

- Not as straightforward
 - 1 rule that works for (most) everything
 - Build a new table to represent relationship
 - Shortcut rules
 - Add columns onto an existing table
 - Special cases
 - Think about: what to use as primary key

Relationships: New Table Approach

- General rule: Should work for every E/R relationship
 - Build a table of N ($N \geq 2$) columns where
 - Column 1: primary key of 1st entity
 - Column 2: primary key of 2nd entity
 - Those primary keys may be > 1 attribute though, so expand as needed
 - Columns 3..N: relationship attributes
 - Primary key of new table is (usually) union of primary keys of entities



+

Student Table
Class Table

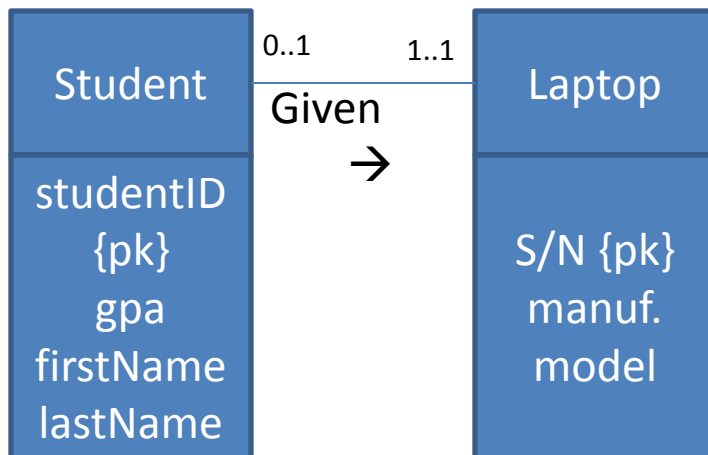
Relationships: Adding Columns Approach

- If you have a 1:1 relationship with partial mandatory participation (one of the 1s in the 1:1) add a column onto the entity table for the entity with mandatory participation with the column representing the 1 it maps to

A 1:1 relationship

Every student given 1 laptop

A laptop can be associated with no students or just 1

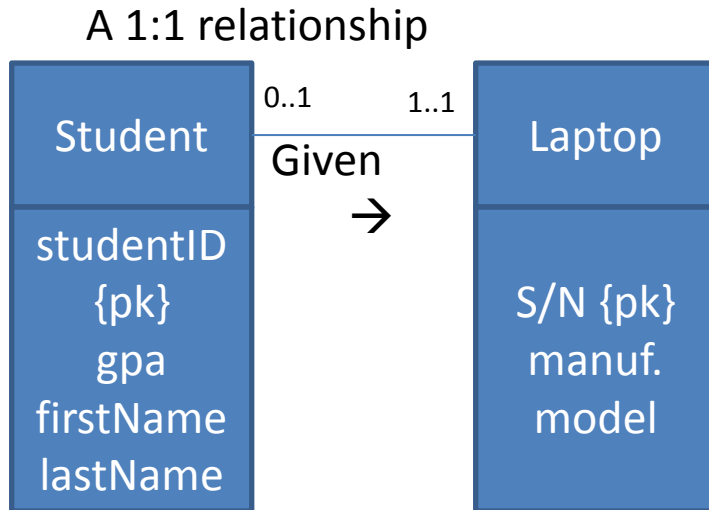


Student				
student Id	gpa	first Name	last Name	S/N

Laptop		
S/N	manuf.	model

Does not change PK of original entities

Relationships: Adding Columns Approach



This also works!

Given

student Id	S/N
---------------	-----

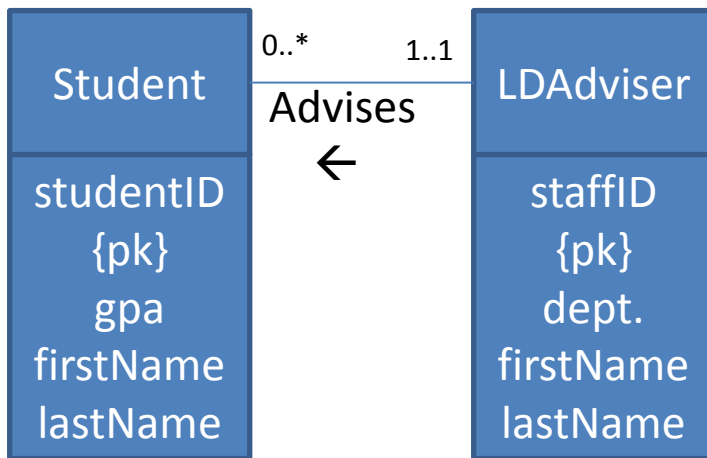
+ original Student relation
+ original Laptop relation

Why might we want to use the first example over the “new table” approach?

Relationships: Adding Columns Approach

If you have a 1:many relationship with partial (on the many side) or full mandatory participation, add a column onto the many, representing the 1

A 1:many relationship w/
mandatory participation on “many” side
(every student must have exactly 1 LDAdviser,
each LDAdviser works with zero or more students)
Student



student Id	gpa	first Name	last Name	staffID
---------------	-----	---------------	--------------	---------

LDAdviser

staffID	dept.	firstName	lastName
---------	-------	-----------	----------

Why don't we want to put this the other way - a studentID on the LDAdviser table?

Relationships: Merging

- If you have a 1:1 relationship with mandatory participation on both sides, and NO other relationships between the two entities
 - You can merge the two entity tables

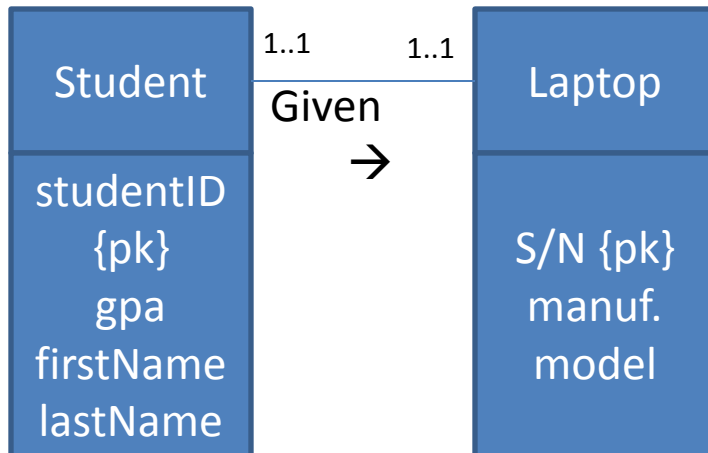
Student

<u>student tld</u>	gpa	first Name	last Name	S/N	manuf	model
------------------------	-----	---------------	--------------	-----	-------	-------

A 1:1 relationship

Every student given 1 laptop

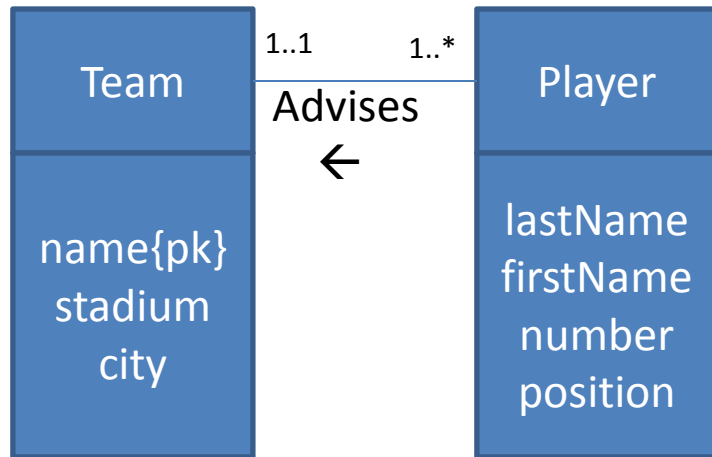
Every laptop given out, each to 1 student



Maintain one of the primary keys

Relationships: Special Cases:

Weak Entity Sets



- Player is a WeakEntity
 - No proper primary key with just Player attributes to uniquely identify Player in system (not all assigned “playerID”)
 - When associated with a team, can uniquely identify a player
 - Number is a discriminator within Team
- Special case of one-to-many
 - One team has many players
 - Mandatory participation of players on a team
 - So, add team name to player entity table
 - PK is StrongEntity primary key + WeakEntity discriminator

Player



last Name	first Name	<u>number</u>	position	<u>name (of team)</u>
-----------	------------	---------------	----------	-----------------------

Team

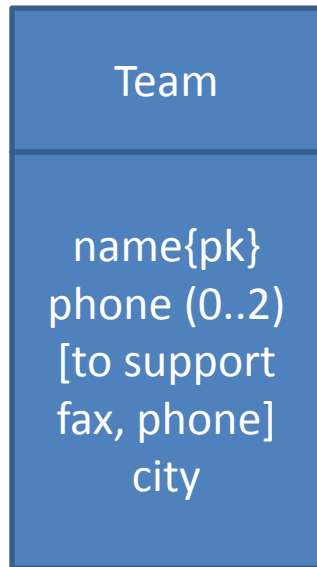
<u>name</u>	stadium	city
-------------	---------	------

Relationships: Special Cases: Composite Attributes

- Decompose a composite attribute into separate attributes and just use those



Relationships: Special Cases: Multi-Valued Attributes



- If an entity has a multi-valued attribute (such as two phone numbers), need to represent this with a separate table
- Build a 2-column table
 - 1st column: primary key of entity
 - 2nd column: setting for the multi-valued attribute
 - Primary key is all attributes
- Remove the multi-valued attribute from the original entity table

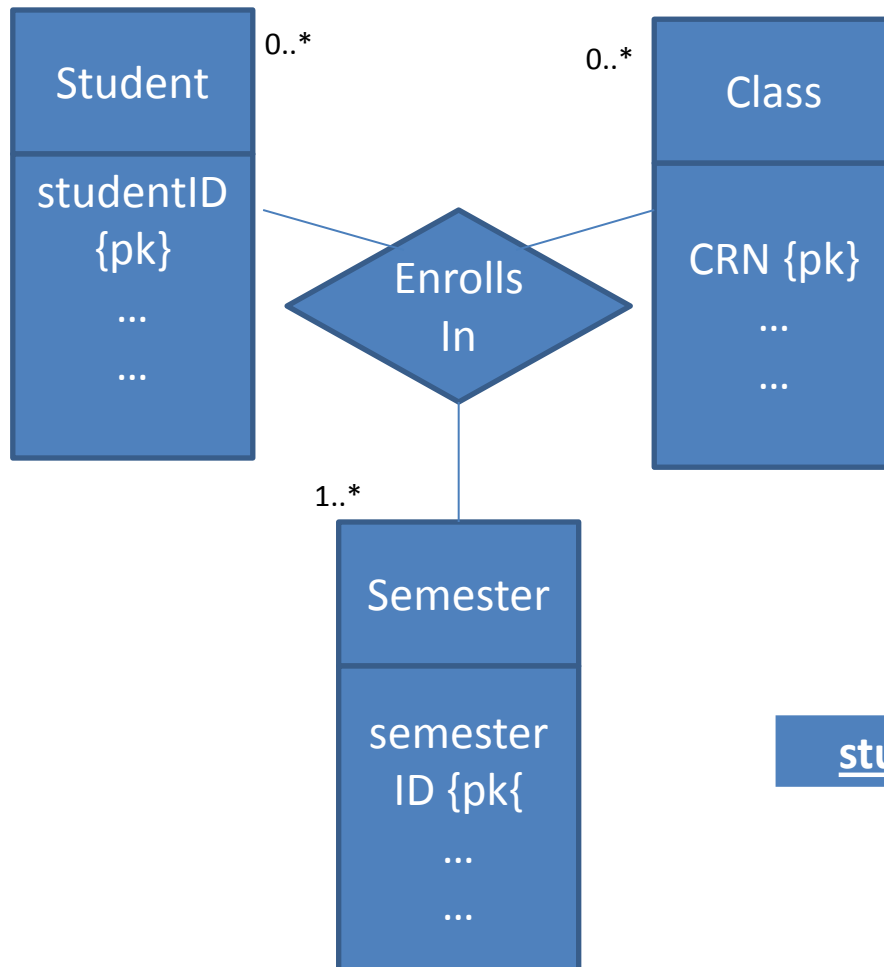
Team

<u>name</u>	city
-------------	------

TeamPhone

<u>name</u>	<u>phone</u>
-------------	--------------

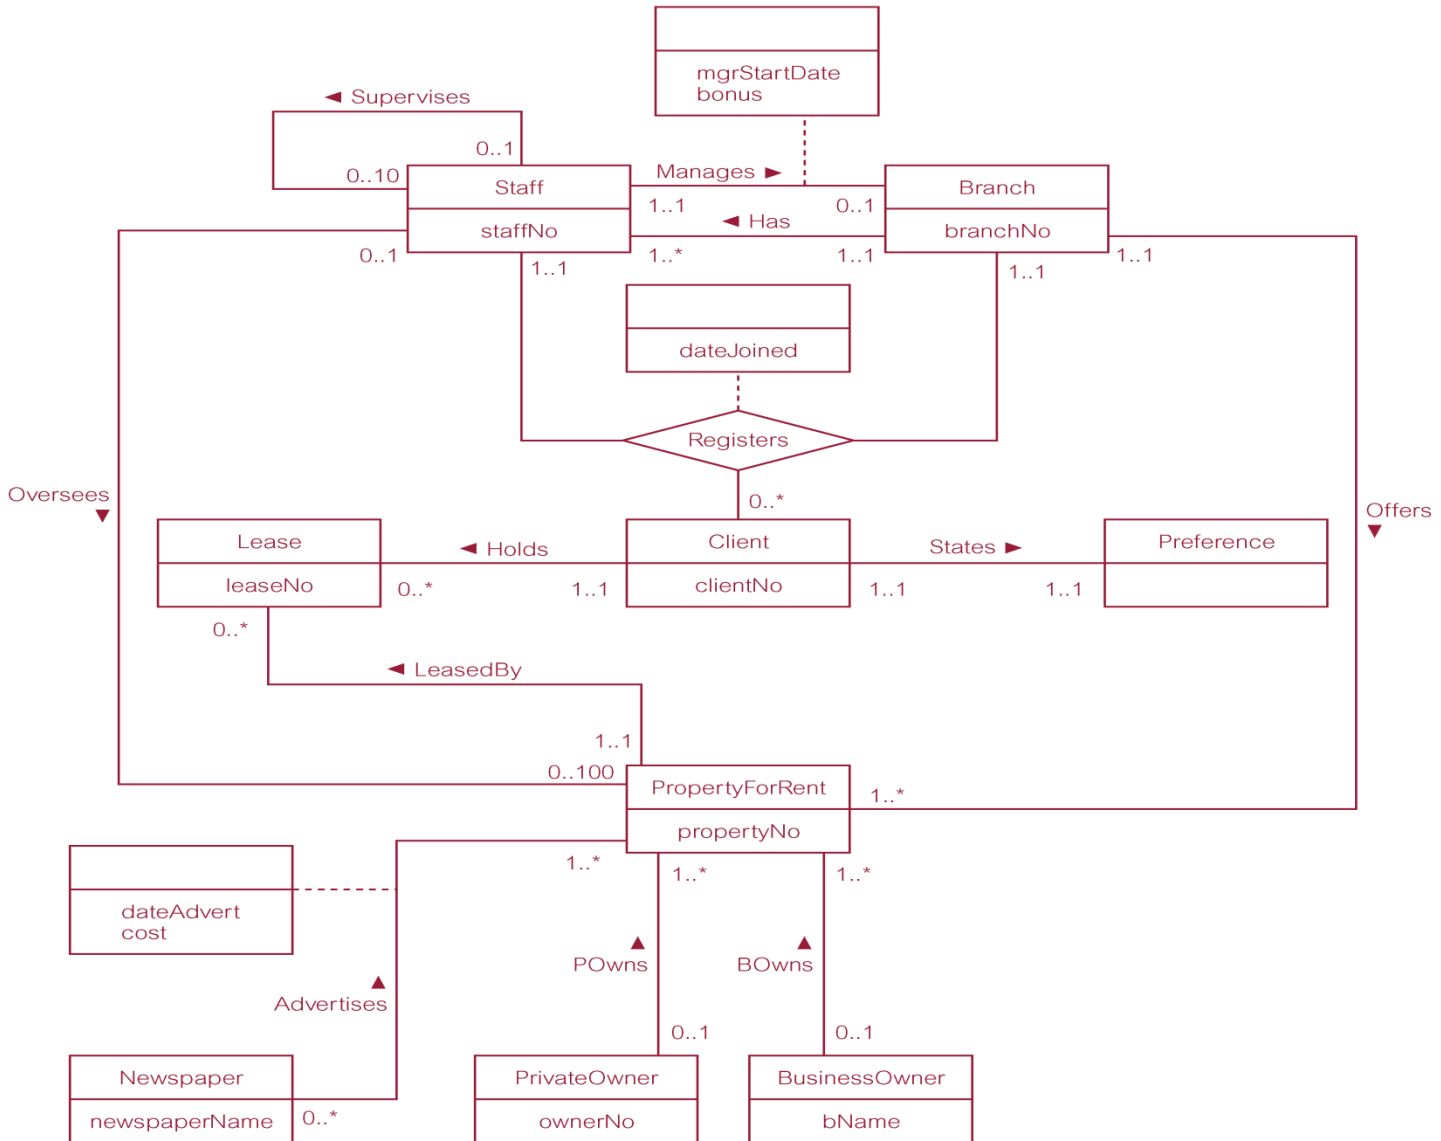
Relationships: N-ary relationships



- For (complex) n-ary relationships, just use “new-table” approach
- Primary key is union of all associated entity primary keys



Validating The Book: DreamHome E/R diagram



Validating The Book: DreamHome Staff & Branch Relationships


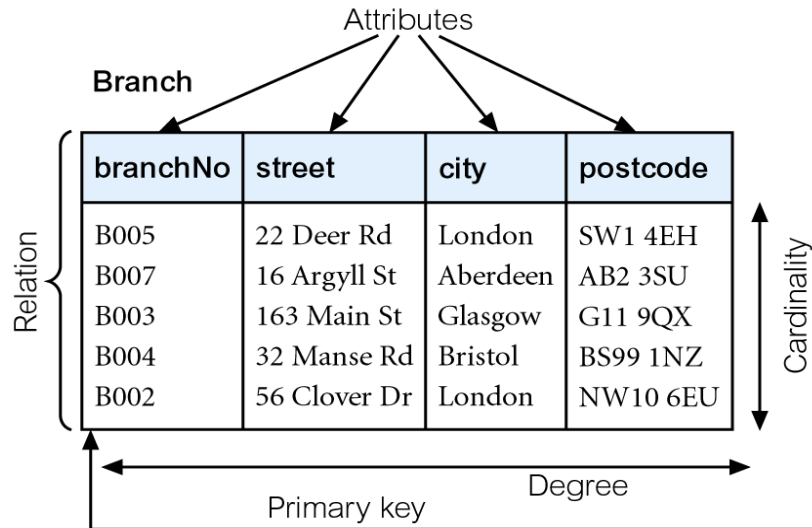


Diagram illustrating the Staff relation structure:

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

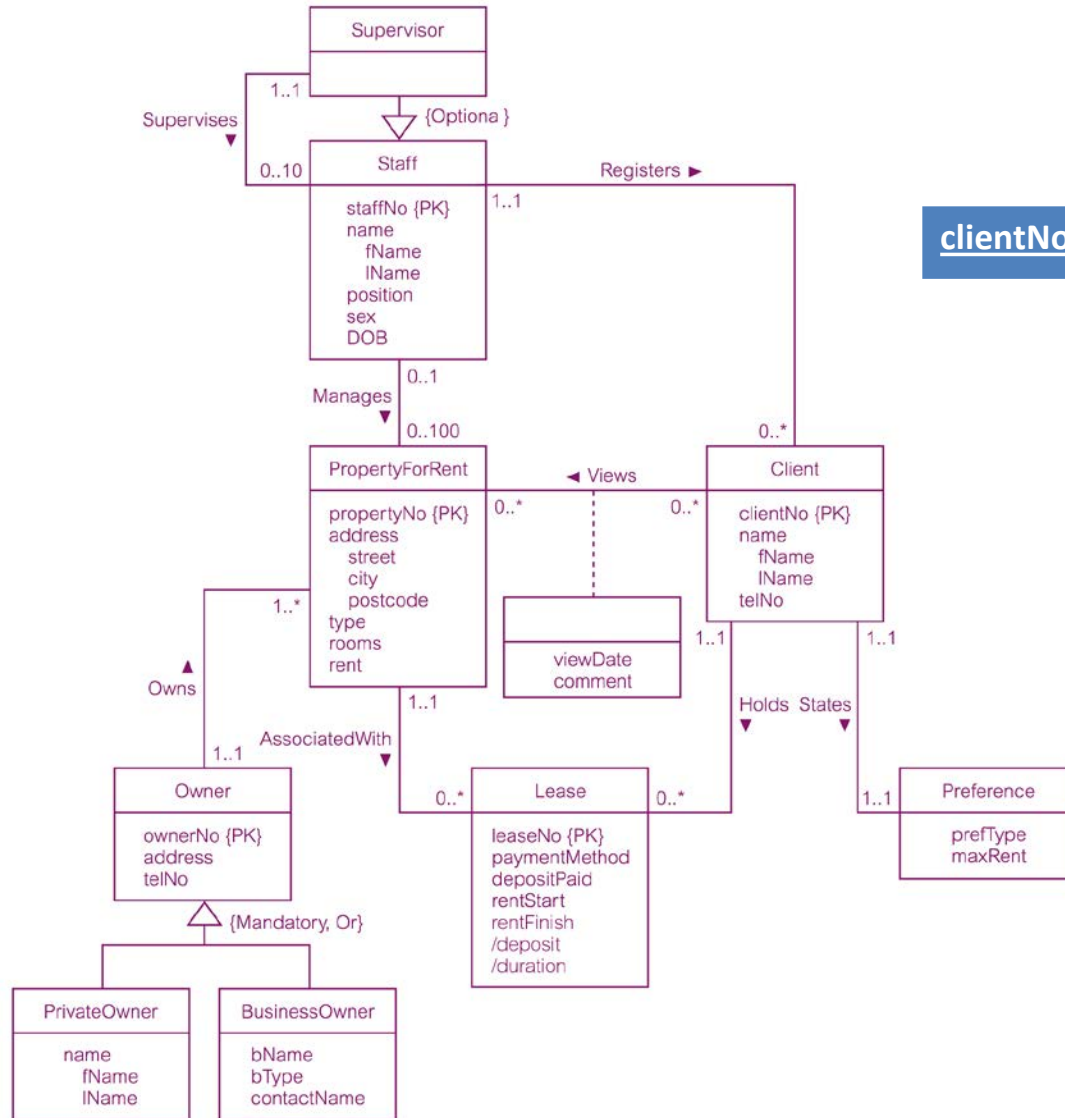
Labels: Relation, Primary key, Degree, Foreign key.

Modeling the “Branch has Staff” relationship

1 to many, mandatory participation by staff (everyone works at a branch)

Validating The Book:

Additional DreamHome Suggestions




Is it reasonable to have?

<u>clientNo</u>	fName	lName	telNo	staffNo
-----------------	-------	-------	-------	---------

Why or why not?

Validating The Book:

Additional DreamHome Suggestions



Is it reasonable to have?

clientNo

propertyNo

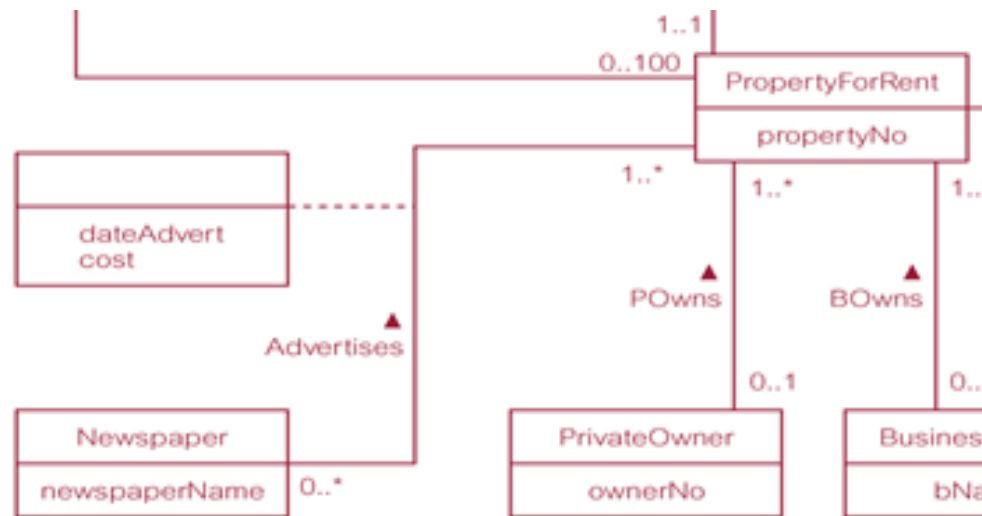
viewDate

comment

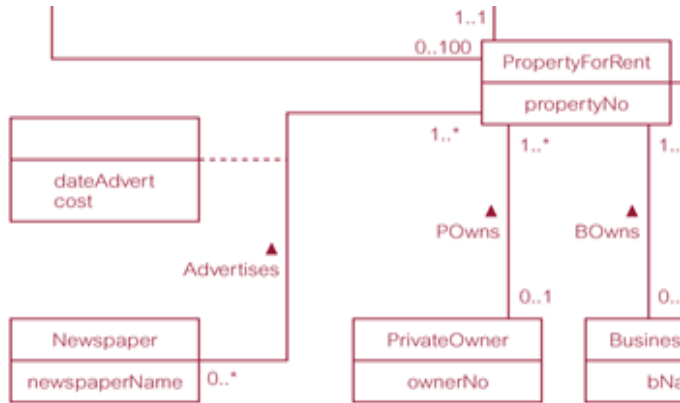
Why or why not?

Practice Problem

Build all the tables necessary to represent the *Advertises* relationship between *Newspaper* and *PropertyForRent* shown in an E/R model below:
(ignore other relationships)



Practice Problem



Many to many relationship
Optional PropertyForRent participation
Mandatory Newspaper participation
Relationship-specific attributes

Newspaper

newspaperName

PropertyForRent

propertyNumber

Advertises

newspaperName

propertyNumber

dateAdvert

cost