

Congestion Control

CSC 343-643



WAKE FOREST
UNIVERSITY


Department of Computer Science

Fall 2013

Congestion Control

Too many packets in a subnet degrades performance \Rightarrow **congestion**

- *What happens to performance under congestion?*
 - Queues at routers increase
 - Delays increase and packets are possibly dropped
- *How does congestion occur?*
 - Network resources are finite *Example resources?*



- Total input rate greater than capacity $\sum r_i > c$ for $D - E$ link

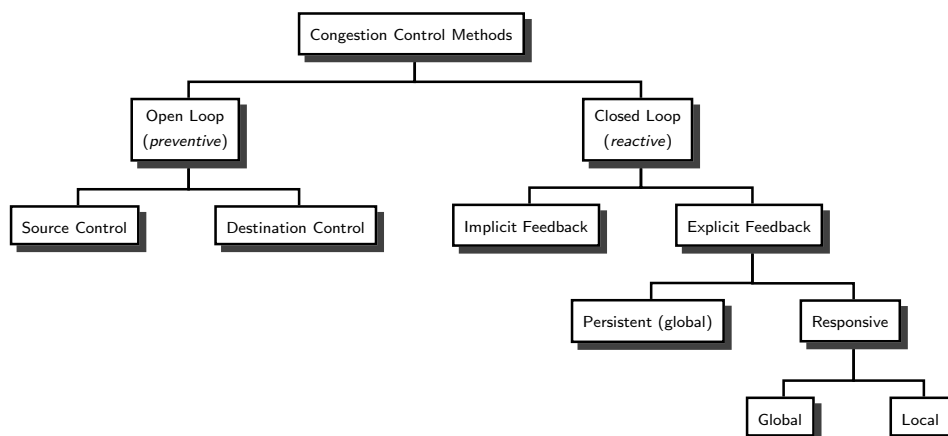
- How do we control congestion?
 - Throttle traffic
 - Increase resource supply (typically not possible)
- **Flow control** relates to point-to-point traffic (sender and receiver)
 - Assure sender does not transmit faster than receiver can accept

What is the difference between congestion and flow control?

If flow control is working properly, will congestion occur?

- We are interested in methods that control congestion

Congestion Control Categories



Control Theory

Control theory point of view: **open loop** or **closed loop**

- Open Loop
 - Preventive method \Rightarrow avoid congestion
 - No **feedback** (current system status) is used
 - Decide if/when a new session can start
- Closed Loop
 - Reactive method
 1. Monitor system
 2. Pass info to control points (feedback)
 3. Adjust system based on feedback

Is congestion control in the Internet open-loop or closed-loop?

Open Loop Congestion Control

- Control decisions do **not** depend on any feedback
 - Do **not** monitor network state
 - Requires knowledge and accounting of network resources
- Before session is allowed to start
 1. Determine session route and resource requirements
 2. Compare against resource availability
 - Form of Connection Admission Control (CAC)
 3. If allowed, must **police** session (use no more than contracted)

What are some problems associated with this approach?

Give an example, actual, network that uses open loop?

Closed Loop Congestion Control

- Control decisions based on network state
 - Monitor congestion
 - Control traffic based on measurements (feedback)
 - Feedback can be **explicit** or **implicit**
- Explicit feedback is sent as separate messages
 - Router may detect congestion and send control messages

Any problems associated with explicit feedback?

- Messages not sent in implicit feedback


If no messages, how is feedback collected?

- In addition, explicit feedback can be **persistent** or **responsive**
 - Persistent - Constant feedback provided (even if no congestion)
 - Responsive - Feedback provided only if congestion occurs
- Classic control problem, *feedback delay*
 - Measurements may be too *old* for accurate decision
 - Acting on old feedback may cause oscillations

What are some problems associated with this approach?

Burstiness

- One main cause of congestion is that traffic is **bursty**
 - Multiple packets/cells sent back-to-back
 - Prefer a predictable and **smooth** arrival pattern
- Consider MPEG compressed data



- A new frame (picture) sent every 1/30 second

We know when to expect a new frame, what is unpredictable?

- So MPEG compression has _____ bit rate and _____ quality

- Transmit 5 cells (10 bits each) during a one second interval
 - Suppose the source can send in two fashions
 1. Send a cell every 1/5 second
 2. Send a cell every 1/10 second
 - Average rate for both is $\frac{5 \times 10 \text{ bits}}{1 \text{ sec}} = 50 \text{ bps}$
 - However, second method sends all the packets in the first half of the second interval (**burst**)

Why are predictable sources more desirable?


Traffic Shaping

- Regulating average rate and burstiness of data transmission
- *Can we use the sliding window protocol?*
 - Limits amount transmitted at once (*helps manage the buffer space at the receiver*)
 - Does not limit the rate at which it is sent
- We will consider two traffic shaping mechanisms,
 - **Leaky bucket** and **token bucket**
- Shaping traffic is also referred to as **smoothing**

What type of congestion control would use traffic shaping? Open loop (preventive) or closed loop (reactive)?

Leaky Bucket Algorithm

- Model - bucket with a hole near the bottom



- Regardless of the rate water enters the bucket
 1. Outflow is constant ρ if any water present
 2. Zero if no water
- Same *bucket* idea can be applied to cells

- Conceptually, each source is connected to the network via an interface containing a leaky bucket (finite queue)
 - If a cell arrives and the queue is full, the cell is discarded
 - Otherwise it joins queue and awaits transmission
 - Server transmits one cell per *clock-tick*

What does this smell like?

- As a result, source can only send one cell per *clock-tick*

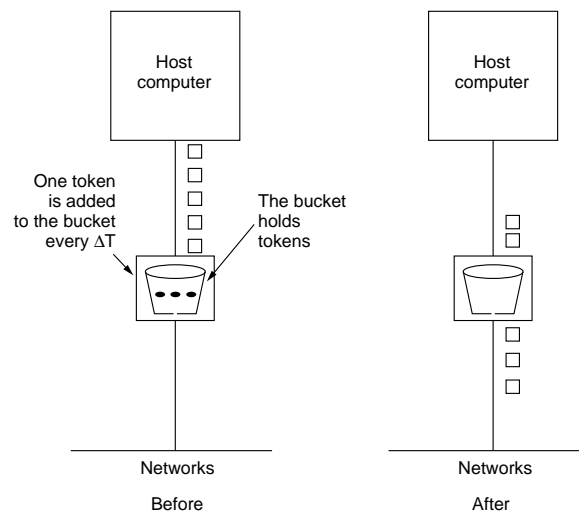
If a source is bounded by when it can send a packet, how could a higher data rate be achieved?

- If packets are the same size the algorithm works as described
 - If packets have variable size, use number of bytes per *clock-tick*
- For example, assume the rule is 1024 bytes per *clock-tick*
 - For every *clock-tick* the source can send: a single 1024 byte packet, or two 512 bytes packets, ...
 - Also called **byte-counting leaky bucket**

Token Bucket Algorithm

- Leaky bucket enforces a rigid output rate ρ
 - However, some applications need a limited *speed-up* in transmission when a burst arrives
 - **Token bucket algorithm** allows limited burst
- Token bucket algorithm design
 - Leaky bucket holds **tokens**
 - Tokens generated by a clock, one token every ΔT seconds
 - Bucket has a maximum size (capacity) of c tokens
 - For a cell to be transmitted, it must remove one token from the bucket

Bursts are allowed, what is the maximum size?



- Token bucket provides a different type of traffic shaping
 - Allows idle source to *save tokens*
 - Up to c cells can be sent at once

- What if packets (variable length) are transmitted, not cells?
 - Let each token represent k bytes
 - Packet transmitted only if enough tokens are present

What is a simple implementation of a token bucket?

- **Composite shapers** \Rightarrow Leaky bucket + token bucket
 - Token bucket regulates the maximum burst size, would like to also bound the *peak rate*
 - Send the output of the token bucket into a leaky bucket (which governs the output rate)

Token Bucket Performance

- Let s = burst length (seconds), c = bucket capacity (bytes), ρ = token arrival rate (bytes/second), and m = maximum source rate (bytes/second)
- *What is the duration of a maximum-rate burst through a token bucket?*
 - Maximum bytes sent from the token bucket during a burst is


$$c + \rho \cdot s$$

- Maximum bytes the source can send during a burst is

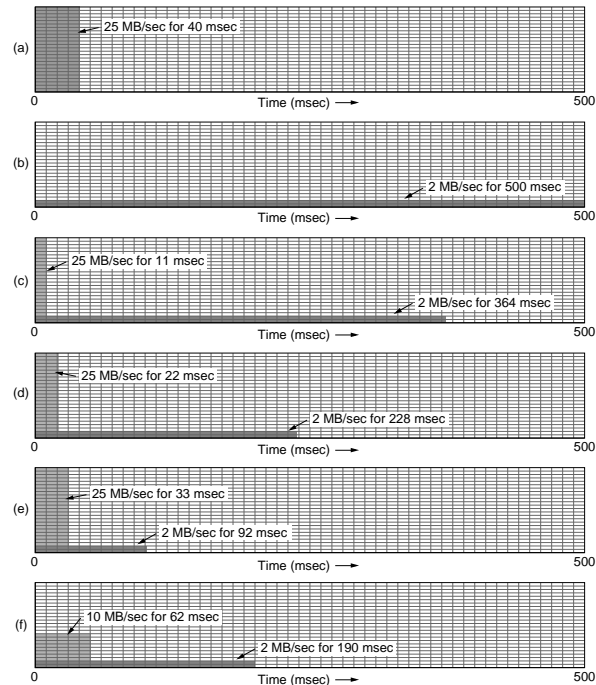
$$m \cdot s$$

- Setting the two equal and solving for s

$$s = \frac{c}{m - \rho}$$



(the larger the bucket, the burstier the output)



(a) Input. (b) Output from a leaky bucket, draining at 2 MB/sec. (c - e) Output from a token bucket with token rate at 2 MB/sec and $c = 250, 500, 750$ KB. (f) Output from 500 KB token bucket feeding a 10 MB/sec leaky bucket.

Current Traffic Shaper Ideas

- Name changes...
 - Leaky buckets are also referred to as *peak rate limiters*
 - Unfortunately, many call a *token bucket* a *leaky bucket* (token bucket is considered a leaky bucket variation).
- Possible traffic shaper modification
 - If a packet arrives and no token is available, it is *marked* then sent (instead of dropped)
 - The *mark* indicates the packet is *out-of-bounds* traffic
 - Routers are allowed to drop *marked* packets if congested

What is the advantage of this modification?

Closed Loop Congestion Control

- Assume a router can monitor the utilization of an output link
- Let s be the sampled utilization and

$$u_n = \alpha \cdot u_{n-1} + (1 - \alpha) \cdot s$$

where u is an estimate of the utilization (based on the current and previous measurements, as given by $0 \leq \alpha \leq 1$)

Why add a history component?

- Once u moves above a threshold (watermark) **choke packets** are transmitted back to the senders
- Sender reduces rate when choke packet received


What are some disadvantages to this approach?

Hop-by-Hop Choke Packets

155 Mbps sent from San Francisco (router A) to New York (router D)

- Assume router D becomes congested and send a choke packet to A
 - It would take 30 msec for the choke packet to reach A
 - During this time 4.6×10^6 bits are sent towards D
 - This data (in the pipeline) is possibly lost
 - *Upstream routers do not react to congestion*
- An alternative is to have the *upstream* routers respond
 - Choke packet reaches F which must immediately reduce its rate
 - F sends a choke packet to E which immediately reduce its rate
 - This repeats until a choke packet reaches the source

So what must be done at each upstream router?



Price-Based Congestion Control


- Choke packets cause senders to reduce at the same rate
 - What if some applications need more link capacity than others?

Why not customize the choke packets?

- Congestion pricing is a *state-less* method to allocate bandwidth
 - If you want more, then pay for it! (competitive market)
 - *A view not shared by the Canadians (AKA French)...*
- Assume we wish to control link usage
 - Iteratively price capacity based on supply and demand
 - Send prices back to users
 - If the price increases demand will fall, and vice versa

- Price calculated as

$$p_{n+1}^i = p_n^i \cdot \frac{d_n^i}{\alpha \cdot s^i}$$



- Advantages
 - No per-connection state information (one price sent)
 - All types of fairness achieved
 - Provides security

What are the disadvantages?

Implicit Closed Loop Congestion Control

- The previous congestion control examples were explicit
 - Control messages sent from the router to the source

While this provides congestion control, there are several implementation problems such as?
- Routers do **not** send messages with implicit feedback
 - Source and destination monitor data sent and received
 - React when congestion is detected

How exactly can congestion be detected?

Any problems?