

Introduction to Network Programming

CSC 343-643



Fall 2013

Socket Application Programming Interface (API)

- Introduced in 1981 Berkeley Software Distribution (BSD 4.1)
- Originally only Unix → WinSock *almost* the same
- Connects *sockets* on two hosts
- Data transfer similar to Unix file operations
- Two services provided
 - Datagram
 - Stream
- Socket interface is **generic**, it can be used with a variety of protocols (not just UDP/TCP/IP)

Review of Unix File Operations

- Unix file operations read/write to a *file descriptor*
 - An integer associated with an open *file*
 - A file can be data-file, device, network connection, ...

```
1 int fileDescriptor; // file descriptor
2 char buffer[256]; // buffer for read/write
3
4 fileDescriptor = open("file.dat", O_RDWR | O_CREAT);
5 // code storing info in buffer ...
6 write(fileDescriptor, buffer, 256);
7 close(fileDescriptor);
```

- Sockets are very similar, except for setup

Generic Data Types

Data types are needed to represent a socket and an address

- A socket is an `int` like a file descriptor
- `sockaddr` is a **generic** struct containing socket address information

```
1 struct sockaddr
2 {
3     unsigned short sa_family; // address family
4     char sa_data[14]; // protocol address
5 };
```

- `sa_family` identifies the address family (`AF_INET`)
 - `sa_data` contains a network address
- Prefer something more IP specific, which is `sockaddr_in`

Internet Data Types

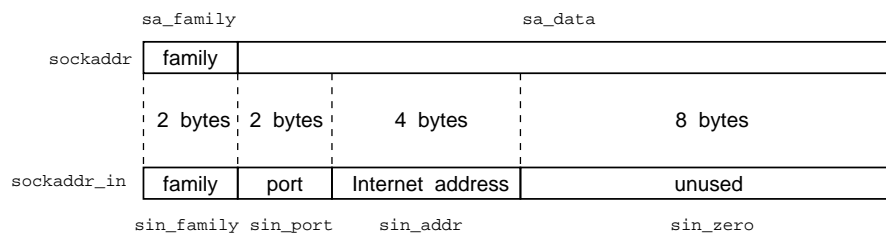
- `sockaddr_in` is a *parallel* struct containing Internet socket address information

```
1 struct sockaddr_in
2 {
3     short int      sin_family; // address family
4     unsigned short int sin_port; // port number
5     struct in_addr sin_addr;    // Internet address
6     char          sin_zero[8]; // filler, no data
7 };
```

- `in_addr` is an Internet address

```
1 struct in_addr
2 {
3     unsigned long int s_addr; // 4 bytes, IPv4 address
4 };
```

What is parallel? Why is it important?



- Socket calls are *generic* (`sockaddr`) but we will use Internet addresses (`sockaddr_in`)
 - Since the structs are the same size (parallel) → **cast**

Byte Order

As you may recall...

- There are two types of byte orderings
 - Most significant byte first (also called **network byte order**)
 - Least significant byte first

How are numbers spoken in different languages, for example 21?

- Computers (based on architecture) will use one **or** the other
 - Therefore we need to be consistent... remember to convert
- How do we convert? Use the following functions
 - `htons()` is *host to network* short
 - `htonl()` is *host to network* long
 - `ntohs()` is *network to host* short
 - `ntohl()` is *network to host* long

Address Example

Create an address for `www.cs.wfu.edu` web server

```
1 #include<arpa/inet.h>    // for sockaddr_in and inet_addr
2 #include<string.h>       // for memset
3
4 struct sockaddr_in wfuAddr; // www.cs.wfu.edu
5
6 wfuAddr.sin_family = AF_INET; // Internet address
7 wfuAddr.sin_port = htons(80); // port (network byte order)
8
9 // copy IP address (ascii to network byte order)
10 inet_aton("152.17.140.92", &(wfuAddr.sin_addr));
11
12 // zero the rest of the struct
13 memset(&(wfuAddr.sin_zero), '\0', 8);
14
15 // print the address to the screen
16 cout << inet_ntoa(wfuAddr.sin_addr) << '\n';
```

Creating a Socket

Before sending/receiving must create a socket

- Must specify, protocol family (PF_INET), type (SOCK_DGRAM or SOCK_STREAM), and protocol (0 for default)

```
1 #include <sys/socket.h> // for socket()
2
3 int sock; // stores the socket descriptor
4 if((sock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP))
5     == -1)
6 {
7     cerr << "Could not create socket\n";
8     exit(1);
9 }
```

- Returns the socket number, -1 if error
- Once done with the socket, your program must close it (like a file)

`close(sock);`

Is the above socket UDP or TCP?

TCP or UDP

- There are two types of services available in the Internet
 - **User Datagram Protocol** (UDP) is an unreliable transmission
 - **Transport Control Protocol** (TCP) is a reliable stream

- Service must be specified once the socket is created

– Creating the socket with SOCK_DGRAM uses UDP

```
1 sock = socket(PF_INET, SOCK_DGRAM, 0)
```

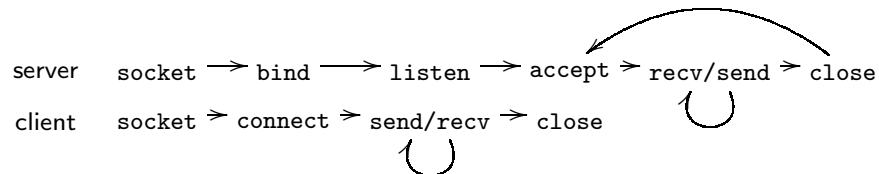
– Creating the socket with SOCK_STREAM used TCP

```
1 sock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)
```

- Functions used to transmit data differ based on the service

TCP Client Server Model

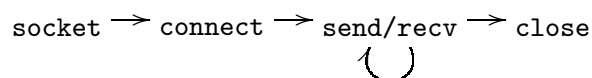
- Network programming typically uses a *client server* model
 - The server passively listens for communication
 - The client initiates the communication
- For TCP, once the server hears a connection request, handshake
 - Determine sequence numbers and buffer space
 - Once connection establish can send and receive data



What is the difference for UDP?

TCP Client

- TCP client must perform the following four steps
 1. Create a TCP socket using `socket`
 2. Establish connection to server using `connect`
 3. Communicate using `send` and `recv`
 4. Close the connection with `close`



- Already discussed how to create a socket
 - Be certain the socket is `SOCK_STREAM` for TCP

TCP Client connect

To connect the client program must specify

- Socket, server address, and address length

```
1  int connect(int sock, struct sockaddr * serverAddr,
2             unsigned int addrLength);
```

- Assume we want to connect to a server

```
1  if(connect(sock, (struct sockaddr *) &serverAddr,
2             sizeof(struct sockaddr)) < 0)
3  {
4      cerr << "Could not connect to server\n";
5      exit(1);
6  }
```

- Connect will perform three-way handshake
- Once connected, client and server can send and recv data

Would UDP use the connect function?

TCP send

To send data (**TCP only**) the program must supply

- Socket, pointer to data, data length, and flags

```
1  int send(int sock, const void *msg,
2           unsigned int msgLength, int flags);
```

- Assume we want to send a C-string

```
1  char str[] = "Pluf still rules";
2  int strLength = strlen(str) + 1;
3  if(send(sock, str, strLength, 0) != strLength)
4  {
5      cerr << "Could not send all the data\n";
6      exit(1);
7  }
```

- send() will return the number of bytes sent

How does send know which address and port to send to?

TCP recv

To recv data (**TCP only**) the program must supply

- Socket, pointer to buffer, maximum buffer length, and flags

```
1 int recv(int sock, void *buffer,  
2         unsigned int bufferLength, int flags);
```

- Assume we want to receive data from a TCP socket

```
1 char buffer[256];  
2 int numBytes;  
3 if((numBytes = recv(sock, buffer, 256, 0)) <= 0)  
4 {  
5     cerr << "Could not recv from socket\n";  
6     exit(1);  
7 }
```

- recv() returns the number of bytes read
 - Blocks until data received or timeout
 - Data is placed in the buffer

Closing a Socket

Once communication over a socket is complete, you must close

- Specify the socket and use the close() function

```
1 #include<unistd.h>  
2  
3 close(sock);
```

- Cannot read/write to the socket after close()

Why is close necessary?

TCP Time Service

TCP/IP defines a service (port 13) that allows one machine to obtain the current data and time from another. The following **client** program will obtain the time from the **server** 152.17.140.3

```
1  #include<iostream>
2  #include<sys/socket.h>    // socket(), send(), recv()
3  #include<arpa/inet.h>     // struct socket sockaddr
4  #include<unistd.h>        // close()
5  #include<string.h>        // memset()
6  #include<cstdlib>         // exit()
7  using namespace std;
8
9  int main()
10 {
11     int sock; // socket for datagram communication
12     if((sock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP))
13        == -1)
14     {
15         cerr << "Could not create socket\n";
16         exit(1);
17     }
18
19     struct sockaddr_in srvAddr;           // time server
20     srvAddr.sin_family = AF_INET;        // Internet addr
21     srvAddr.sin_port = htons(13);        // port 13
22     srvAddr.sin_addr.s_addr = inet_addr("152.17.140.3");
23     memset(&(srvAddr.sin_zero), '\0', 8); // set to zero
24
25     if(connect(sock, (struct sockaddr *) &srvAddr,
26                sizeof(struct sockaddr)) == -1)
27     {
28         cerr << "Could not connect to server\n";
29         exit(2);
30     }
31
32     // to receive the time, just ask the question...
33     char* msg = "What time is it?";
34     if(send(sock, msg, strlen(msg) + 1, 0) == -1)
35     {
36         cerr << "Could not send to socket\n";
37         exit(3);
38     }
39 }
```

```

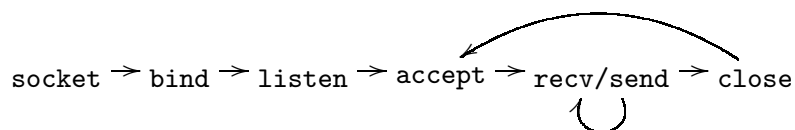
1  char buffer[256]; // stores data received from socket
2  // will block until datagram received from socket
3  if(recv(sock, buffer, 256, 0) == -1)
4  {
5      cerr << "Could not receive from socket\n";
6      exit(4);
7  }
8
9  cout << buffer << '\n';
10 close(sock);
11 return 0;
12 }

```



TCP Server

- TCP server must complete the following six steps
 1. Create a TCP socket using `socket`
 2. Assign a port number to the socket using `bind`
 3. Tell system to allow connections to port using `listen`
 4. Accept connection request using `accept`
 5. Communicate using `recv` and `send` via **new** socket
 6. Close the connection using `close`



- Already know about `socket`, `recv`, `send`, and `close`

TCP bind

To bind a port to an address the program must supply

- Socket, address, and address length

```
1  int bind(int sock, struct sockaddr *localAddr,  
2          unsigned int addrLength);
```

- Assume we wanted to bind to port 1122

```
1  struct sockaddr_in localAddr;           // local addr  
2  memset(&localAddr, 0, sizeof(localAddr)); // zero addr  
3  localAddr.sin_family = AF_INET;         // Internet  
4  localAddr.sin_addr.s_addr = htonl(INADDR_ANY); // any  
5  localAddr.sin_port = htons(1122);       // local port  
6  
7  if(bind(sock, (struct sockaddr *) &localAddr,  
8          sizeof(struct sockaddr)) < 0)  
9  {  
10     cerr << "Could not bind to socket\n";  
11     exit(1);  
12 }
```

TCP listen

To listen for a connection request the program must supply

- Socket and queue limit

```
1  int listen(int sock, int queueLimit);  
    – queueLimit is the maximum outstanding requests
```

- Assume we wanted to listen for incoming connections

```
1  #define MAX_PENDING 5 // maximum outstanding req  
2  if(listen(sock, MAX_PENDING) < 0)  
3  {  
4      cerr << "Could not listen on socket\n";  
5      exit(1);  
6  }
```

- listen() returns 0 on success, -1 otherwise
 – The queue is used by accept

TCP accept

To accept an incoming connection the program must supply

- Socket, client address, and address length

```
1  int accept(int sock, struct sockaddr *clientAddress,
2             unsigned int *addrLength);
```

- Assume we wanted to accept an incoming connection

```
1  int clientSock;
2  // socket for connection
3  struct sockaddr_in clientAddr; // incoming address
4  unsigned int addrLength = sizeof(struct sockaddr_in);
5  if((clientSock = accept(sock, (struct sockaddr *)
6                          &clientAddr, &addrLength)) < 0)
7  {
8      cerr << "Could not accept connection\n";
9      exit(1);
10 }
```

- `accept()` returns a descriptor for a *new* socket
 - Dequeues the next connection on the socket queue
 - Creates a **new** socket for arriving connection
 - Sets the address and address length variables
 - If the queue is empty, then **block**
 - If error then `accept` returns -1
- After `accept()` program can `recv()` and `send()`
 - As described there are two sockets, one for receiving a connection request and another created for the connection

Which socket is used for sending and which for receiving?
 - The socket that has been bound to a port and marked *listening* is **never** used for sending and receiving in TCP

Example TCP Client Server Program

Write a TCP client and a server that will send/receive C-strings

- Client will send the message "Hello from client" to the server
 - Assume the destination port is 1848
- Server will reply to the client "Hello w.x.y.z"
 - Listen for a connection on port 1848
 - Must be started before client program

Why must the server be running before the client starts?

TCP C-String Client Program

```
1  #include<iostream>
2  #include<sys/socket.h>      // socket, send(), recv()
3  #include<arpa/inet.h>      // struct sockaddr
4  #include<unistd.h>         // close()
5  #include<string.h>         // memset()
6  #include<cstdlib>          // exit()
7  using namespace std;
8
9  int main()
10 {
11     int sock; // socket for stream communication
12     if((sock = socket(PF_INET, SOCK_STREAM,
13                       IPPROTO_TCP)) == -1)
14     { cerr << "Could not create socket\n"; exit(1); }
15
16     struct sockaddr_in srvAddr; // address of server
17     srvAddr.sin_family = AF_INET; // Internet address
18     srvAddr.sin_port = htons(1848); // port 1848
19     srvAddr.sin_addr.s_addr = inet_addr("152.17.140.17");
20     memset(&(srvAddr.sin_zero), '\0', 8); // set to zero
21
22
23
```

```

24     // establish connection with server
25     if(connect(sock, (struct sockaddr *) &srvAddr,
26               sizeof(struct sockaddr)) < 0)
27     { cerr << "Could not connect\n"; exit(2); }
28
29     char* msg = "Hello from client"; // c-string to send
30     if(send(sock, msg, strlen(msg) + 1, 0) == -1)
31     { cerr << "Could not send to socket\n"; exit(3); }
32
33     char buffer[256]; // stores data received from socket
34     unsigned int addrLength = sizeof(sockaddr);
35     int numBytes = 0;
36     if((numBytes = recv(sock, buffer, 256, 0)) <= 0)
37     { cerr << "Could not receive from socket\n"; exit(4); }
38     cout << buffer << '\n';
39
40     close(sock);
41     return 0;
42 }

```

TCP C-String Server Program

```

1  #include<iostream>
2  #include<sys/socket.h> // socket, send(), recv()
3  #include<arpa/inet.h> // struct sockaddr
4  #include<unistd.h> // close()
5  #include<string.h> // memset()
6  #include<cstdlib> // exit()
7  using namespace std;
8
9  int main()
10 {
11     int sock; // socket for datagram communication
12     if((sock = socket(PF_INET, SOCK_STREAM,
13                       IPPROTO_TCP)) == -1)
14     { cerr << "Could not create socket\n"; exit(1); }
15
16     struct sockaddr_in myAddr; // my (server) addr
17     myAddr.sin_family = AF_INET; // Internet addr
18     myAddr.sin_port = htons(1848); // port 1848
19     myAddr.sin_addr.s_addr = INADDR_ANY; // my address
20     memset(&(myAddr.sin_zero), '\0', 8); // set to zero

```

```

1  if(bind(sock, (struct sockaddr*) &myAddr,
2      sizeof(struct sockaddr)) == -1)
3  { cerr << "Could not bind to port\n"; exit(2); }
4
5  if(listen(sock, 0) < 0)
6  { cerr << "Could not listen to socket\n"; exit(3); }
7
8  int clientSock; // client socket
9  struct sockaddr_in clientAddr; // client addr
10 unsigned int clientAddrLength =
11     sizeof(struct sockaddr_in);
12
13 if((clientSock = accept(sock, (struct sockaddr *)
14     &clientAddr, &clientAddrLength)) < 0)
15 { cerr << "Could not accept connection\n"; exit(4); }
16
17 char buffer[256]; // data received from socket
18 if(recv(clientSock, buffer, 256, 0) == -1)
19 { cerr << "Could not receive from socket\n"; exit(3); }
20 cout << buffer << '\n';

```

```

1  char *reply = new char[strlen("Hello ") +
2      strlen(inet_ntoa(clientAddr.sin_addr)) + 1];
3  strcpy(reply, "Hello ");
4  strcat(reply, inet_ntoa(clientAddr.sin_addr));
5  if(send(clientSock, reply, strlen(reply) + 1, 0) == -1)
6  { cerr << "Could not send to socket\n"; exit(4); }
7
8  close(sock);
9  return 0;
10 }

```

Terminal

```

> tcpStrSrv 1
Hello from client 3
>

```

Terminal

```

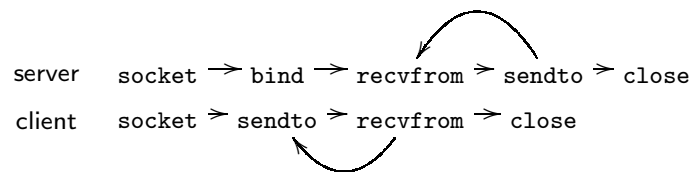
> tcpStrCli 2
Hello 152.17.140.12 4
>

```

UDP Sockets

Datagram sockets are connectionless and *do not have connection setup*

- Client
 - socket, create new end point
 - sendto and recvfrom, transmit data
- Server
 - socket, creates a new end point
 - bind, binds an *address* to the connection
 - sendto and recvfrom, transmit data



- Client does **not** establish a connection
 - connect is not needed
- Server does not *accept* a connection
 - listen is not needed
 - accept is not needed

What is the implication of no established connection?

Does the server still wait for datagrams?

*In TCP **accept** assigned a new socket to the incoming data, how will communication take place in UDP?*

Datagram sendto()

To sendto() the program must specify

- Socket, data, data length, flags, *to* address, *to* address length

```
1 int sendto(int sock, const void *msg, int len,
2           unsigned int flags,
3           const struct sockaddr *toAddr,
4           int toLength);
```

- Assume we want to send a C-string

```
1 char msg[] = "What time is it?";
2 sendto(sock, msg, strlen(msg) + 1, 0,
3        (struct sockaddr*) &srvAddr,
4        sizeof(struct sockaddr));
```

- sendto() will return the number of data bytes sent, -1 if error

Datagram recvfrom()

Similarly, to recvfrom() the program must specify

- Socket, buffer (to store data), buffer length, flags, *from* address, and *from* address length

```
1 int recvfrom(int sock, void *buffer, int bufferLength,
2             unsigned int flags,
3             struct sockaddr *fromAddr,
4             int *fromLength);
```

- Assume we want to receive data from a datagram

```
1 char buffer[256];
2 struct sockaddr_in fromAddr;
3 unsigned int addrLength = sizeof(struct sockaddr);
4 recvfrom(sock, buffer, 256, 0, (struct sockaddr*)
5        &fromAddr, &addrLength);
```

- recvfrom() returns the number of data bytes, -1 if error
 - fromAddr stores the datagram address
 - addrLength stores the address length *must be initialized*

UDP Time Service

TCP/IP defines a service (port 13) that allows one machine to obtain the current data and time from another. The following program will obtain the time from 152.17.140.3

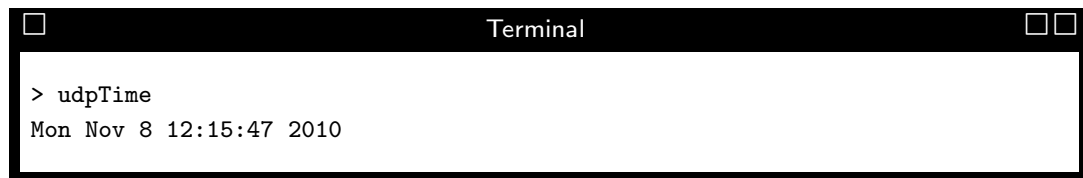
```
1  #include<iostream>
2  #include<sys/socket.h>      // socket, sendto(), recvfrom()
3  #include<arpa/inet.h>       // struct sockaddr
4  #include<unistd.h>          // close()
5  #include<string.h>          // memset()
6  #include<cstdlib>           // exit()
7  using namespace std;
8
9  int main()
10 {
11     int sock; // socket for datagram communication
12     if((sock = socket(PF_INET, SOCK_DGRAM, 0)) == -1)
13     {
14         cerr << "Could not create socket\n";
15         exit(1);
16     }
```

```
1     struct sockaddr_in srvAddr;          // address of server
2     srvAddr.sin_family = AF_INET;        // Internet address
3     srvAddr.sin_port = htons(13);        // port 13
4     srvAddr.sin_addr.s_addr = inet_addr("152.17.140.13");
5     memset(&(srvAddr.sin_zero), '\0', 8); // set to zero
6
7     // to receive the time, just ask the question...
8     char* msg = "What time is time?";
9     if(sendto(sock, msg, strlen(msg) + 1, 0,
10               (struct sockaddr*) &srvAddr,
11               sizeof(struct sockaddr)) == -1)
12     {
13         cerr << "Could not send to socket\n";
14         exit(2);
15     }
```

```

1  char buffer[256]; // stores data received from socket
2  unsigned int addrLength = sizeof(sockaddr);
3
4  // will block until datagram received from socket
5  if(recvfrom(sock, buffer, 256, 0, (struct sockaddr*)
6      &srvAddr, &addrLength) == -1)
7  {
8      cerr << "Could not receive from socket\n";
9      exit(3);
10 }
11 close(sock);
12 cout << buffer << '\n';
13 return 0;
14 }

```



```

Terminal
> udpTime
Mon Nov 8 12:15:47 2010

```

Example UDP Client Server Program

Write a UDP client and a server that will send/receive C-strings

- Client will send the message "Hello from client" to the server
 - Assume the destination port is 1848
- Server will reply to the client "Hello *w.x.y.z*"
 - Listen for datagrams on port 1848
 - Must be started before client program

UDP C-String Client Program

```
1  #include<iostream>
2  #include<sys/socket.h>    // socket, sendto(), recvfrom()
3  #include<arpa/inet.h>    // struct sockaddr
4  #include<unistd.h>       // close()
5  #include<string.h>       // memset()
6  #include<cstdlib>        // exit()
7  using namespace std;
8  int main()
9  {
10     int sock; // socket for datagram communication
11     if((sock = socket(PF_INET, SOCK_DGRAM, 0)) == -1)
12     {
13         cerr << "Could not create socket\n";
14         exit(1);
15     }
16     struct sockaddr_in srvAddr; // address of server
17     srvAddr.sin_family = AF_INET; // Internet address
18     srvAddr.sin_port = htons(1848); // port 1848
19     srvAddr.sin_addr.s_addr = inet_addr("152.17.140.17");
20     memset(&(srvAddr.sin_zero), '\0', 8); // set to zero
```

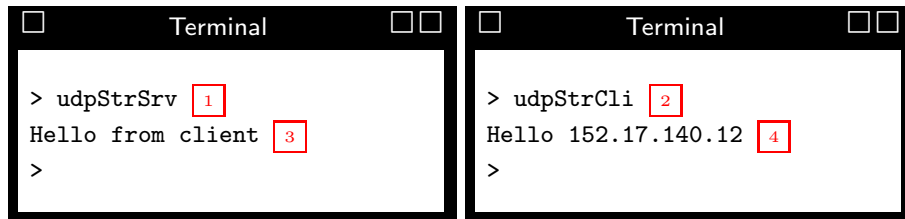
```
1  char* msg = "Hello from client"; // c-string to send
2  if(sendto(sock, msg, strlen(msg) + 1, 0,
3           (struct sockaddr*) &srvAddr,
4           sizeof(struct sockaddr)) == -1)
5  {
6      cerr << "Could not send to socket\n";
7      exit(2);
8  }
9
10 char buffer[256]; // stores data received from socket
11 unsigned int addrLength = sizeof(sockaddr);
12 if(recvfrom(sock, buffer, 256, 0,
13            (struct sockaddr*) &srvAddr,
14            &addrLength) == -1)
15 {
16     cerr << "Could not receive from socket\n";
17     exit(3);
18 }
19 cout << buffer << '\n';
20
21 close(sock);
22 return 0;
23 }
```

UDP C-String Server Program

```
1  #include<iostream>
2  #include<sys/socket.h>      // socket, sendto(), recvfrom()
3  #include<arpa/inet.h>      // struct sockaddr
4  #include<unistd.h>         // close()
5  #include<string.h>         // memset()
6  #include<cstdlib>          // exit()
7  using namespace std;
8  int main()
9  {
10     int sock; // socket for datagram communication
11     if((sock = socket(PF_INET, SOCK_DGRAM, 0)) == -1)
12     {
13         cerr << "Could not create socket\n";
14         exit(1);
15     }
16     struct sockaddr_in myAddr; // my (server) addr
17     myAddr.sin_family = AF_INET; // Internet address
18     myAddr.sin_port = htons(1848); // port 1848
19     myAddr.sin_addr.s_addr = INADDR_ANY; // my address
20     memset(&(myAddr.sin_zero), '\0', 8); // set to zero
21
22
23
```

```
24     if(bind(sock, (struct sockaddr*) &myAddr, sizeof(struct
25         sockaddr)) == -1)
26     {
27         cerr << "Could not bind to port\n";
28         exit(1);
29     }
30
31     struct sockaddr_in destAddr; // dest (client) address
32     unsigned int destAddrLength = sizeof(struct sockaddr);
33
34     char buffer[256]; // data from socket
35     if(recvfrom(sock, buffer, 256, 0, (struct sockaddr*)
36         &destAddr, &destAddrLength) == -1)
37     {
38         cerr << "Could not receive from socket\n";
39         exit(3);
40     }
41
42     cout << buffer << '\n';
43
44     char *reply = new char[strlen("Hello ") +
45         strlen(inet_ntoa(destAddr.sin_addr)) + 1];
46     strcpy(reply, "Hello ");
47     strcat(reply, inet_ntoa(destAddr.sin_addr));
48
49     if(sendto(sock, reply, strlen(reply) + 1, 0,
```

```
50         (struct sockaddr*) &destAddr,  
51         sizeof(struct sockaddr)) == -1)  
52     {  
53         cerr << "Could not send to socket\n";  
54         exit(4);  
55     }  
56     close(sock);  
57     return 0;  
58 }
```



The image shows two terminal windows side-by-side. The left window, titled 'Terminal', shows a prompt '>' followed by the command 'udpStrSrv' (with a red box around the '1' in 'Srv'). The output is 'Hello from client' (with a red box around the '3' in 'client'). The prompt '>' is shown again. The right window, also titled 'Terminal', shows a prompt '>' followed by the command 'udpStrCli' (with a red box around the '2' in 'Cli'). The output is 'Hello 152.17.140.12' (with a red box around the '4' in '12'). The prompt '>' is shown again.