

Association Rule Mining

- Given a set of transactions, find rules that will predict the occurrence of an item based on the occurrences of other items in the transaction

Market-Basket transactions

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Example of Association Rules

$\{\text{Diaper}\} \rightarrow \{\text{Beer}\}$,
 $\{\text{Milk, Bread}\} \rightarrow \{\text{Eggs, Coke}\}$,
 $\{\text{Beer, Bread}\} \rightarrow \{\text{Milk}\}$,

Implication means co-occurrence, not causality!

Definition: Frequent Itemset

- Itemset**
 - A collection of one or more items
 - Example: {Milk, Bread, Diaper}
 - k-itemset
 - An itemset that contains k items
- Support count (σ)**
 - Frequency of occurrence of an itemset
 - E.g. $\sigma(\{\text{Milk, Bread, Diaper}\}) = 2$
- Support**
 - Fraction of transactions that contain an itemset
 - E.g. $s(\{\text{Milk, Bread, Diaper}\}) = 2/5$
- Frequent Itemset**
 - An itemset whose support is greater than or equal to a *minsup* threshold

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Definition: Association Rule

- Association Rule**
 - An implication expression of the form $X \rightarrow Y$, where X and Y are itemsets
 - Example: $\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Example:

$\{\text{Milk, Diaper}\} \Rightarrow \text{Beer}$

Rule Evaluation Metrics

- Support (s)**
 - Fraction of transactions that contain both X and Y
- Confidence (c)**
 - Measures how often items in Y appear in transactions that contain X

$$s = \frac{\sigma(\text{Milk, Diaper, Beer})}{|T|} = \frac{2}{5} = 0.4$$

$$c = \frac{\sigma(\text{Milk, Diaper, Beer})}{\sigma(\text{Milk, Diaper})} = \frac{2}{3} = 0.67$$

Association Rule Mining Task

- Given a set of transactions T, the goal of association rule mining is to find all rules having
 - support \geq *minsup* threshold
 - confidence \geq *minconf* threshold
- Brute-force approach:**
 - List all possible association rules
 - Compute the support and confidence for each rule
 - Prune rules that fail the *minsup* and *minconf* thresholds

\Rightarrow **Computationally prohibitive!**

Mining Association Rules

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Example of Rules:

$\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$ ($s=0.4, c=0.67$)
 $\{\text{Milk, Beer}\} \rightarrow \{\text{Diaper}\}$ ($s=0.4, c=1.0$)
 $\{\text{Diaper, Beer}\} \rightarrow \{\text{Milk}\}$ ($s=0.4, c=0.67$)
 $\{\text{Beer}\} \rightarrow \{\text{Milk, Diaper}\}$ ($s=0.4, c=0.67$)
 $\{\text{Diaper}\} \rightarrow \{\text{Milk, Beer}\}$ ($s=0.4, c=0.5$)
 $\{\text{Milk}\} \rightarrow \{\text{Diaper, Beer}\}$ ($s=0.4, c=0.5$)

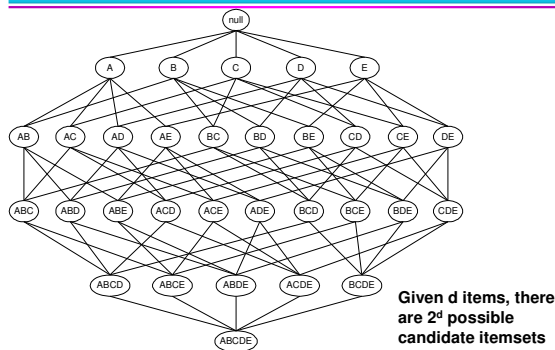
Observations:

- All the above rules are binary partitions of the same itemset: {Milk, Diaper, Beer}
- Rules originating from the same itemset have identical support but can have different confidence
- Thus, we may decouple the support and confidence requirements

Mining Association Rules


- Two-step approach:**
 - Frequent Itemset Generation**
 - Generate all itemsets whose support \geq *minsup*
 - Rule Generation**
 - Generate high confidence rules from each frequent itemset, where each rule is a binary partitioning of a frequent itemset
- Frequent itemset generation is still computationally expensive

Frequent Itemset Generation



Frequent Itemset Generation

- Brute-force approach:
 - Each itemset in the lattice is a **candidate** frequent itemset
 - Count the support of each candidate by scanning the database



- Match each transaction against every candidate
- Complexity $\sim O(NMw) \Rightarrow$ **Expensive since $M = 2^d$!!!**

Frequent Itemset Generation Strategies


- Reduce the **number of candidates** (M)
 - Complete search: $M = 2^d$
 - Use pruning techniques to reduce M
- Reduce the **number of transactions** (N)
 - Reduce size of N as the size of itemset increases
 - Used by DHP and vertical-based mining algorithms
- Reduce the **number of comparisons** (NM)
 - Use efficient data structures to store the candidates or transactions
 - No need to match every candidate against every transaction

Reducing Number of Candidates

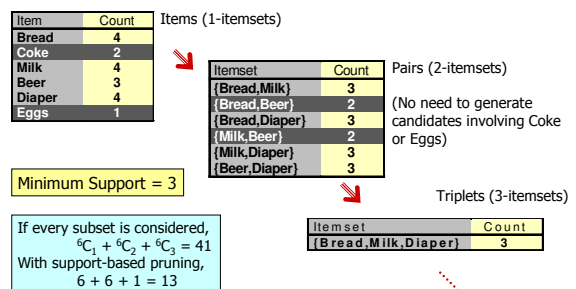
- Apriori principle:**
 - If an itemset is frequent, then all of its subsets must also be frequent
- Apriori principle holds due to the following property of the support measure:

$$\forall X, Y : (X \subseteq Y) \Rightarrow s(X) \geq s(Y)$$
 - Support of an itemset never exceeds the support of its subsets
 - This is known as the **anti-monotone** property of support

Illustrating Apriori Principle



Illustrating Apriori Principle



Frequent Itemset Mining


2 strategies:

- Breadth-first: Apriori
 - ◆ Exploit monotonicity to the maximum
- Depth-first strategy: Eclat
 - ◆ Prune the database
 - ◆ Do not fully exploit monotonicity

Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D


minsup=2



Apriori

1	B, C
2	B, C
3	A, C, D
4	A, B, C, D
5	B, D

minsup=2



Apriori Algorithm

Method:

- Let $k=1$
- Generate frequent itemsets of length 1
- Repeat until no new frequent itemsets are identified
 - Generate length $(k+1)$ candidate itemsets from length k frequent itemsets
 - Prune candidate itemsets containing subsets of length k that are infrequent
 - Count the support of each candidate by scanning the DB
 - Eliminate candidates that are infrequent, leaving only those that are frequent

Frequent Itemset Mining

2 strategies:

- Breadth-first: Apriori
 - Exploit monotonicity to the maximum
- Depth-first strategy: Eclat
 - Prune the database
 - Do not fully exploit monotonicity

Depth-First Algorithms

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

Depth-First Algorithms

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

without D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

with D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

Depth-First Algorithms

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

without D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC, BC

with D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC

Depth-First Algorithms

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

without D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC, BC

with D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC

add D again

AD, BD, CD, ACD

A, B, C, AC, BC

Depth-First Algorithms

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC, BC, AD, BD, CD, ACD

without D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC, BC

with D

Find all frequent itemsets

1		B, C
2		B, C
3	A,	C, D
4	A, B, C,	D
5		B, D

minsup=2

A, B, C, AC

AD, BD, CD, ACD

+ A, B, C, AC, BC

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, D
4	A, B, C, D
5	B, D

A: 2
 B: 4
 C: 4
 D: 3

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, D
4	A, B, C, D
5	B, D

A: 2
 B: 4
 C: 4
 D: 3

DB[D]	
3	A, C
4	A, B, C
5	B, C

A: 2
 B: 4
 C: 2
 AC: 2

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, D
4	A, B, C, D
5	B, D

A: 2
 B: 4
 C: 4
 D: 3

DB[D]	
3	A, B, C
4	A, B, C
5	B, C

A: 2
 B: 4
 C: 2

DB[CD]	
3	A, B
4	A, B

A: 2

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, D
4	A, B, C, D
5	B, D

A: 2
 B: 4
 C: 4
 D: 3

DB[D]	
3	A, B, C
4	A, B, C
5	B, C

A: 2
 B: 4
 C: 2
 AC: 2

DB[CD]	
3	A, B
4	A, B

A: 2

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, D
4	A, B, C, D
5	B, D

A: 2
 B: 4
 C: 4
 D: 3

DB[D]	
3	A, B, C
4	A, B, C
5	B, C

A: 2
 B: 4
 C: 2
 AC: 2

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, D
4	A, B, C, D
5	B, D

A: 2
 B: 4
 C: 4
 D: 3

DB[D]	
3	A, B, C
4	A, B, C
5	B, C

A: 2
 B: 4
 C: 2
 AC: 2

DB[BD]	
4	A

A: 1

Depth-First Algorithm

DB		DB[D]	
1	B, C	3	A, A
2	B, C	4	A, A
3	A, B, C, C	5	
4	A, B, C, C		
5	B, B		

A: 2
 B: 4
 C: 4
 D: 3
 AC: 2

Depth-First Algorithm

DB		DB[D]	
1	B, C	3	A, A
2	B, C	4	A, A
3	A, B, C, C	5	
4	A, B, C, C		
5	B, B		

A: 2
 B: 4
 C: 4
 D: 3
 AD: 2
 BD: 4
 CD: 2
 ACD: 2

Depth-First Algorithm

DB	
1	B, C
2	B, C
3	A, B, C, C
4	A, B, C, C
5	B, B

A: 2
 B: 4
 C: 4
 D: 3
 AD: 2
 BD: 4
 CD: 2
 ACD: 2

Depth-First Algorithm

DB		DB[C]	
1	B, C	1	B
2	B, C	2	B
3	A, B, C, C	3	A
4	A, B, C, C	4	A, B
5	B, B		

A: 2
 B: 4
 C: 4
 D: 3
 AD: 2
 BD: 4
 CD: 2
 ACD: 2

Depth-First Algorithm

DB		DB[C]		DB[BC]	
1	B, C	1	B	1	A
2	B, C	2	B	2	A
3	A, B, C, C	3	A	3	A
4	A, B, C, C	4	A, B	4	A
5	B, B				


A: 2
 B: 4
 C: 4
 D: 3
 AD: 2
 BD: 4
 CD: 2
 ACD: 2

Depth-First Algorithm


DB		DB[C]	
1	B, C	1	B
2	B, C	2	B
3	A, B, C, C	3	A
4	A, B, C, C	4	A, B
5	B, B		

A: 2
 B: 4
 C: 4
 D: 3
 AD: 2
 BD: 4
 CD: 2
 ACD: 2


Depth-First Algorithm




Depth-First Algorithm




Depth-First Algorithm



Depth-First Algorithm



Depth-First Algorithm



ECLAT

- For each item, store a list of transaction ids (tids)

Horizontal Data Layout

TID	Items
1	A,B,E
2	B,C,D
3	C,E
4	A,C,D
5	A,B,C,D
6	A,E
7	A,B
8	A,B,C
9	A,C,D
10	B


Vertical Data Layout

	A	B	C	D	E
1	1	2	2	1	
4	2	3	4	3	
5	5	4	5	6	
6	7	8	9		
7	8	10			
8					
9					

TID-list

ECLAT

- Determine support of any k-itemset by intersecting tid-lists of two of its (k-1) subsets.



- Depth-first traversal of the search lattice
- Advantage: very fast support counting
- Disadvantage: intermediate tid-lists may become too large for memory

Rule Generation

- Given a frequent itemset L, find all non-empty subsets $f \subset L$ such that $f \rightarrow L - f$ satisfies the minimum confidence requirement
 - If $\{A, B, C, D\}$ is a frequent itemset, candidate rules:

$ABC \rightarrow D,$	$ABD \rightarrow C,$	$ACD \rightarrow B,$	$BCD \rightarrow A,$
$A \rightarrow BCD,$	$B \rightarrow ACD,$	$C \rightarrow ABD,$	$D \rightarrow ABC$
$AB \rightarrow CD,$	$AC \rightarrow BD,$	$AD \rightarrow BC,$	$BC \rightarrow AD,$
$BD \rightarrow AC,$	$CD \rightarrow AB,$		
- If $|L| = k$, then there are $2^k - 2$ candidate association rules (ignoring $L \rightarrow \emptyset$ and $\emptyset \rightarrow L$)

Rule Generation

- How to efficiently generate rules from frequent itemsets?
 - In general, confidence does not have an anti-monotone property
 - $c(ABC \rightarrow D)$ can be larger or smaller than $c(AB \rightarrow D)$
 - But confidence of rules generated from the same itemset has an anti-monotone property
 - e.g., $L = \{A, B, C, D\}$:

$$c(ABC \rightarrow D) \geq c(AB \rightarrow CD) \geq c(A \rightarrow BCD)$$
 - Confidence is anti-monotone w.r.t. number of items on the RHS of the rule

Rule Generation for Apriori Algorithm

