

## Problems of Varying Size

Small-scale problems:

- “anything goes,”
- no problem to use SVD (recommended).

Medium-size problems:

- cannot ignore computing time,
- other factorizations, sparse matrix aspects, etc.

Large-scale problems:

- storage and computing time set the limitations,
- factorizations are not possible in general,
- if possible, use matrix structure (Toeplitz, Kronecker, ...),
- otherwise must use iterative methods!

## Advantages of Iterative Methods

Iterative methods produce a sequence  $x^{[0]} \rightarrow x^{[1]} \rightarrow x^{[2]} \rightarrow \dots$  of iterates that (hopefully) converge to the desired solution, solely through the use of matrix-vector multiplications.

- The matrix  $A$  is never altered, only “touched” via matrix-vector multiplications  $Ax$  and  $A^T y$ .
- The matrix  $A$  is not explicitly required – we only need a “black box” that computes the action of  $A$  or the underlying operator.
- Atomic operations of iterative methods (mat-vec product, saxpy, norm) suited for high-performance computing.
- Often produce a natural sequence of regularized solutions; stop when the solution is “satisfactory” (parameter choice).

## Two Types of Iterative Methods

1. Iterative solution of a regularized problem, such as Tikhonov

$$(A^T A + \lambda^2 L^T L) x = A^T b .$$

Challenge: to construct a good preconditioner!

2. Iterate on the un-regularized system, e.g., on

$$A x = b \quad \text{or} \quad A^T A x = A^T b$$


and use the iteration number as the regularization parameter.

The latter approach relies on *semi-convergence*:

- initial convergence towards  $x^{\text{exact}}$ ,
- followed by (slow) convergence to  $x_{\text{naive}}$ .

Must stop at the end of the first stage!

## Illustration of Semi-Convergence



## Landweber Iteration

A classical stationary iterative method:

$$x^{[k+1]} = x^{[k]} + \omega A^T(b - Ax^{[k]}), \quad k = 0, 1, 2, \dots$$

where  $0 < \omega < 2 \|A^T A\|_2^{-1} = 2 \sigma_1^{-2}$ .

Where does this come from? Consider the function


$$\phi(x) = \frac{1}{2} \|b - Ax\|_2^2$$

associated with the least squares problem  $\min_x \phi(x)$ . It is straightforward (but perhaps a bit tedious) to show that the gradient of  $\phi$  is

$$\nabla \phi(x) = -A^T(b - Ax).$$

Thus, each step in Landweber's method is a step in the direction of steepest descent. See next slide for an example of iterations.

## The Geometry of Landweber Iterations



## SVD Analysis of Landweber's Method

SVD analysis shows that the filter factors (see next page) are:

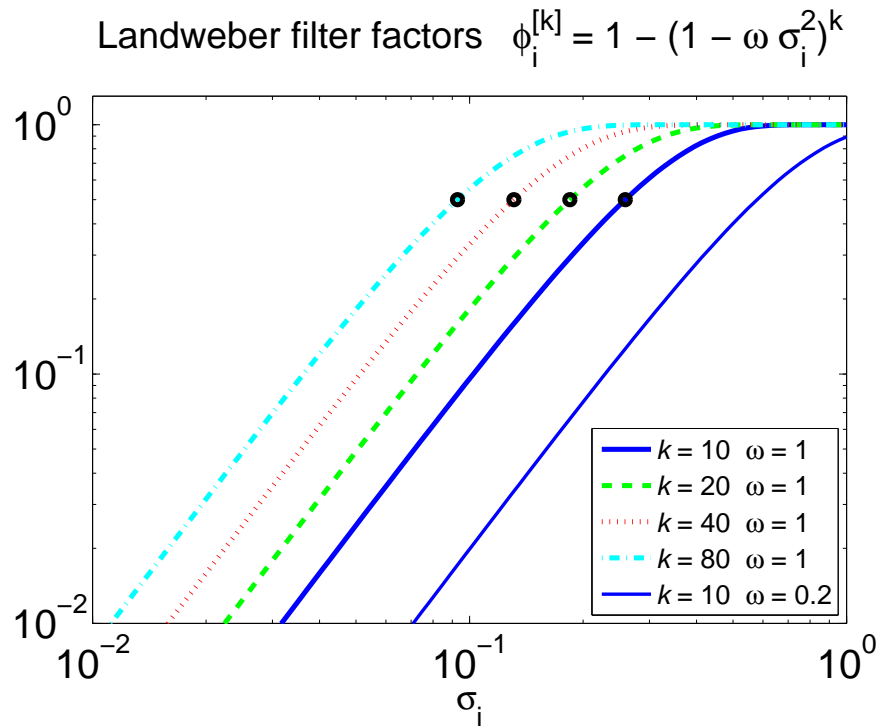
$$\phi_i^{[k]} = 1 - (1 - \omega \sigma_i^2)^k.$$

Let  $\sigma_{\text{break}}^{[k]}$  denote the value of  $\sigma_i$  for which  $f_i^{[k]} = 0.5$ . Then

$$\frac{\sigma_{\text{break}}^{[k]}}{\sigma_{\text{break}}^{[2k]}} = \sqrt{1 + \left(\frac{1}{2}\right)^{\frac{1}{2k}}} \rightarrow \sqrt{2} \quad \text{for } k \rightarrow \infty.$$

Hence, as  $k$  increases, the breakpoint tends to be reduced by a factor  $\sqrt{2} \approx 1.4$  each time the number of iterations  $k$  is doubled.

## Landweber Filter Factors



## Cimmino Iteration

Cimmino's method is a variant of Landweber's method, with a diagonal scaling:

$$x^{[k+1]} = x^{[k]} + \omega A^T D (b - A x^{[k]}), \quad k = 1, 2, \dots$$

in which  $D = \text{diag}(d_i)$  is a diagonal matrix whose elements are defined in terms of the rows  $a_i^T = A(i, :)$  of  $A$  as

$$d_i = \begin{cases} \frac{1}{m} \frac{1}{\|a_i\|_2^2}, & a_i \neq 0 \\ 0, & a_i = 0. \end{cases}$$

Landweber and Cimmino belong to a class of iterative methods called Simultaneous Iterative Reconstruction Techniques (SIRT).

## ... and the prize for best acronym goes to “ART”

Kaczmarz's method = algebraic reconstruction technique (ART).

Let  $a_i^T = A(i, :) = i$ th row of  $A$ , and  $b_i = i$ th component  $b$ .

Each iteration of ART involves the following “sweep” over all rows:

$$x^{[k^{(0)}]} = x^{[k]}$$

for  $i = 1, \dots, m$

$$x^{[k^{(i)}]} = x^{[k^{(i-1)}]} + \frac{b_i - a_i^T x^{[k^{(i-1)}]}}{\|a_i\|_2^2} a_i$$


end

$$x^{[k+1]} = x^{[k^{(m)}]}$$


This method is not “simultaneous” because each row must be processed sequentially.

In general: fast initial convergence, then slow. See next slides.

## The Geometry of ART Iterations



## Slow Convergence of SIRT and ART Methods



The test problem is shaw.

## Projection Methods

As an important step towards the faster *Krylov subspace methods*, we consider projection methods.

Assume the columns of  $W_k = (w_1, \dots, w_k) \in \mathbb{R}^{n \times k}$  form a “good basis” for an approximate regularized solution, obtained by solving

$$\min_x \|A x - b\|_2 \quad \text{s.t.} \quad x \in \mathcal{W}_k = \text{span}\{w_1, \dots, w_k\}.$$

This solution takes the form

$$x^{(k)} = W_k y^{(k)}, \quad y^{(k)} = \operatorname{argmin}_y \|(A W_k) y - b\|_2,$$

and we refer to the least squares problem  $\|(A W_k) y - b\|_2$  as the *projected problem*, because it is obtained by projecting the original problem onto the  $k$ -dimensional subspace  $\text{span}(w_1, \dots, w_k)$ .

If  $W_k = V_k$  then we obtain the TSVD method, and  $x^{(k)} = x_k$

But we want to work with computationally simpler basis vectors.

## Computations with DCT Basis

Note that

$$\hat{A}_k = A W_k = (W_k^T A^T)^T = [(W^T A^T)^T]_{:,1:k}.$$

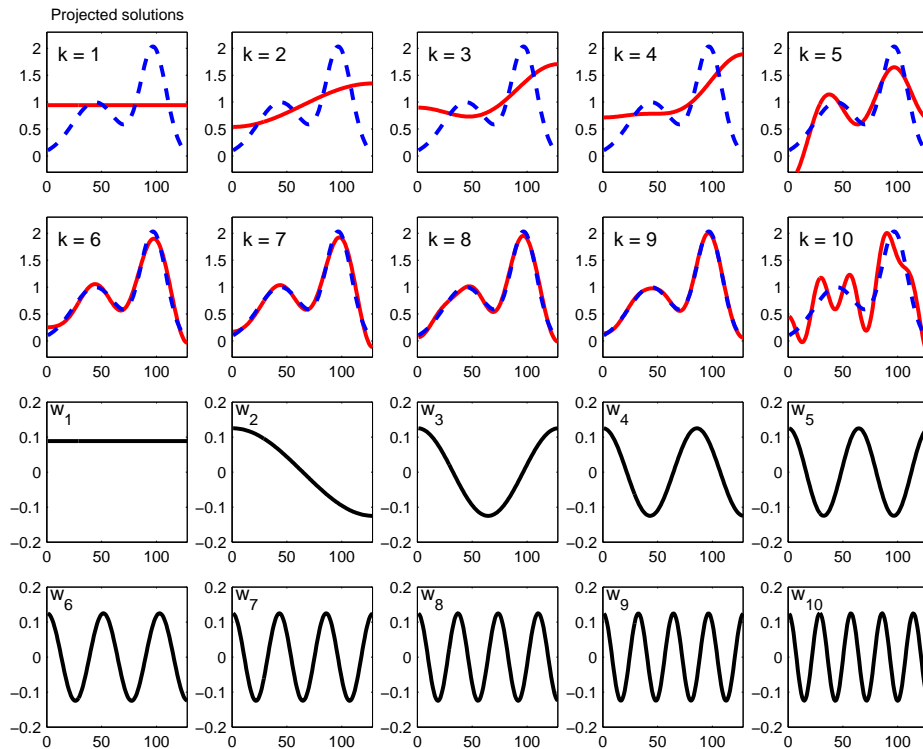
In the case of the discrete cosine basis, multiplication with  $W^T$  is equivalent to a DCT. The algorithm takes the form:

```
Akhat = dct(A')';
Akhat = Akhat(:,1:k);
y = Akhat\b;
xk = idct([y;zeros(n-k,1)]);
```

Next page:

- Top: solutions  $x^{(k)}$  for  $k = 1, \dots, 10$ .
- Bottom: cosine basis  $w_i$ ,  $i = 1, \dots, 10$ .

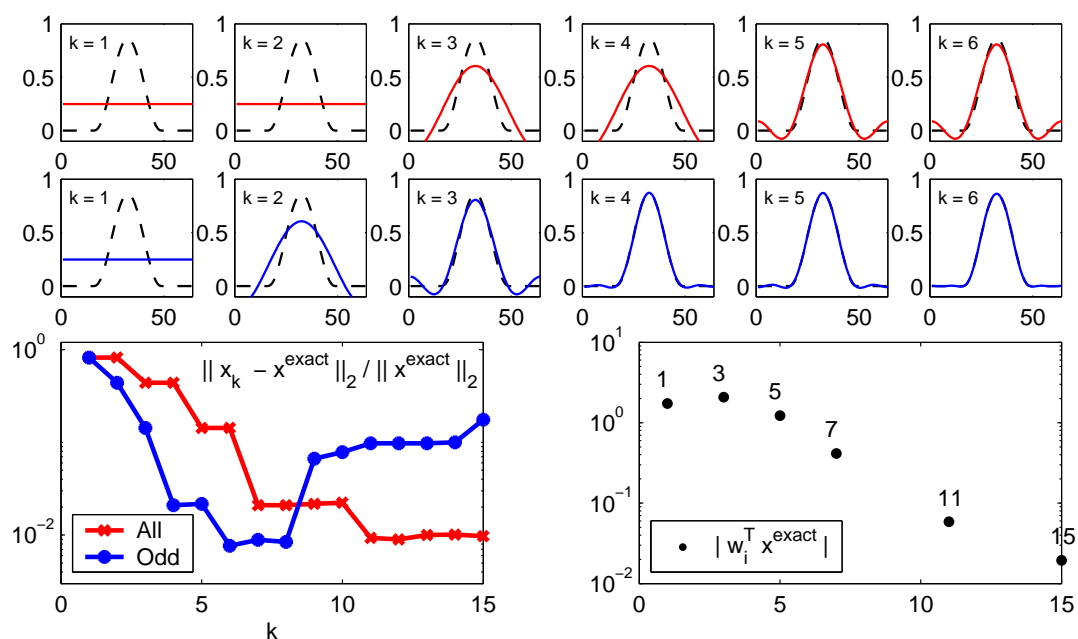
## Example Using Discrete Cosine Basis (shaw)



## Symmetric Solution and DCT (phillips)

Red: Using all the DCT basis vectors  $w_1, w_2, w_3, w_4, w_5, w_6, \dots$

Blue: Using only the odd-numbered DCT vectors  $w_1, w_3, w_5, \dots$





## The Krylov Subspace

The DCT basis is sometimes a good basis – but not always.

The **Krylov subspace**, defined as

$$\mathcal{K}_k \equiv \text{span}\{A^T b, A^T A A^T b, (A^T A)^2 A^T b, \dots, (A^T A)^{k-1} A^T b\},$$

always *adapts* itself to the problem at hand! But the “naive” basis,

$$p_i = (A^T A)^{i-1} A^T b / \|(A^T A)^{i-1} A^T b\|_2, \quad i = 1, 2, \dots$$

are NOT useful:  $p_i \rightarrow v_1$  as  $i \rightarrow \infty$ . Use modified Gram-Schmidt:

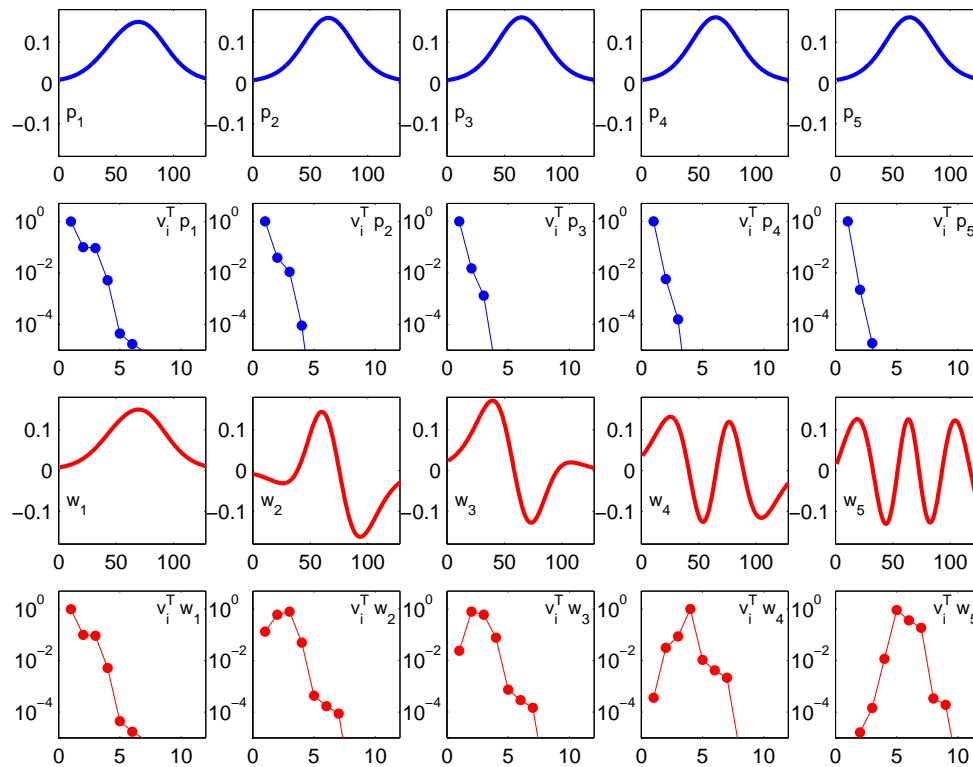
$$w_1 \leftarrow A^T b; \quad w_1 \leftarrow w_1 / \|w_1\|_2$$

$$w_2 \leftarrow A^T A w_1; \quad w_2 \leftarrow w_2 - w_1^T w_2 w_1; \quad w_2 \leftarrow w_2 / \|w_2\|_2$$

$$w_3 \leftarrow A^T A w_2; \quad w_3 \leftarrow w_3 - w_1^T w_3 w_1;$$

$$w_3 \leftarrow w_3 - w_2^T w_3 w_2; \quad w_3 \leftarrow w_3 / \|w_3\|_2$$

Comparison of basis vectors  $p_i$  (blue) and  $w_i$  (red).



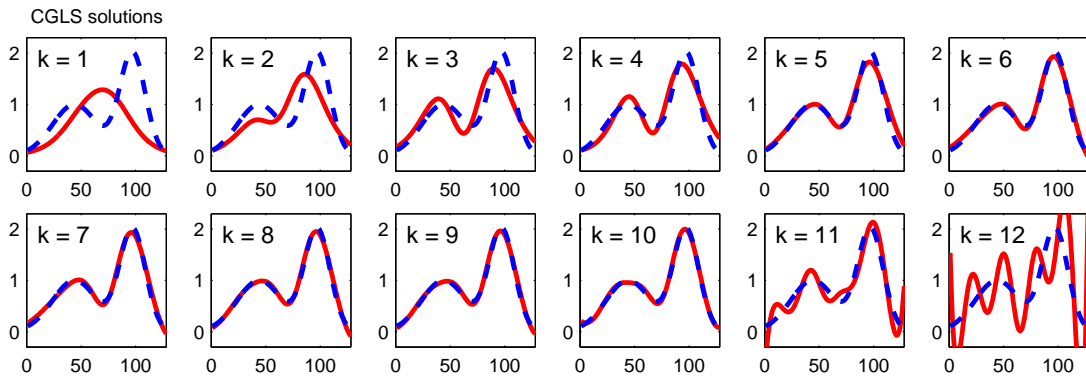
## Can We Compute $x^{(k)}$ Without Storing $W_k$ ?

Yes – the CGLS algorithm computes iterates given by

$$x^{(k)} = \operatorname{argmin}_x \|Ax - b\|_2 \quad \text{s.t.} \quad x \in \mathcal{K}_k.$$

The algorithm eventually converges to the least squares solution.

But since  $\mathcal{K}_k$  is a good subspace for approximate regularized solutions, CGLS exhibits semi-convergence.



## CGLS = Conjugate Gradients for Least Squares

The CGLS algorithm takes the following form:

$x^{(0)}$  = starting vector (e.g., zero)

$$r^{(0)} = b - Ax^{(0)}$$

$$d^{(0)} = A^T r^{(0)}$$

for  $k = 1, 2, \dots$

$$\bar{\alpha}_k = \|A^T r^{(k-1)}\|_2^2 / \|A d^{(k-1)}\|_2^2$$

$$x^{(k)} = x^{(k-1)} + \bar{\alpha}_k d^{(k-1)}$$

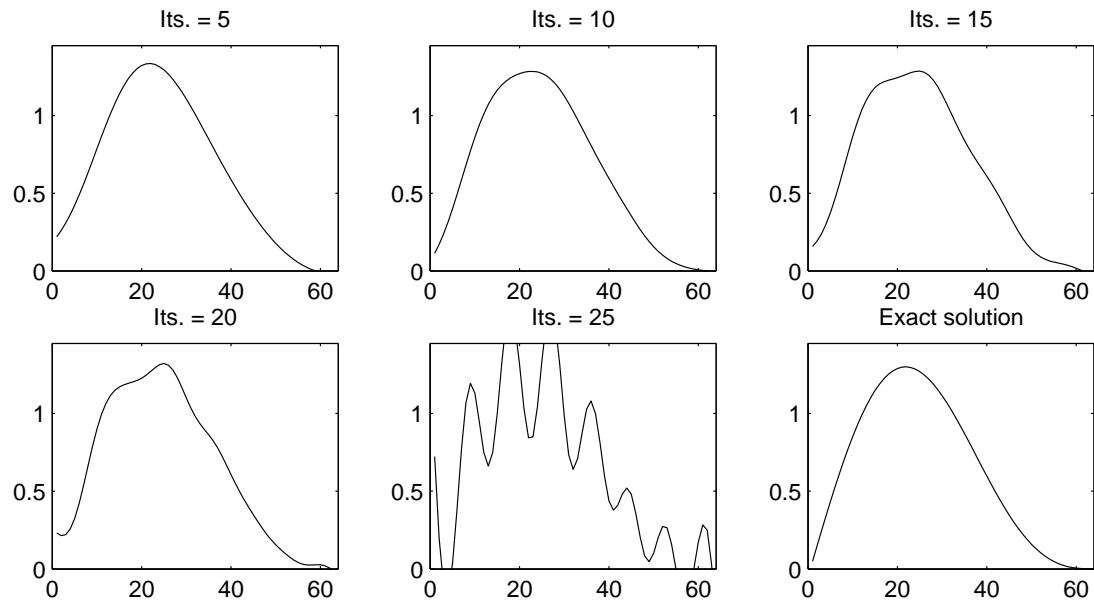
$$r^{(k)} = r^{(k-1)} - \bar{\alpha}_k A d^{(k-1)}$$

$$\bar{\beta}_k = \|A^T r^{(k)}\|_2^2 / \|A^T r^{(k-1)}\|_2^2$$

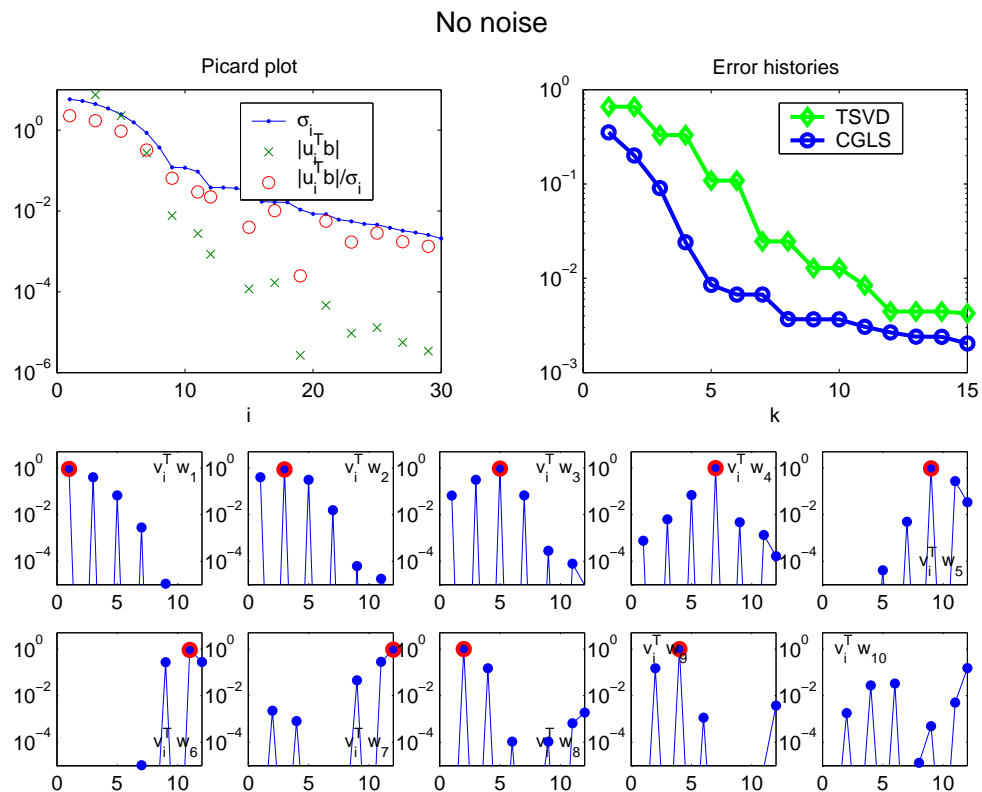
$$d^{(k)} = A^T r^{(k)} + \bar{\beta}_k d^{(k-1)}$$

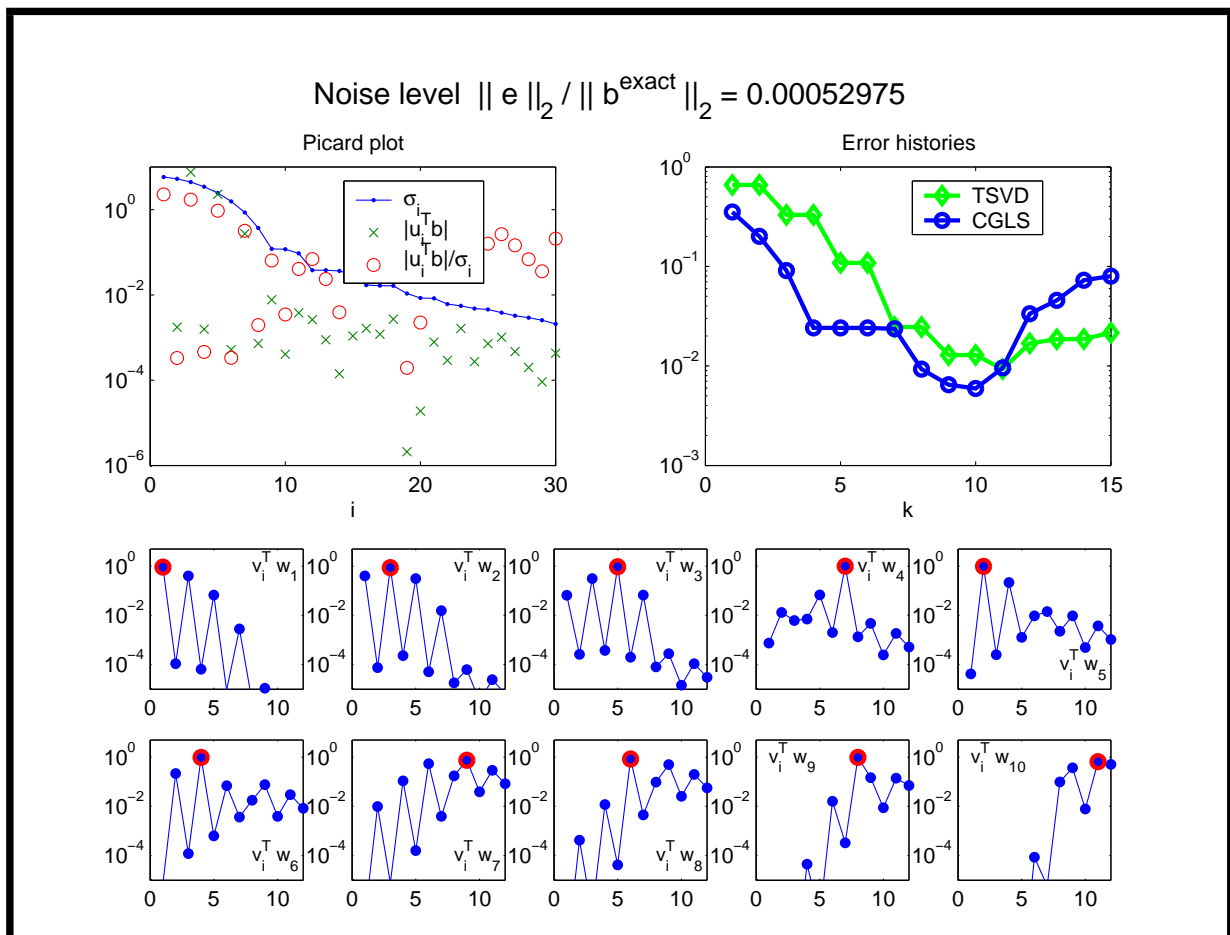
end

## CGLS Solutions to the Gravity Problem



## CGLS Focuses on the Significant Components





## Other Iterations – GMRES and RRGMRES

Sometimes difficult or inconvenient to write a matrix-free black-box function for multiplication with  $A^T$ . Can we avoid this?

The GMRES method for square nonsymmetric matrices is based on the Krylov subspace

$$\mathcal{K}_k = \text{span}\{b, Ab, A^2b, \dots, A^{k-1}b\}.$$

The presence of the noisy data  $b = b^{\text{exact}} + e$  in this subspace is unfortunate: the solutions include the noise component  $e$ !

A better subspace, underlying the RRGMRES method:

$$\vec{\mathcal{K}}_k = \text{span}\{Ab, A^2b, \dots, A^k b\}.$$

Now the noise vector is multiplied with  $A$  (smoothing) at least once.

Symmetric matrices: use MR-II (a simplified variant).

## Tomography (a Case for Iterative Methods) §7.7

Tomography is the science of computing reconstructions from projections, i.e., data obtained by integrations along rays (typically straight lines) that penetrate a domain  $\Omega$ .

The unknown function  $f(\mathbf{t}) = f(t_1, t_2)$  represents some material parameter, and the damping of a signal penetrating a part  $d\tau$  of a ray at position  $\mathbf{t}$  is proportional to the product to  $f(\mathbf{t}) d\tau$ .

The data consist of measurements of the damping of signals following well-defined rays through the domain  $\Omega$ .

## Formulation of Tomography Problem

The  $i$ th observation  $b_i$ ,  $i = 1, \dots, m$ , represents the damping of a signal that penetrates  $\Omega$  along a straight line, ray <sup>$i$</sup> .

All the point  $\mathbf{t}^i$  on ray <sup>$i$</sup>  are given by

$$\mathbf{t}^i(\tau) = \mathbf{t}^{i,0} + \tau \mathbf{d}^i, \quad \tau \in \mathbb{R},$$

where  $\mathbf{t}^{i,0}$  is an arbitrary point on the ray, and  $\mathbf{d}^i$  is a (unit) vector that points in the direction of the ray.

The damping is proportional to the integral of the function  $f(\mathbf{t})$  along the ray. Specifically, for the  $i$ th observation, the damping associated with the  $i$ th ray is given by

$$b_i = \int_{-\infty}^{\infty} f(\mathbf{t}^i(\tau)) d\tau, \quad i = 1, \dots, m,$$

where  $d\tau$  denotes the integration along the ray.

## Discretization of 2D Tomography Problem

We consider a square domain  $\Omega = [0, 1] \times [0, 1]$ .

We can discretize this problem by dividing  $\Omega$  into an  $N \times N$  array of pixels, and in each pixel  $(k, \ell)$  we assume that the function  $f(\mathbf{t})$  is a constant  $f_{k\ell}$ :


$$f(\mathbf{t}) = f_{k\ell} \quad \text{for} \quad t_1 \in I_k \text{ \& } t_2 \in I_\ell,$$

where we have defined the interval  $I_k = [(k-1)/N, k/N]$ ,  $k = 1, \dots, N$  (and similarly for  $I_\ell$ ).

With this assumption about  $f(\mathbf{t})$  being piecewise constant, the expression for the  $k$ th measurement takes the simpler form

$$b_i = \sum_{(k,\ell) \in \text{ray}^i} f_{k\ell} \Delta L_{k\ell}^{(i)}, \quad \Delta L_{k\ell}^{(i)} = \text{length of ray}_i \text{ in pixel } (k, \ell)$$

for  $i = 1, \dots, m$ .



## Arriving at the System of Linear Equations

The above equation is, in fact, a system of linear equations in the  $N^2$  unknowns  $f_{k\ell}$ . We introduce the vector  $x$  of length  $n = N^2$  whose elements are the (unknown) function values  $f_{k\ell}$ :

$$x_\ell = f_{k\ell}, \quad \ell = (k-1)N + \ell.$$

This corresponds to stacking the columns of the  $N \times N$  matrix  $F$ . Moreover we organize the measurements  $b_i$  into a vector  $b$ .

There is clearly a linear relationship between the data  $b_k$  and the unknowns in the vector  $x$ , meaning that we can always write

$$b_i = \sum_{j=1}^n a_{ij} x_j, \quad i = 1, \dots, m.$$

This is a system of linear equations  $Ax = b$  with an  $m \times n$  matrix.

The elements of  $A$  are given by

$$a_{ij} = \begin{cases} \Delta L_{k\ell}^{(i)}, & (k, \ell) \in \text{ray}_i \\ 0 & \text{else} \end{cases}$$

where index  $i$  denotes the  $i$ th observation and  $j$  denotes the pixel number in an ordering with  $j = (k-1)N + \ell$ . The matrix  $A$  is very sparse.

